

# CS150: Database & Datamining

## Lecture 8: The IO Model

ShanghaiTech-SIST


Spring 2019

*Acknowledgement: Slides are adopted from the Berkeley course CS186 by Joey Gonzalez and Joe Hellerstein, Stanford CS145 by Peter Bailis.*

# Transition to **Mechanisms**

1. So you can **understand** what the database is doing!
  1. Understand the CS challenges of a database and how to use it.
  2. Understand how to optimize a query
  
2. Many **mechanisms** have become **stand-alone systems**
  - **Indexing** to Key-value stores
  - Embedded join processing
  - SQL-like languages take some aspect of what we discuss (PIG, Hive)

# Block diagram of a DBMS




# Today's Lecture

1. The Disk and Files
2. The Buffer

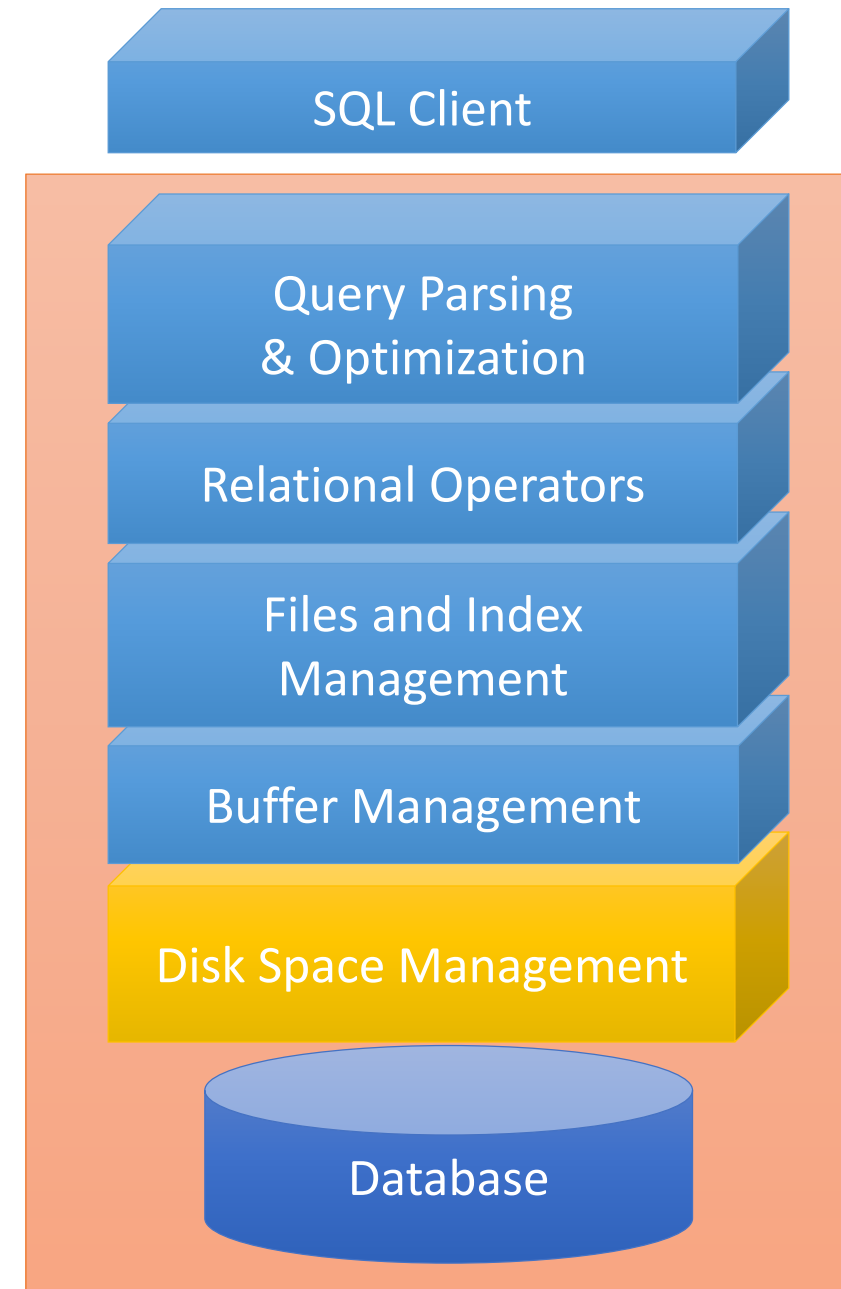

# 1. The Disk and Files

# Architecture of a DBMS

Translates page requests into physical bytes on one or more device(s)



Disk Space Mngmt.




# Overview

## Record


Bob	Harmon	M	32	94703
Varchar	Varchar	Char	Int	Int



## Byte Rep. Record



## Slotted Page




## Table

Name	Addr	Sex	Age	Zip
Bob	Harmon	M	32	94703
Alice	Mabel	F	33	94703
Jose	Chavez	M	31	94110
Jane	Chavez	F	30	94110



## File




# Overview: Files of Pages of Records

- Tables stored as a *logical files* consisting of *pages* each containing a collection of *records*
- Pages are managed
  - *in memory* by the *buffer manager*: higher levels of database only operate in memory
  - *on disk* by the *disk space manager*: reads and writes pages to physical disk/files



# A Simplified Filesystem Model

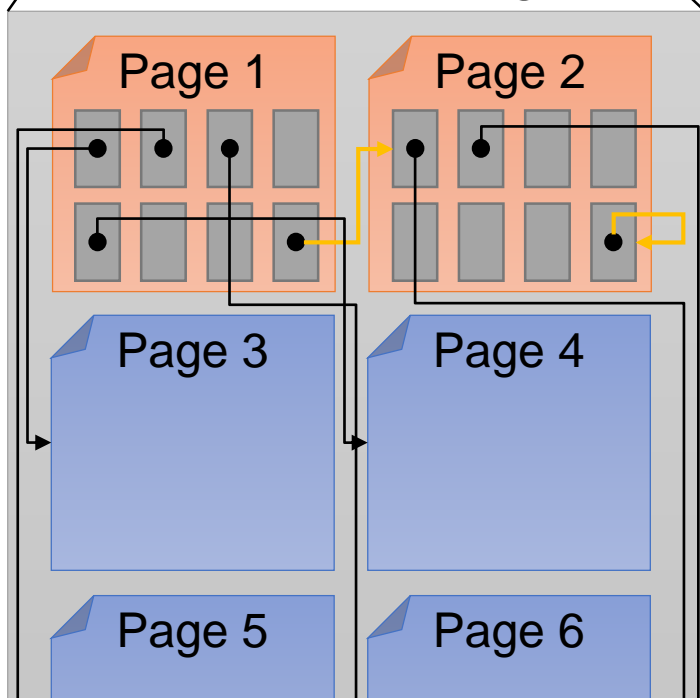
- For us, a **page** is a ***fixed-sized array*** of memory
  - Think: One or more disk blocks
  - Interface:
    - write to an entry (called a **slot**) or set to “None”
  - DBMS also needs to handle variable length fields
    - Page layout is important for good hardware utilization as well (see next next lecture)
- And a **file** is a ***variable-length list*** of pages
  - Interface: create / open / close; next\_page(); etc.




# Disk Space Management



Disk Space Mngmt.




## Buffer Management



Database  
operates  
on **in  
memory  
pages.**

# Recall: Disks and Files

- DBMS stores information on Disks and SSDs.
  - Disks are a mechanical anachronism (slow!)
    - 10ms Random Read Latency



# Recall: Arranging Pages on Disk

- “*Next*” page concept:
  - pages on same track, followed by
  - pages on same cylinder, followed by
  - pages on adjacent cylinder
- Arrange file pages sequentially on disk
  - minimize seek and rotational delay.
- For a *sequential scan*, *pre-fetch*
  - several pages at a time!
- Read large consecutive blocks

# Disks and Files

- DBMS stores information on **Disks** and **SSDs**.
  - Disks are a mechanical **anachronism** (slow!)
  - SSDs faster, **slow relative to memory**, costly writes
- DBMS operate at **Block Level**
  - Read and Write large **chunks seq. bytes**
    - Leverage cache hierarchy and HW pre-fetch
    - Amortize seek delays on HDDs and Writes on SSD
  - *Sequentially: Next* disk block is fastest
  - Maximize usage of data per R/W
- Organize data for fast in memory processing (i.e., mapping)

# Disk Space Management

Lowest layer of DBMS, manages space on disk

- Mapping pages to locations on disk
- Loading pages from disk to memory
- Saving pages back to disk & ensuring writes

Higher levels call upon this layer to:

- read/write a pages
- allocate/de-allocate logical pages

Request for a *sequence* of pages best satisfied by pages stored sequentially on disk

- Physical details hidden from higher levels of system
- Higher levels may assume **Next Page** is fast!

# Disk Space Management Implementation

## Proposal 1: Talk to the device directly

- Could be very fast if you knew the device well
- What happens when devices change?

## Proposal 2: Run over filesystem (FS)

- Allocate single large “contiguous” file and assume sequential / nearby byte access are fast
- Most FS optimize for sequential access and temporal locality (buffer cache on hot items)
  - Sometimes disable FS buffering
- May span multiple files on multiple disks / machines

# Typically sits on top of local file system


Get Page 4

Get Page 5






# Disk Space Management



Disk Space Mngmt.

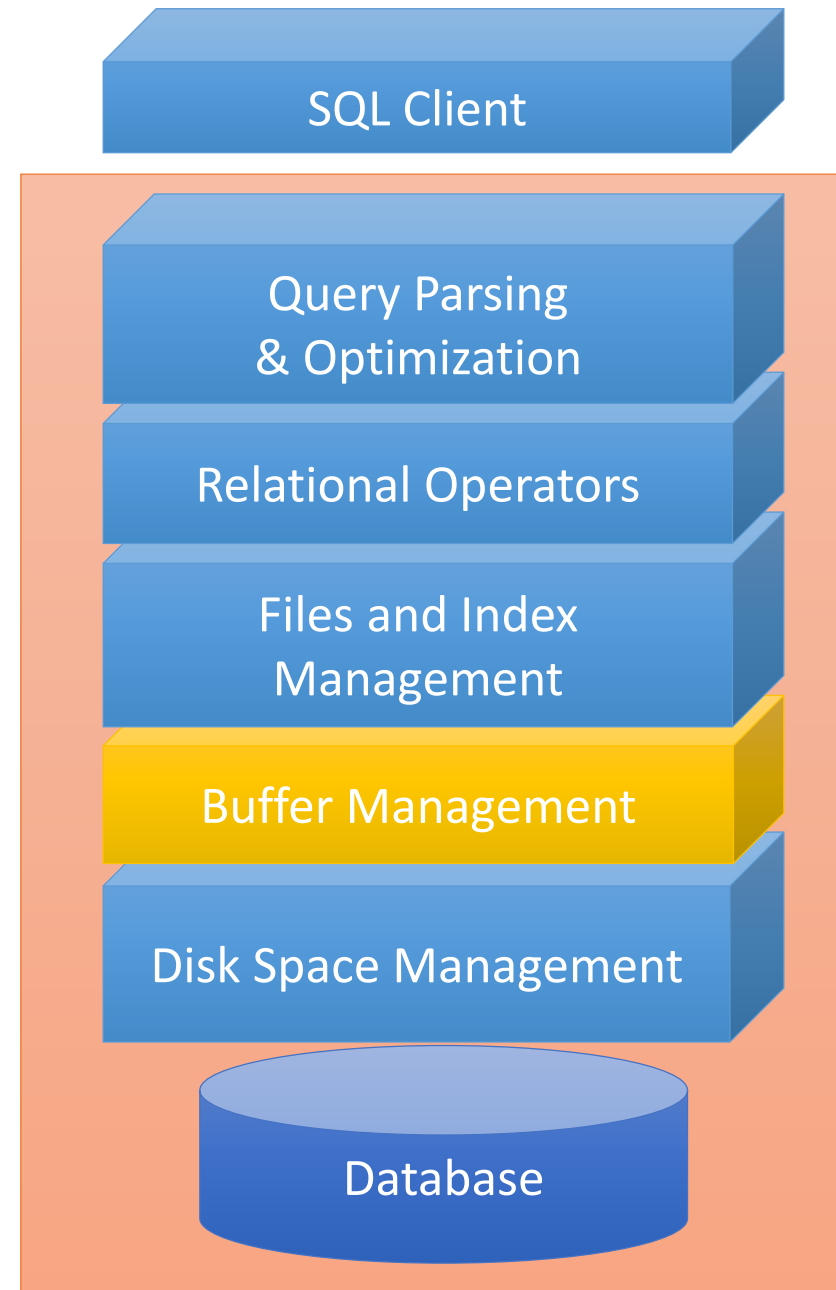



- Provide API to read and write pages to device
- Pages: block level organization of bytes on disk
- Ensures next locality and abstracts FS/Device details

## 2. The Buffer

# Architecture of a DBMS


Illusion of operating in memory



# What you will learn about in this section

1. RECAP: Storage and memory model
2. Buffer primer

# High-level: Disk vs. Main Memory



## Disk:


- **Slow:** Sequential *block* access
  - Read a blocks (not byte) at a time, so sequential access is cheaper than random
  - **Disk read / writes are expensive!**
- **Durable:** We will assume that once on disk, data is safe!
- **Cheap**

## Random Access Memory (RAM) or Main Memory:

- **Fast:** Random access, byte addressable
  - ~10x faster for sequential access
  - ~100,000x faster for random access!
- **Volatile:** Data can be lost if e.g. crash occurs, power goes out, etc!
- **Expensive:** For \$100, get 16GB of RAM vs. 2TB of disk!


# The Buffer

- A **buffer** is a region of physical memory used to store *temporary data*
  - *In this lecture:* a region in main memory used to store **intermediate data between disk and processes**
- *Key idea:* Reading / writing to disk is slow - need to cache data!



# The (Simplified) Buffer


- In this class: We'll consider a buffer located in **main memory** that operates over **pages** and **files**:
- **Read(page)**: Read page from disk -> buffer *if not already in buffer*



# The (Simplified) Buffer

- In this class: We'll consider a buffer located in **main memory** that operates over **pages** and **files**:
- **Read(page)**: Read page from disk -> buffer *if not already in buffer*


Processes can then read from / write to the page in the buffer






# The (Simplified) Buffer

- In this class: We'll consider a buffer located in **main memory** that operates over **pages** and **files**:
  - **Read(page)**: Read page from disk -> buffer *if not already in buffer*
  - **Flush(page)**: Evict page from buffer & write to disk




# The (Simplified) Buffer

- In this class: We'll consider a buffer located in **main memory** that operates over **pages** and **files**:
  - **Read(page)**: Read page from disk -> buffer *if not already in buffer*
  - **Flush(page)**: Evict page from buffer & write to disk
  - **Release(page)**: Evict page from buffer *without* writing to disk



# Managing Disk: The DBMS Buffer

- Database maintains its own buffer
  - Why? The OS already does this...
  - DB knows more about access patterns.
    - Watch for how this shows up! (cf. *Sequential Flooding*)
  - Recovery and logging require ability to **flush** to disk.



# When a Page is Requested ...

- Buffer pool information “table” contains:  
    <frame#, pageid, pin\_count, dirty>

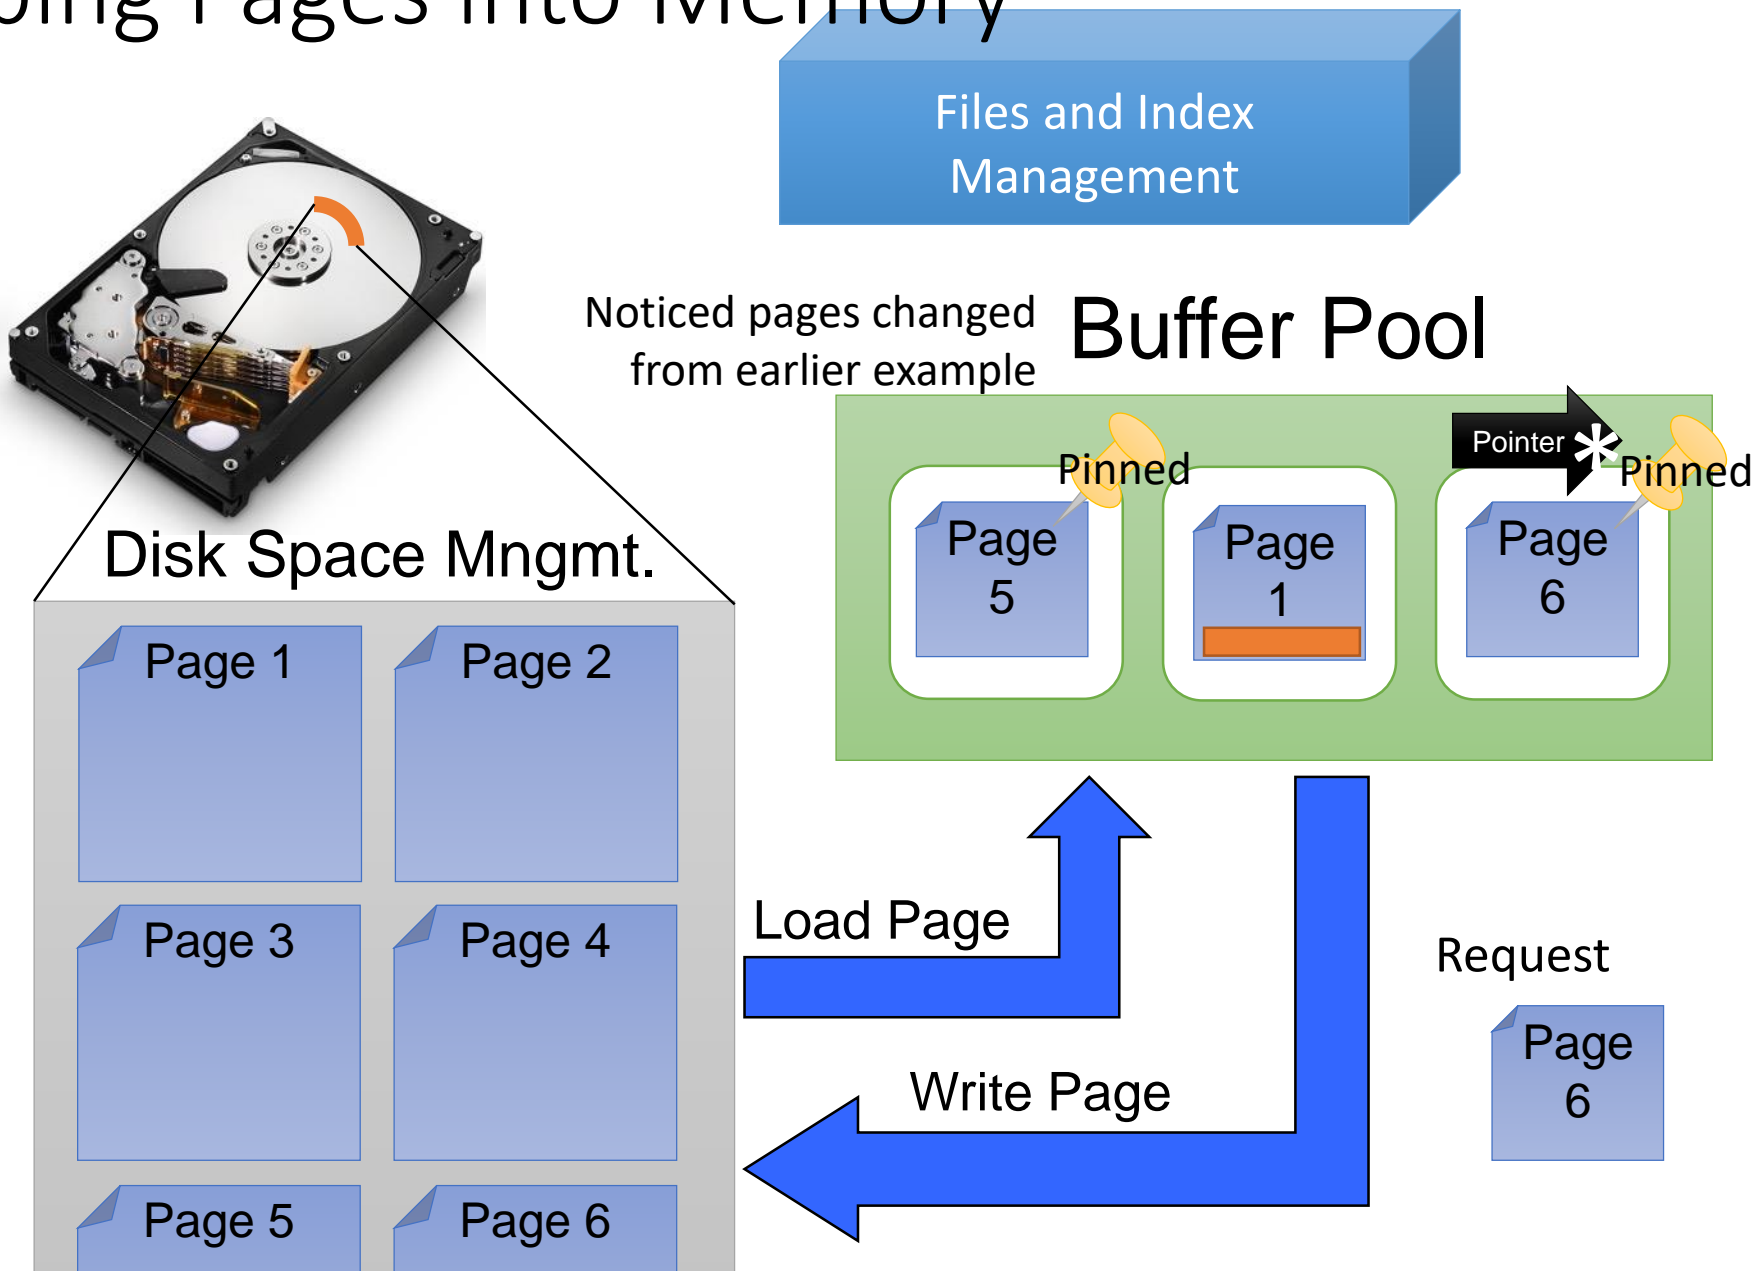
1. If requested page is not in pool:

- a. Choose a frame for *replacement*.  
    *Only “un-pinned” pages are candidates!*
- b. If frame “dirty”, write current page to disk
- c. Read requested page into frame

2. *Pin* the page and return its address.

If requests can be predicted (e.g., sequential scans)  
pages can be pre-fetched several pages at a time!


# Mapping Pages into Memory



# After Requestor Finishes


- Requestor of page must:
  1. indicate whether page was modified via *dirty* bit.
  2. *unpin* it (soon preferably!) why?
- Page in pool may be requested many times,
  - a *pin count* is used.
  - To pin a page: `pin_count++`
  - A page is a candidate for replacement iff *pin count* == 0 (“unpinned”)
- CC & recovery may do additional I/Os upon replacement.
  - *Write-Ahead Log* protocol; more later!

# Mapping Pages into Memory





Noticed pages changed  
from earlier example

## Buffer Pool



## Disk Space Mngmt.



Finished



# Page Replacement Policy

- Page is chosen for replacement by a *replacement policy*:
  - Least-recently-used (LRU), Clock
  - Most-recently-used (MRU)
- Policy can have big impact on #I/O's;
  - Depends on the *access pattern*.



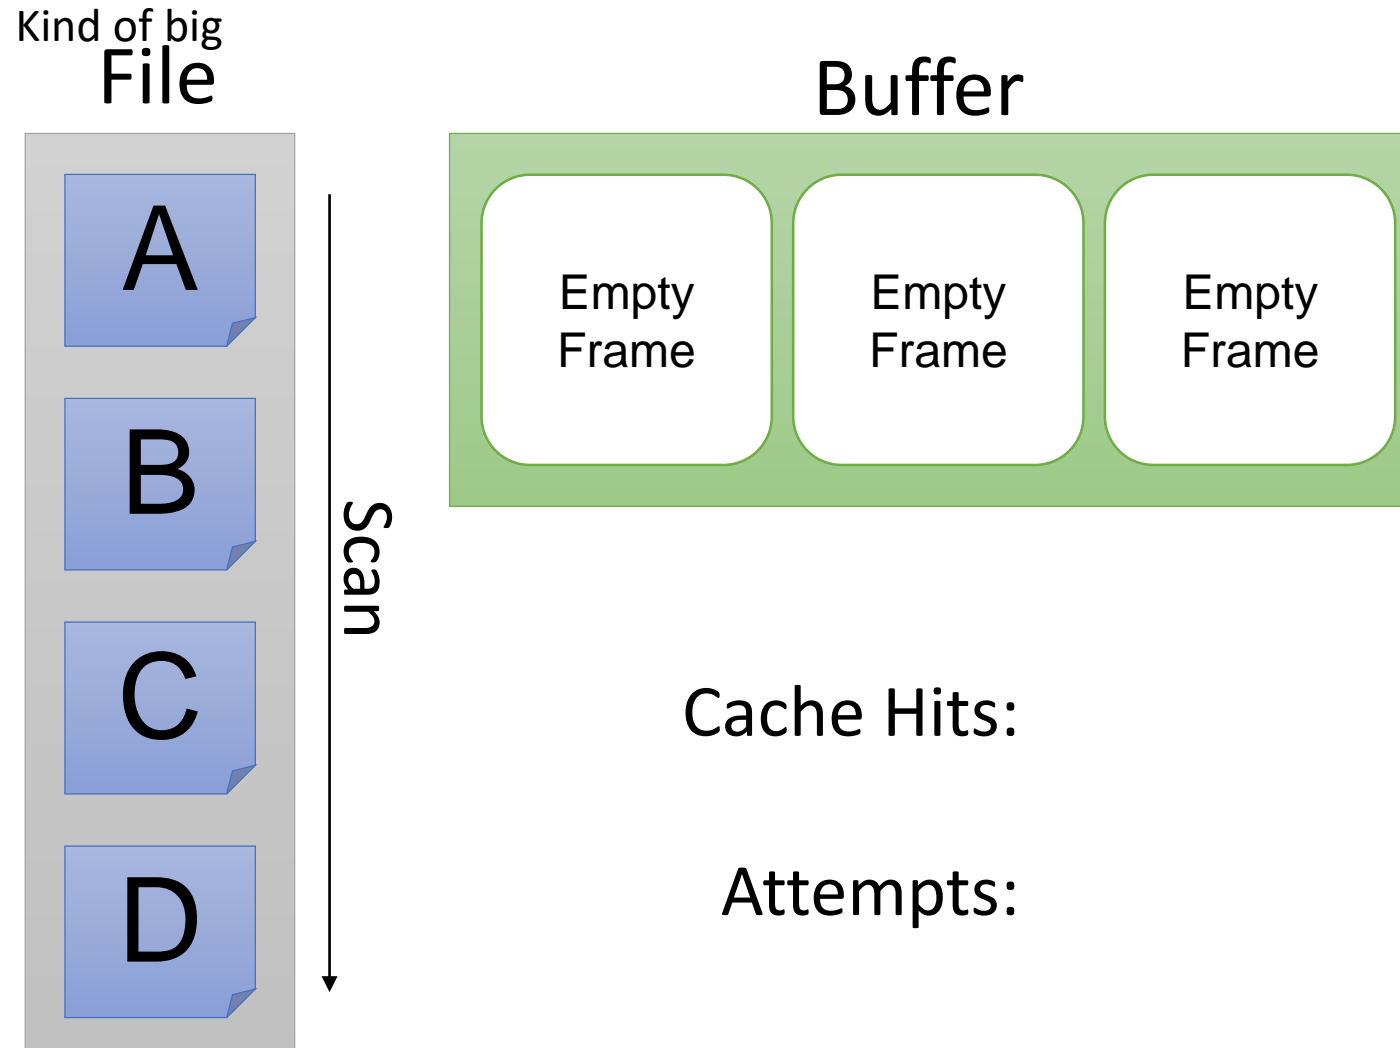
# LRU Replacement Policy

- Least Recently Used (LRU)
  - Pinned Frame: not available to replace
  - track time each frame last *unpinned* (end of use)
  - replace the frame which least recently unpinned
- Very common policy: intuitive and simple
  - Works well for repeated accesses to popular pages (temporal locality)
  - Can be costly. Why?
    - Need to maintain heap data-structure
    - Solution?

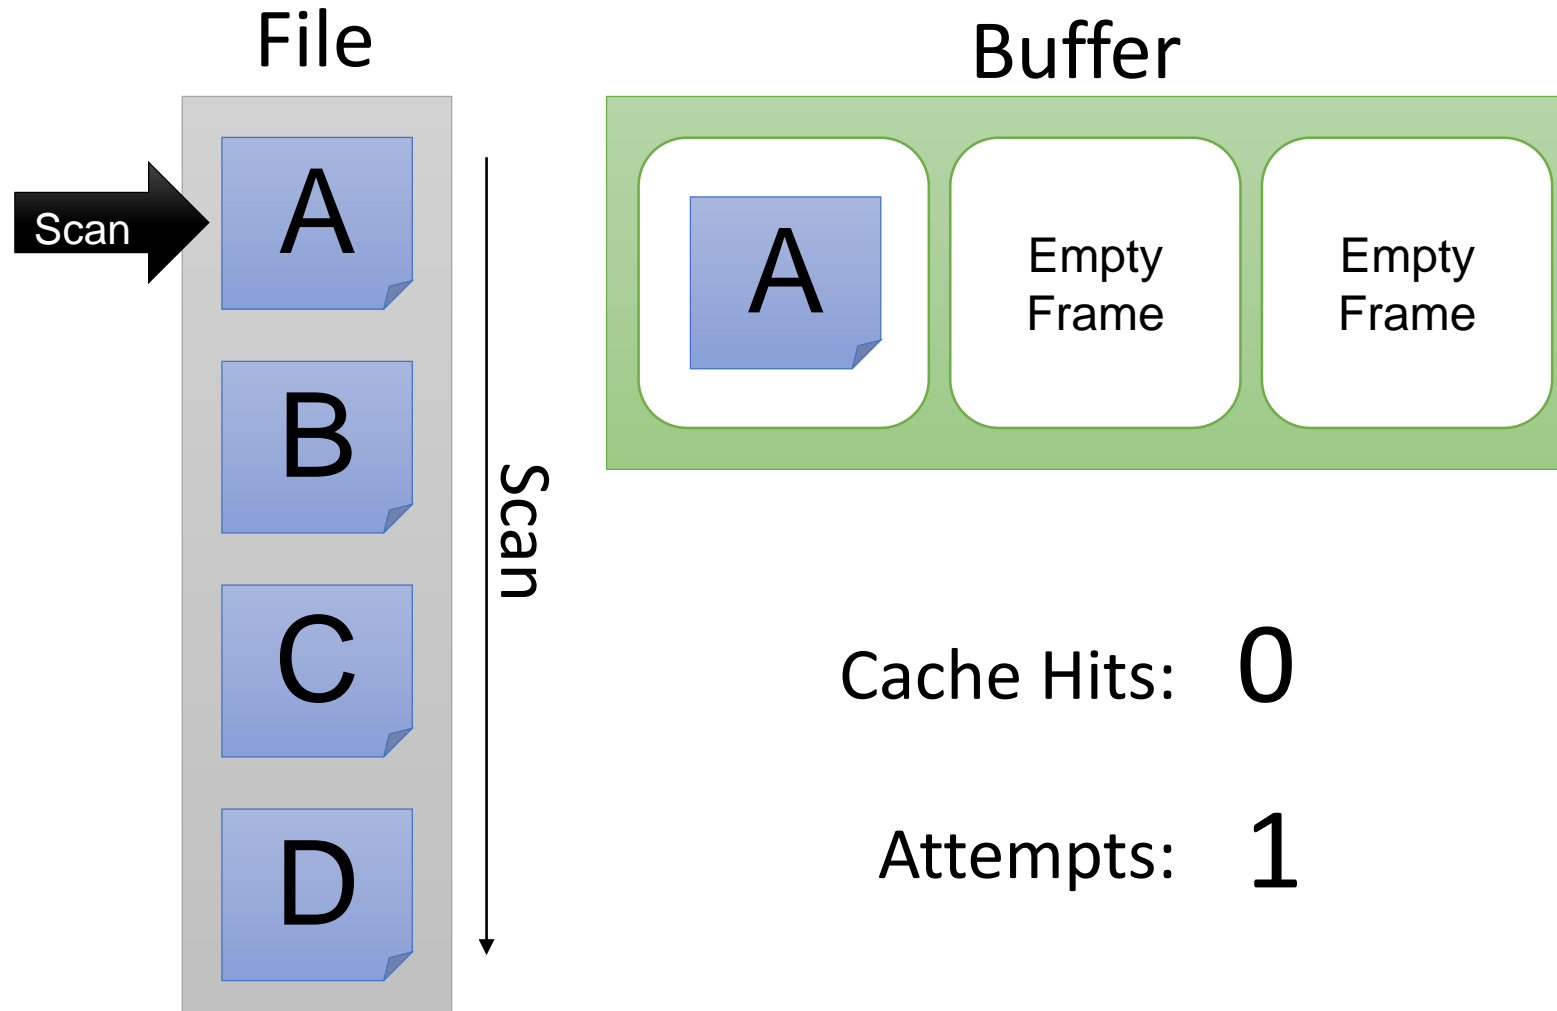
# Is LRU/Clock Always Best?

- Very common policy: *intuitive* and *simple*
- Works well for repeated accesses to popular pages
  - temporal locality
- LRU can be costly → Clock policy is cheap
- When might it perform poorly
  - What about repeated scans of big files?

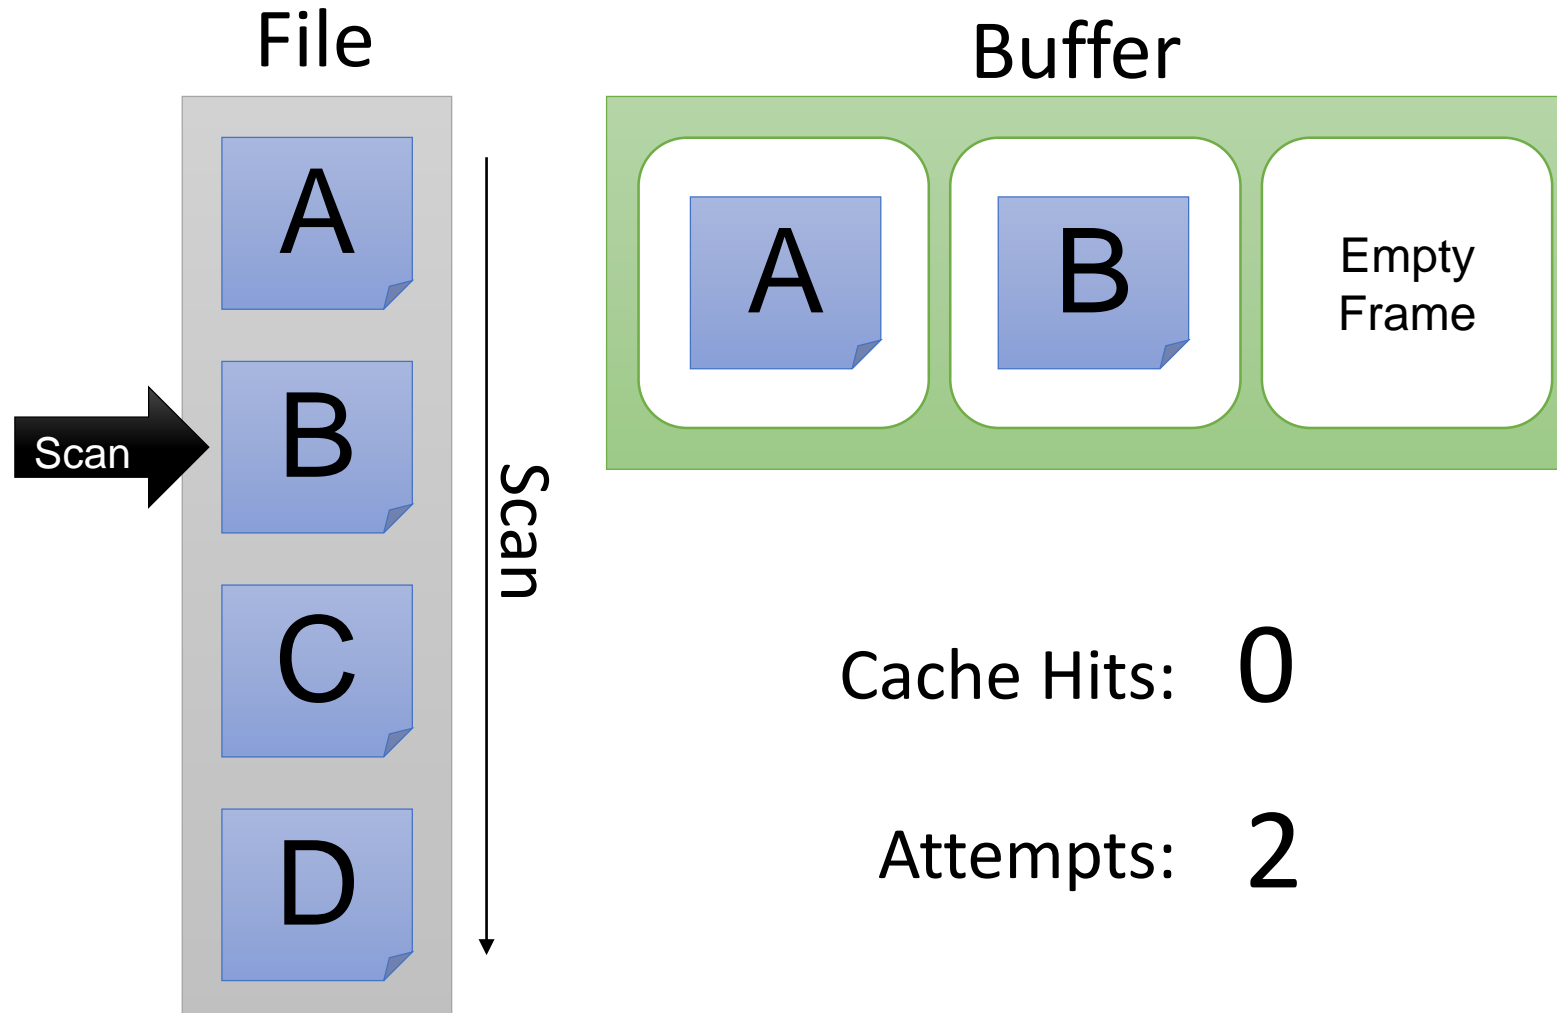
# Repeated Scan of Big File (LRU)



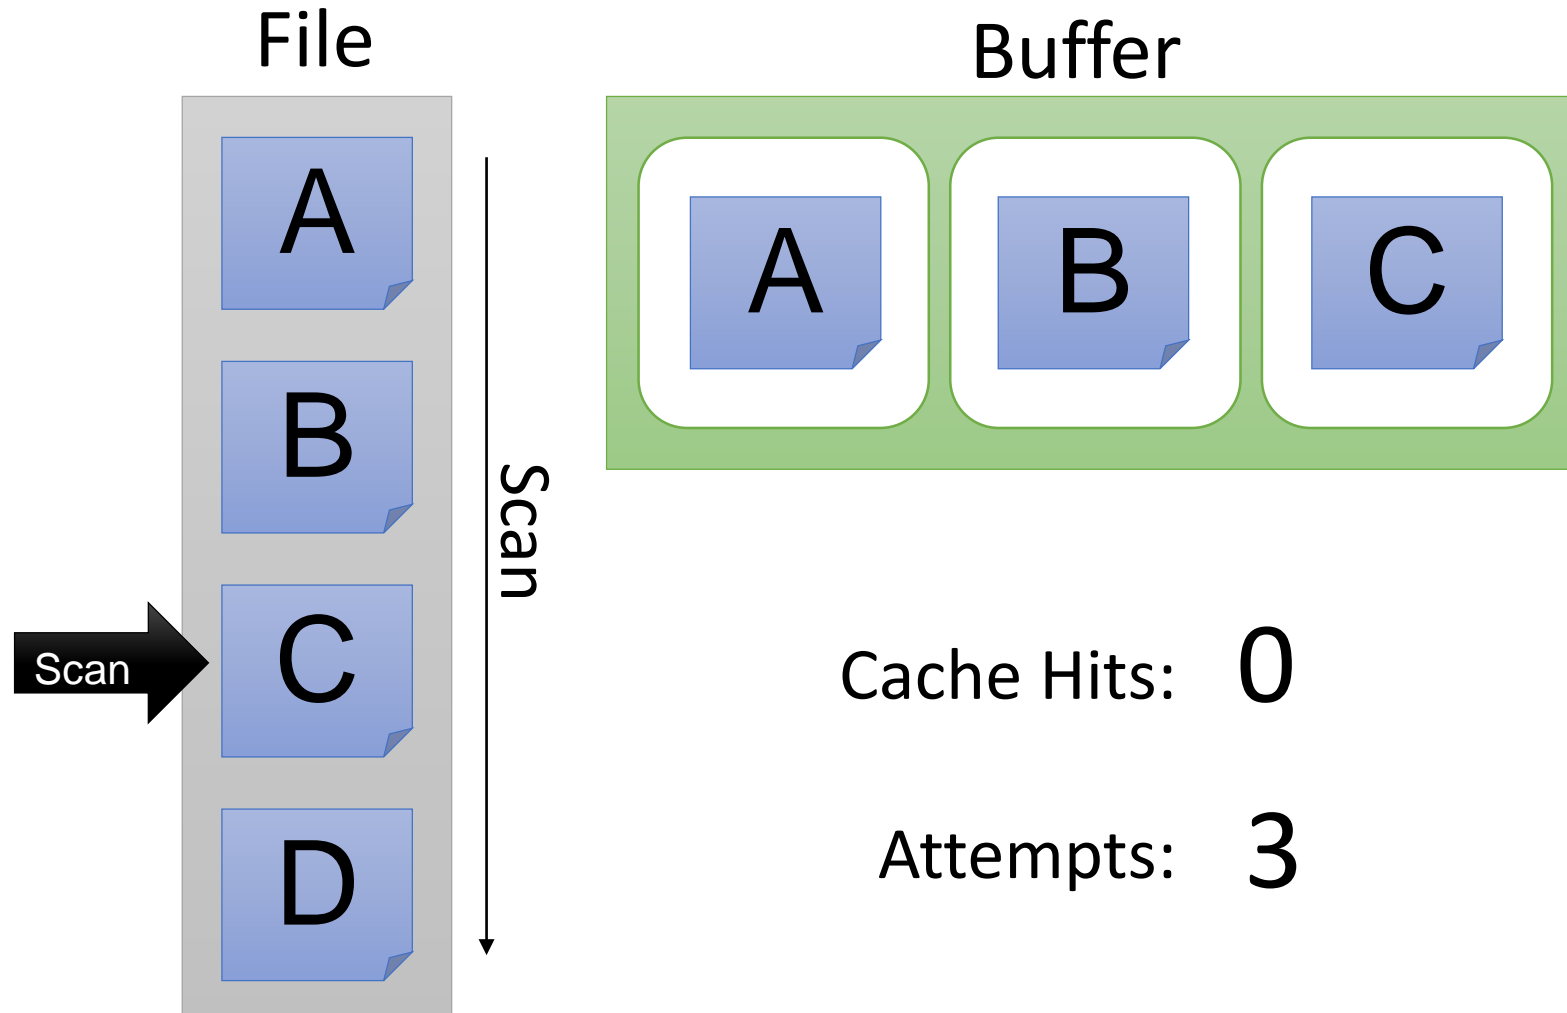
# Repeated Scan of Big File (LRU)



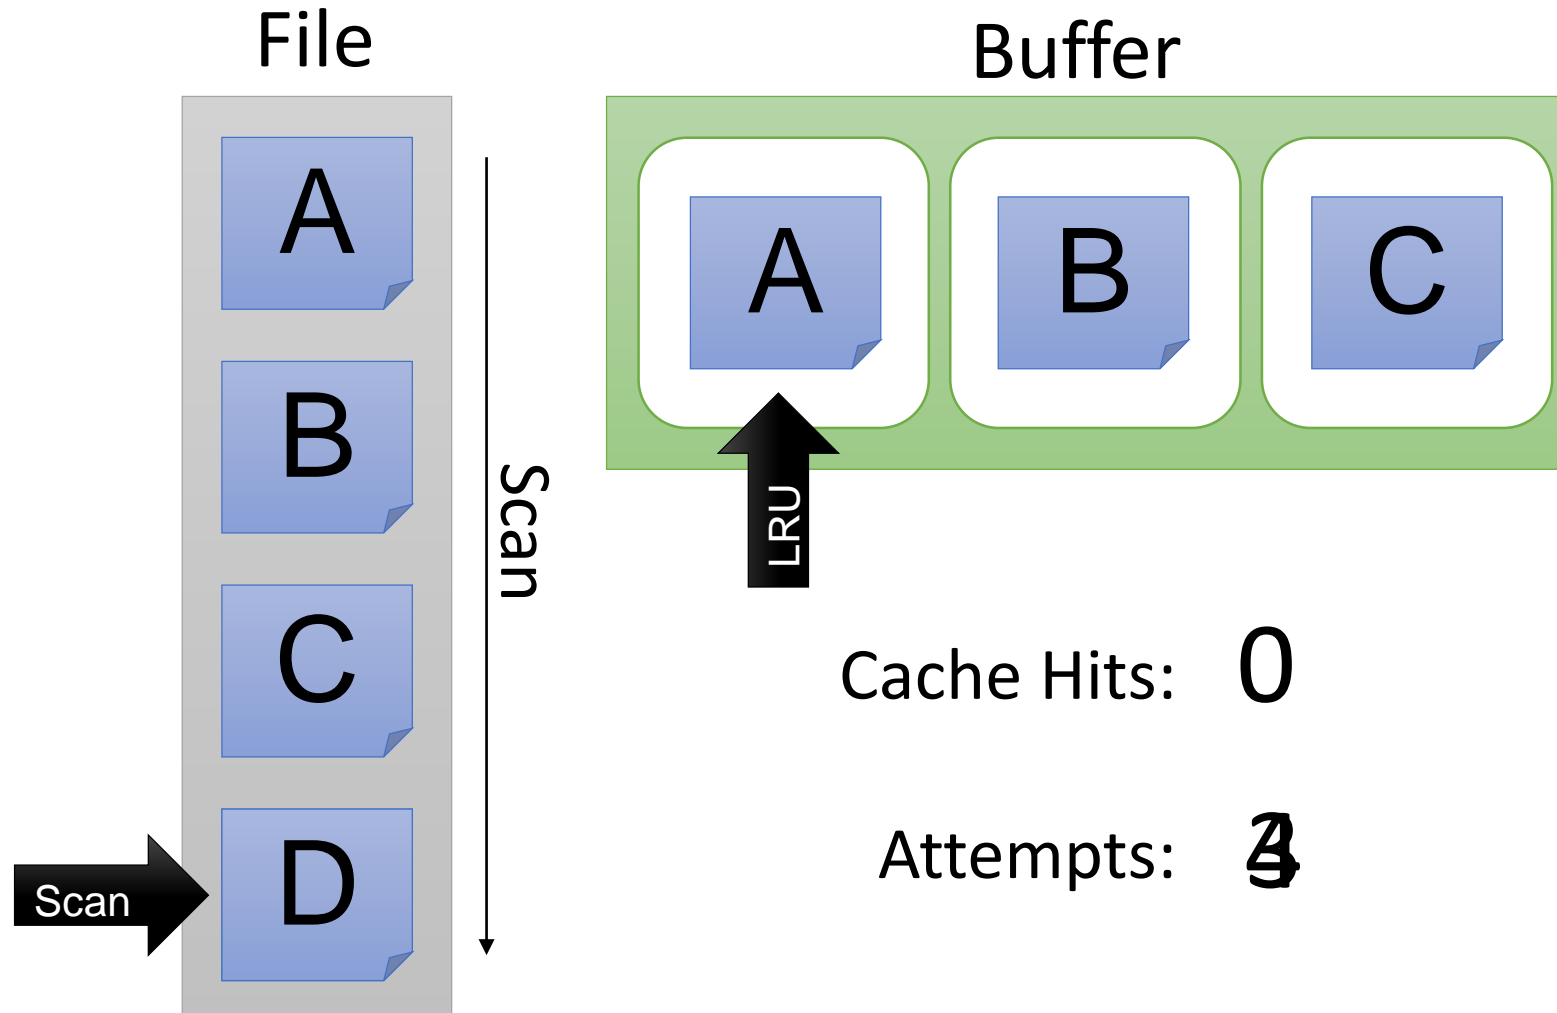
# Repeated Scan of Big File (LRU)



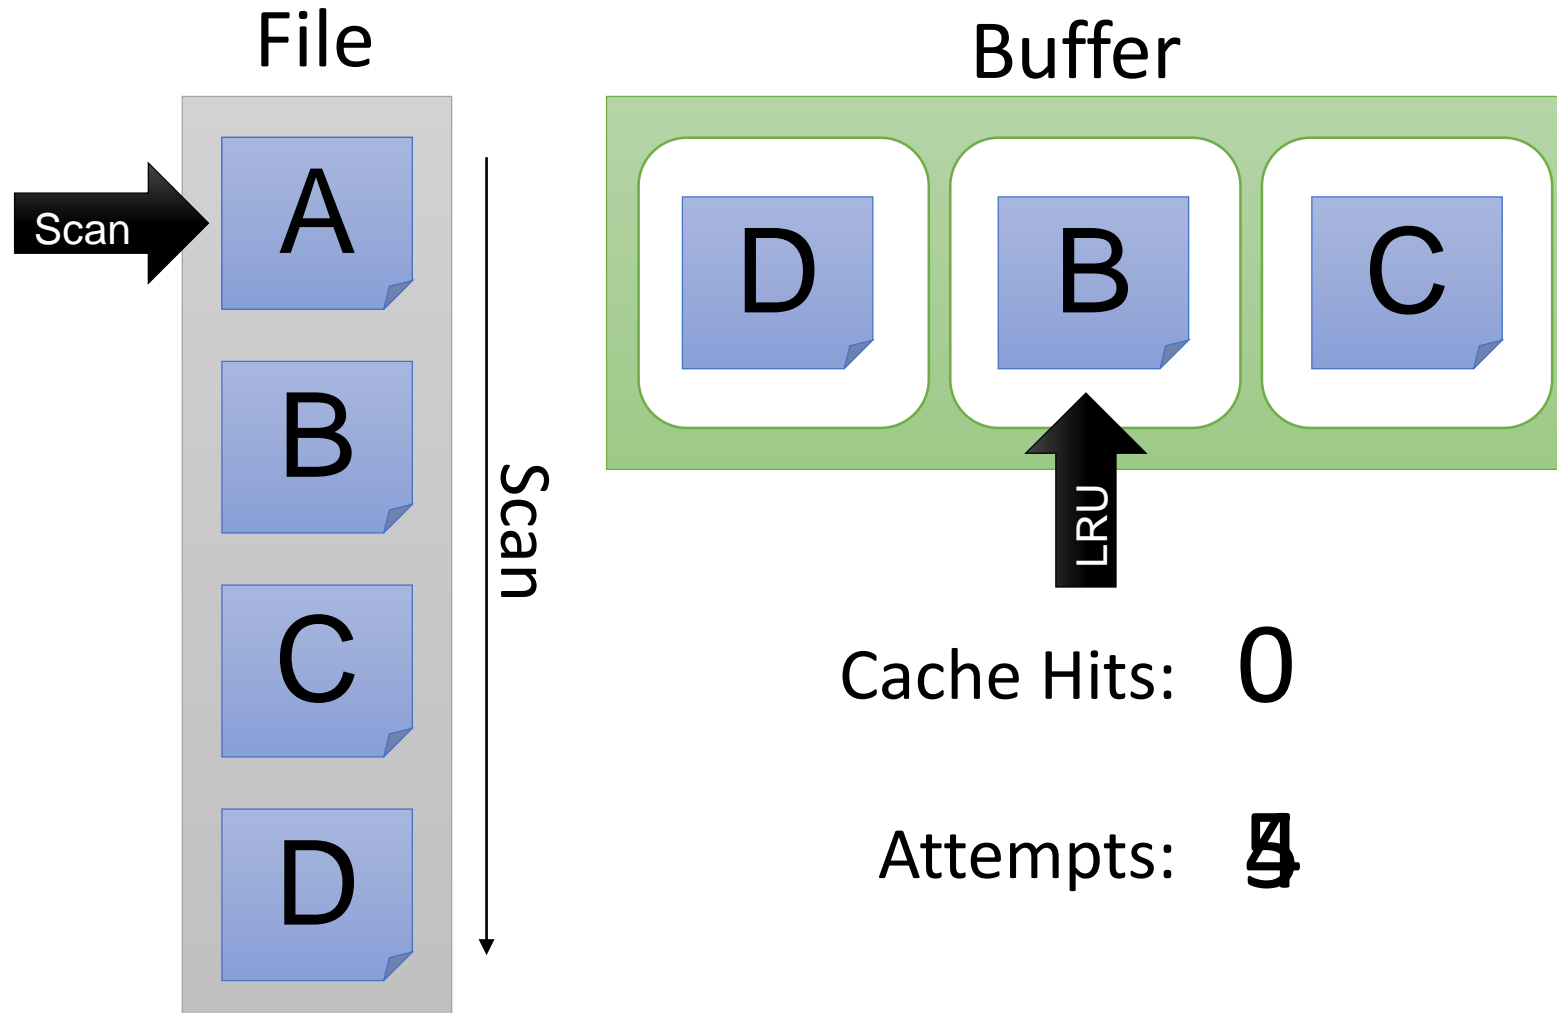
# Repeated Scan of Big File (LRU)



# Repeated Scan of Big File (LRU)

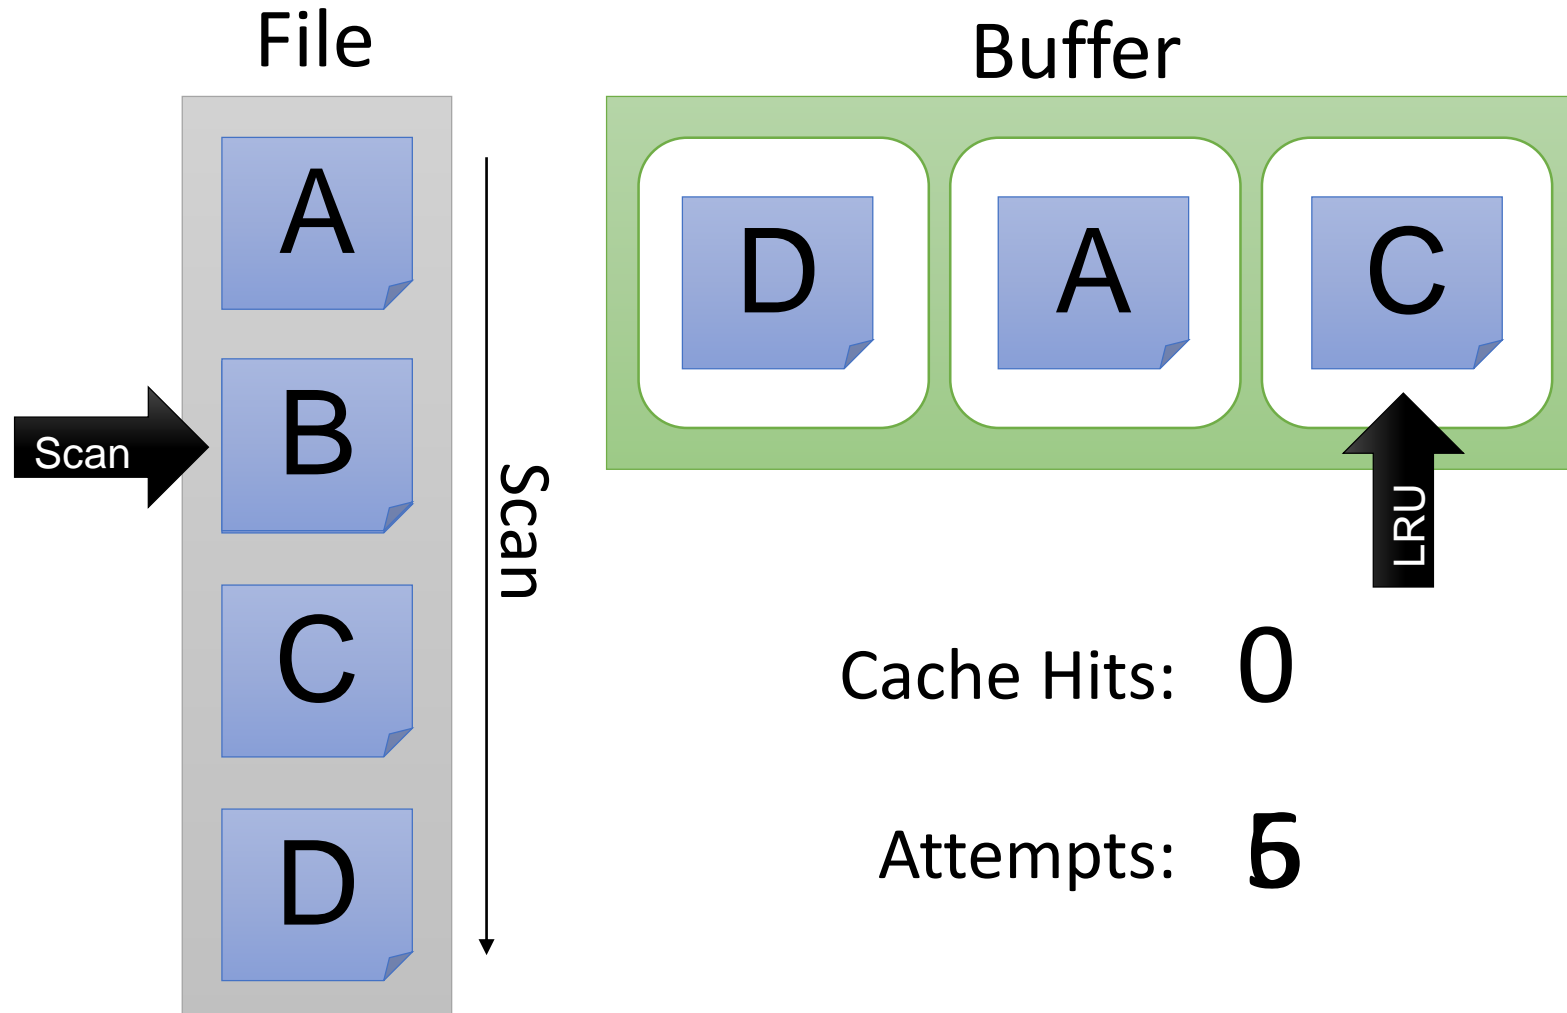


# Repeated Scan of Big File (LRU)

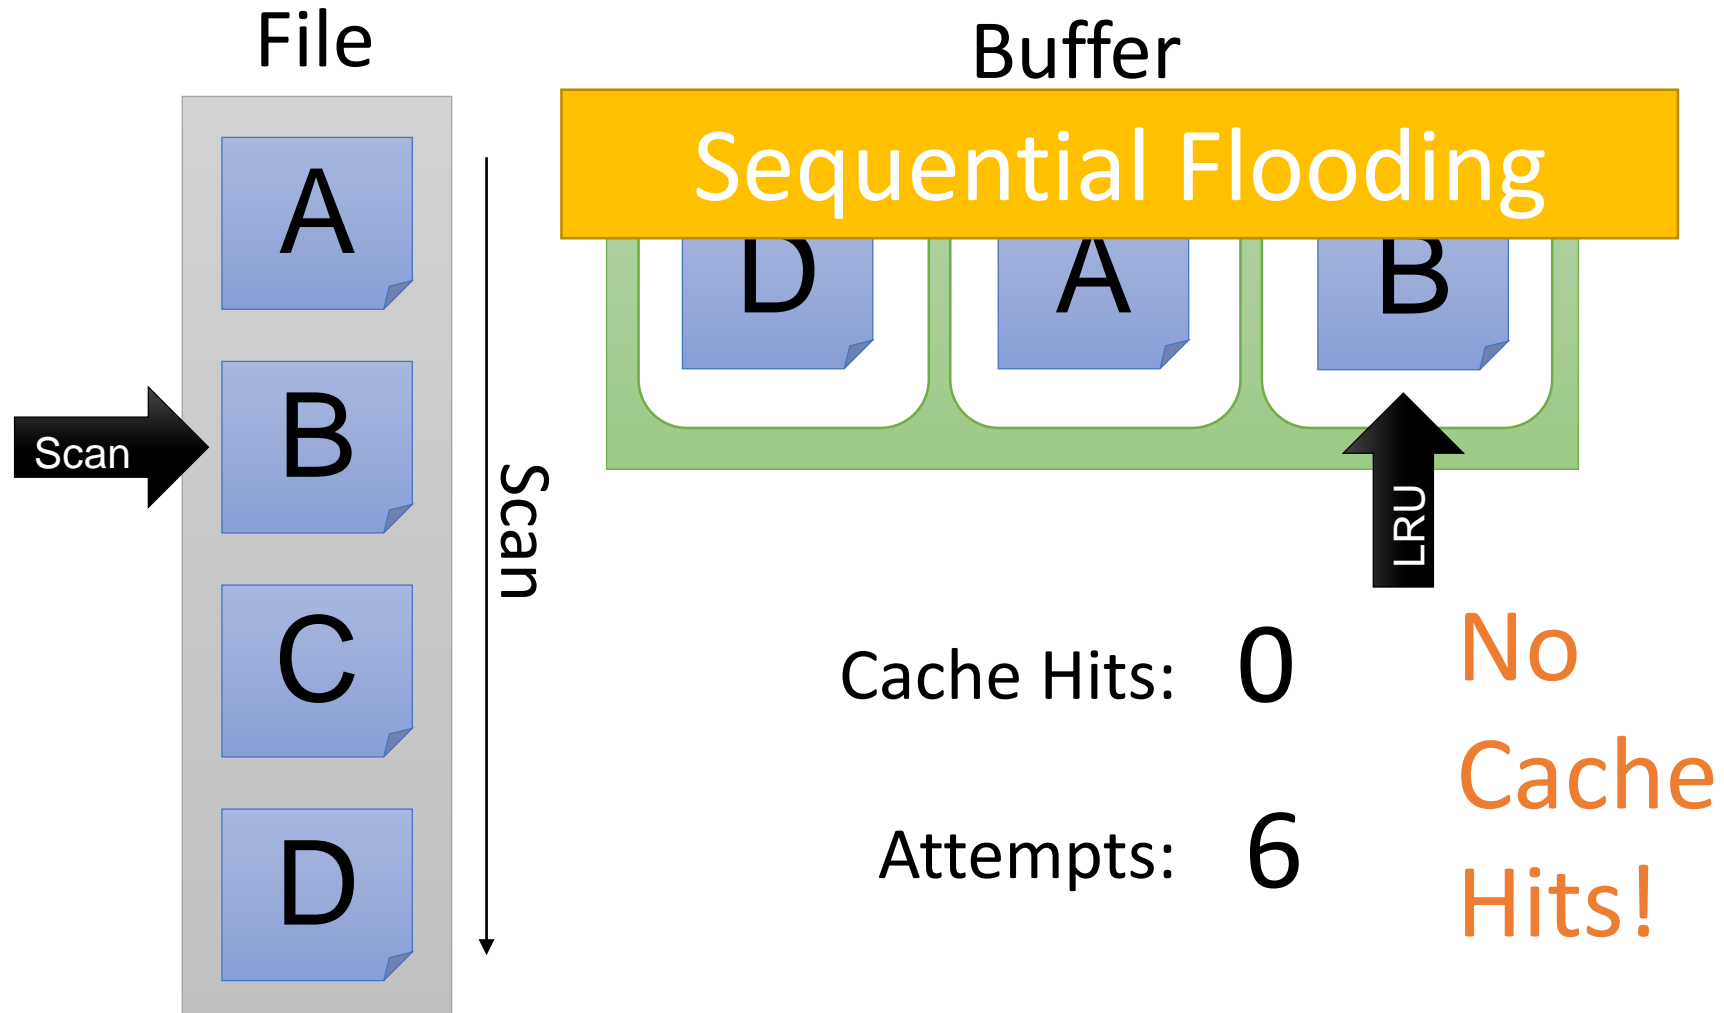




# Repeated Scan of Big File (LRU)




# Repeated Scan of Big File (LRU)

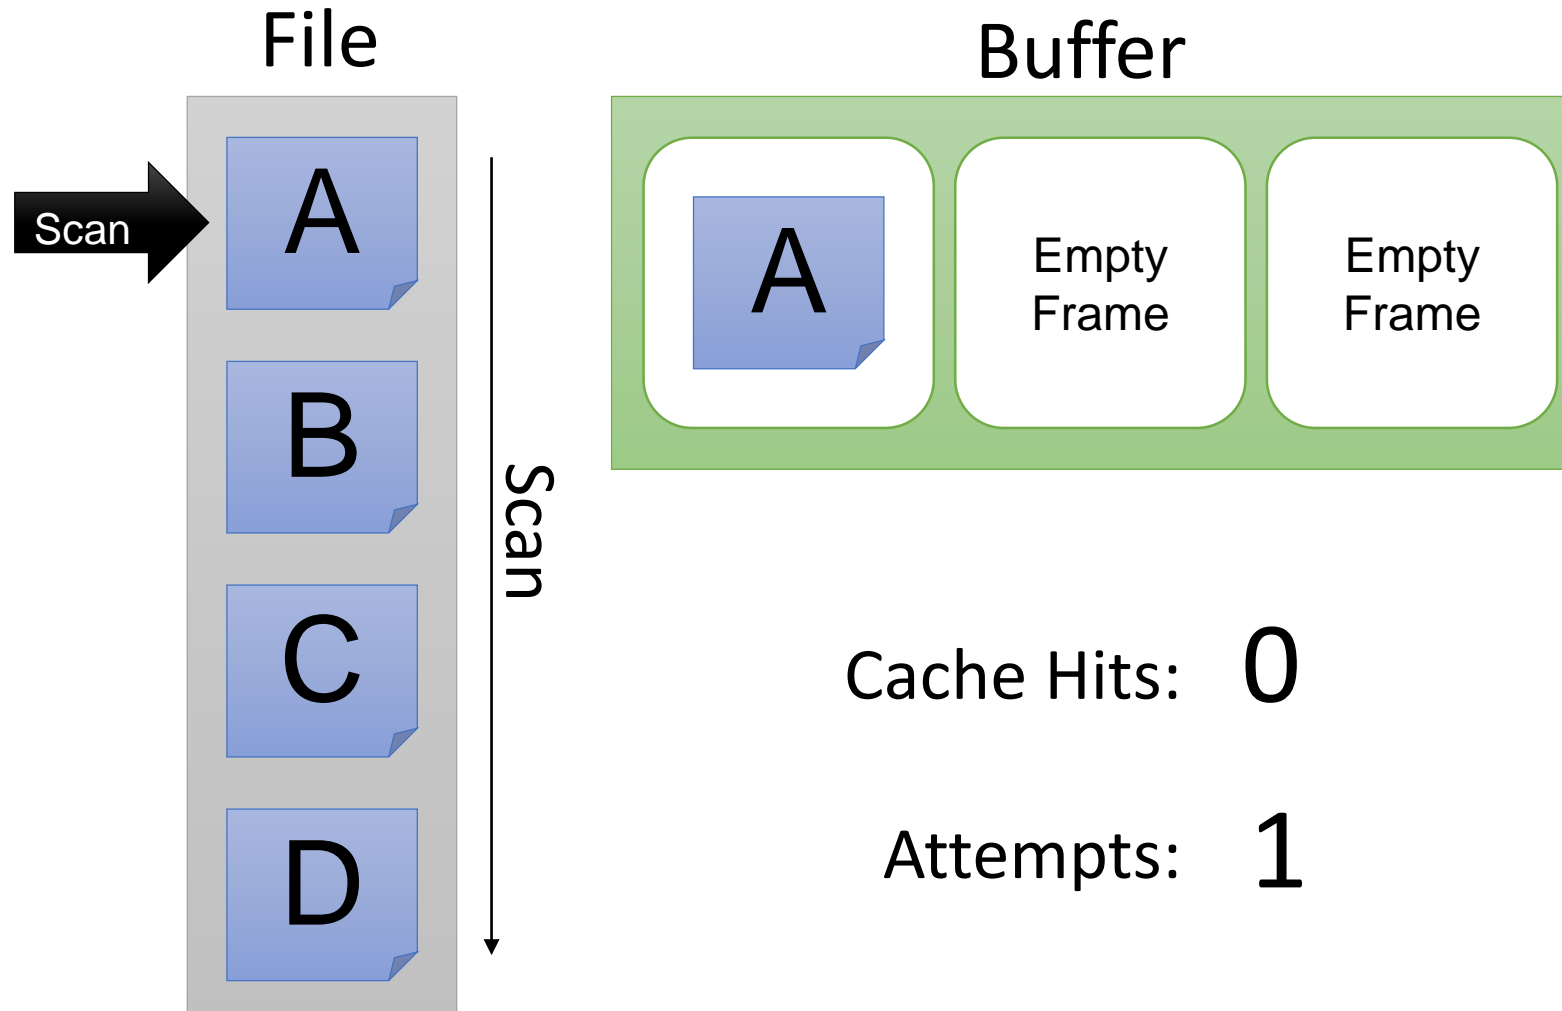


# Repeated Scan of Big File (MRU)


Most Recently Used




# Repeated Scan of Big File (MRU)



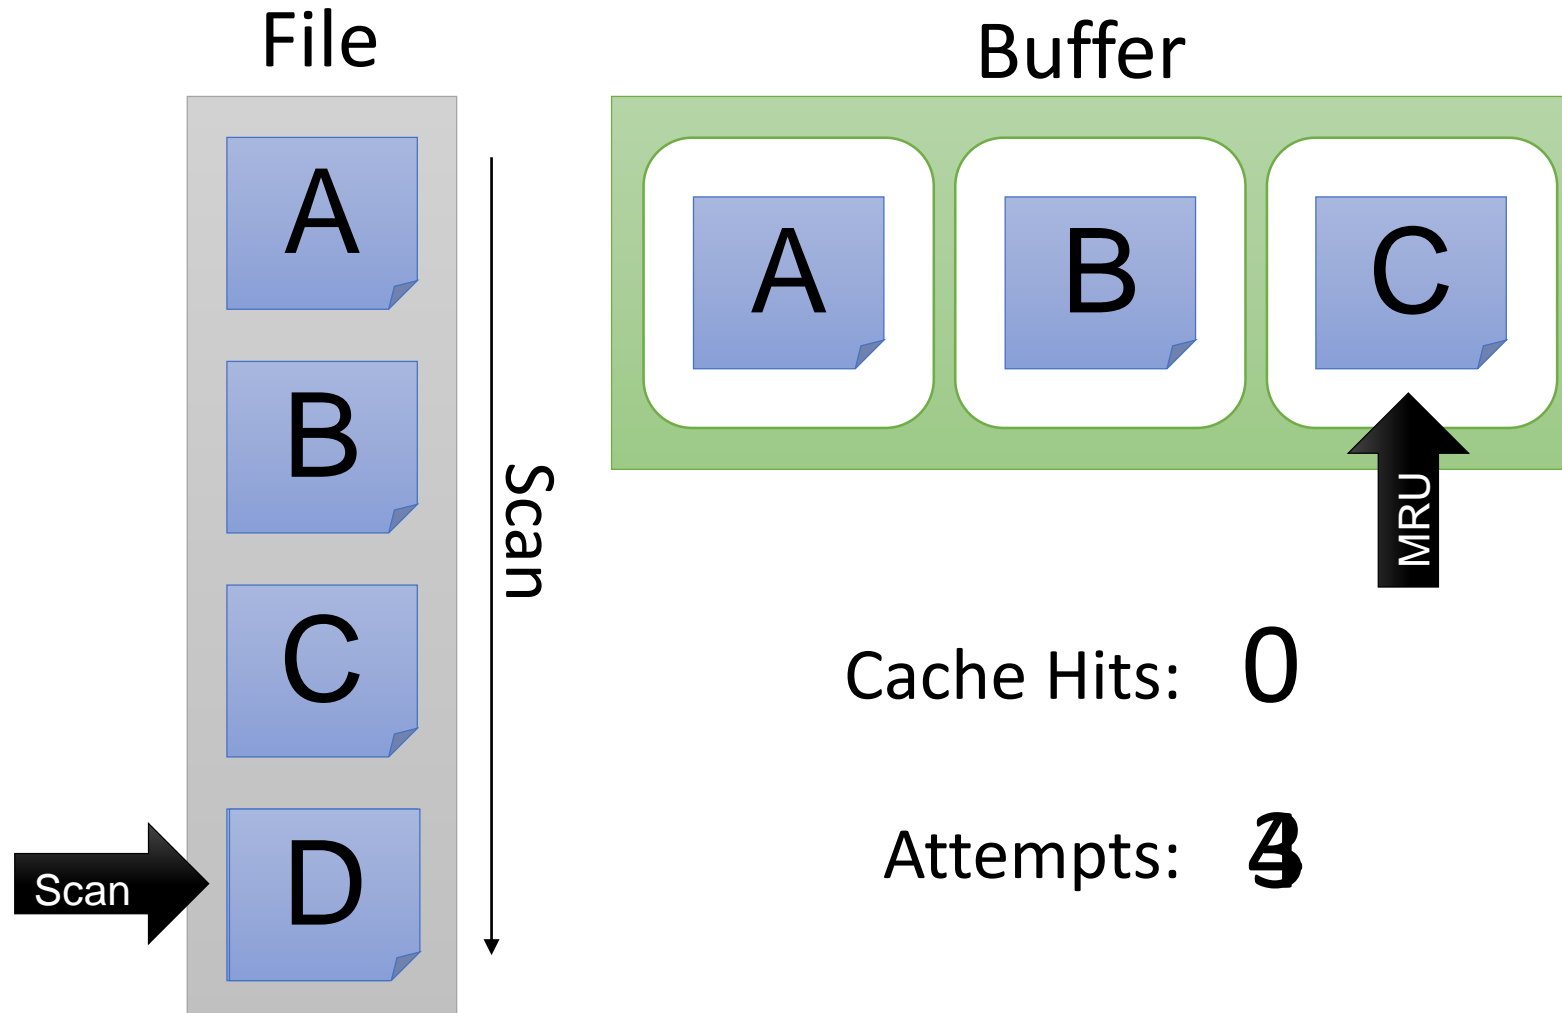
# Repeated Scan of Big File (MRU)




# Repeated Scan of Big File (MRU)



# Repeated Scan of Big File (MRU)

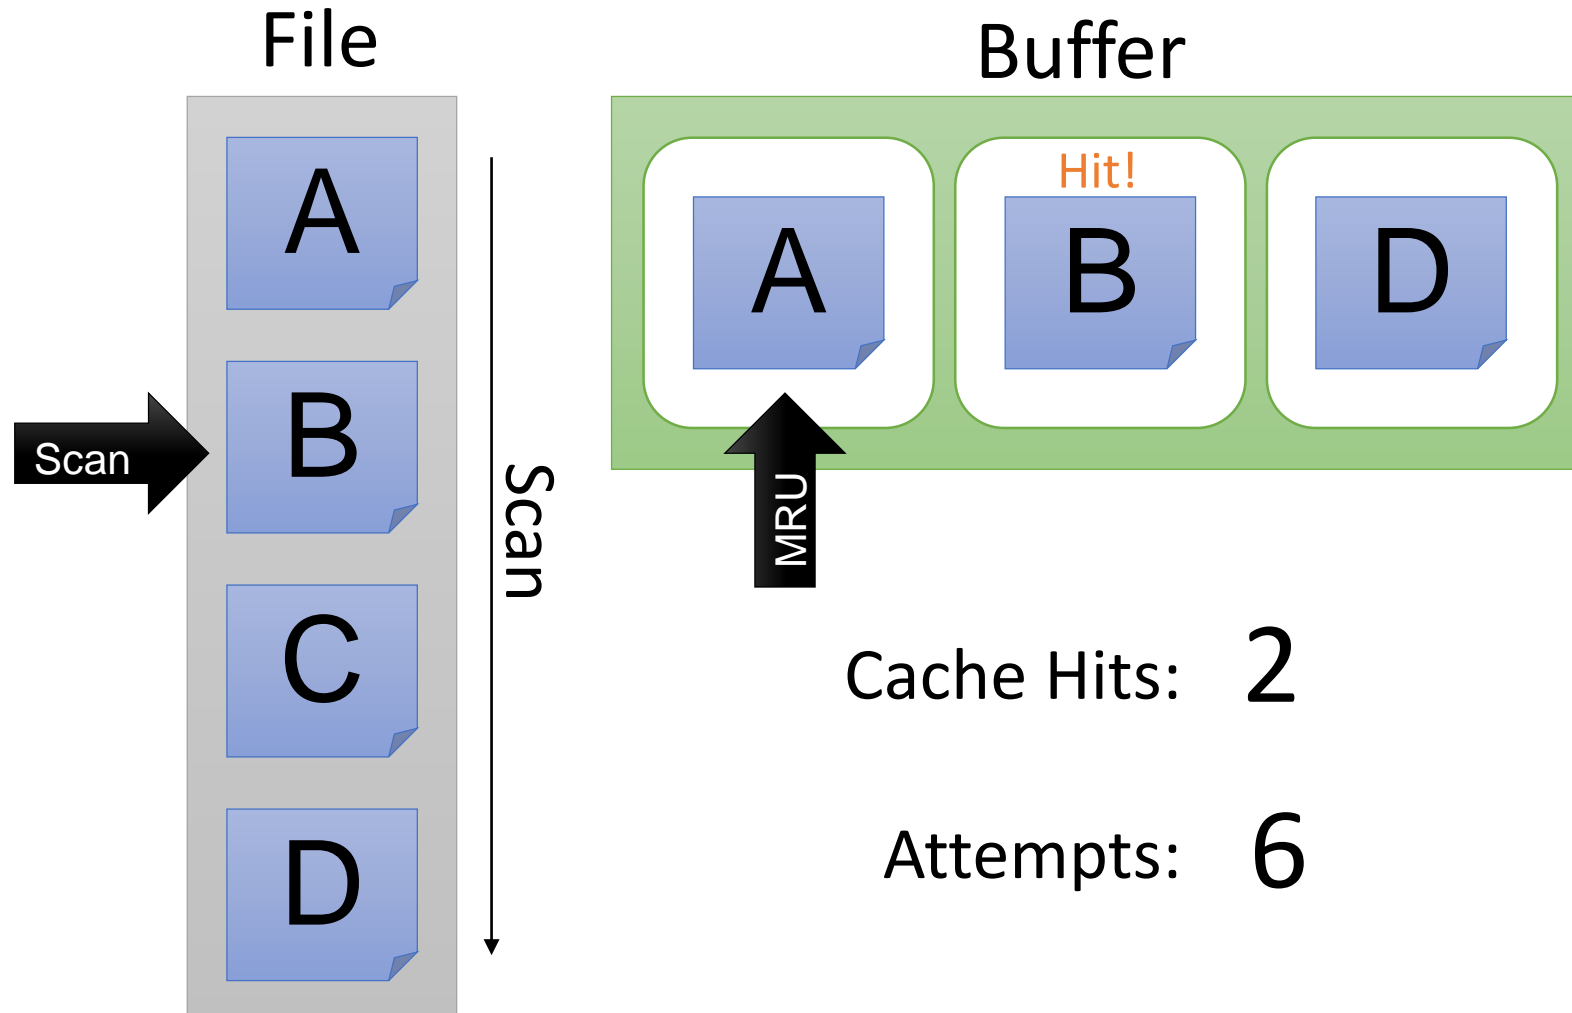


# Repeated Scan of Big File (MRU)

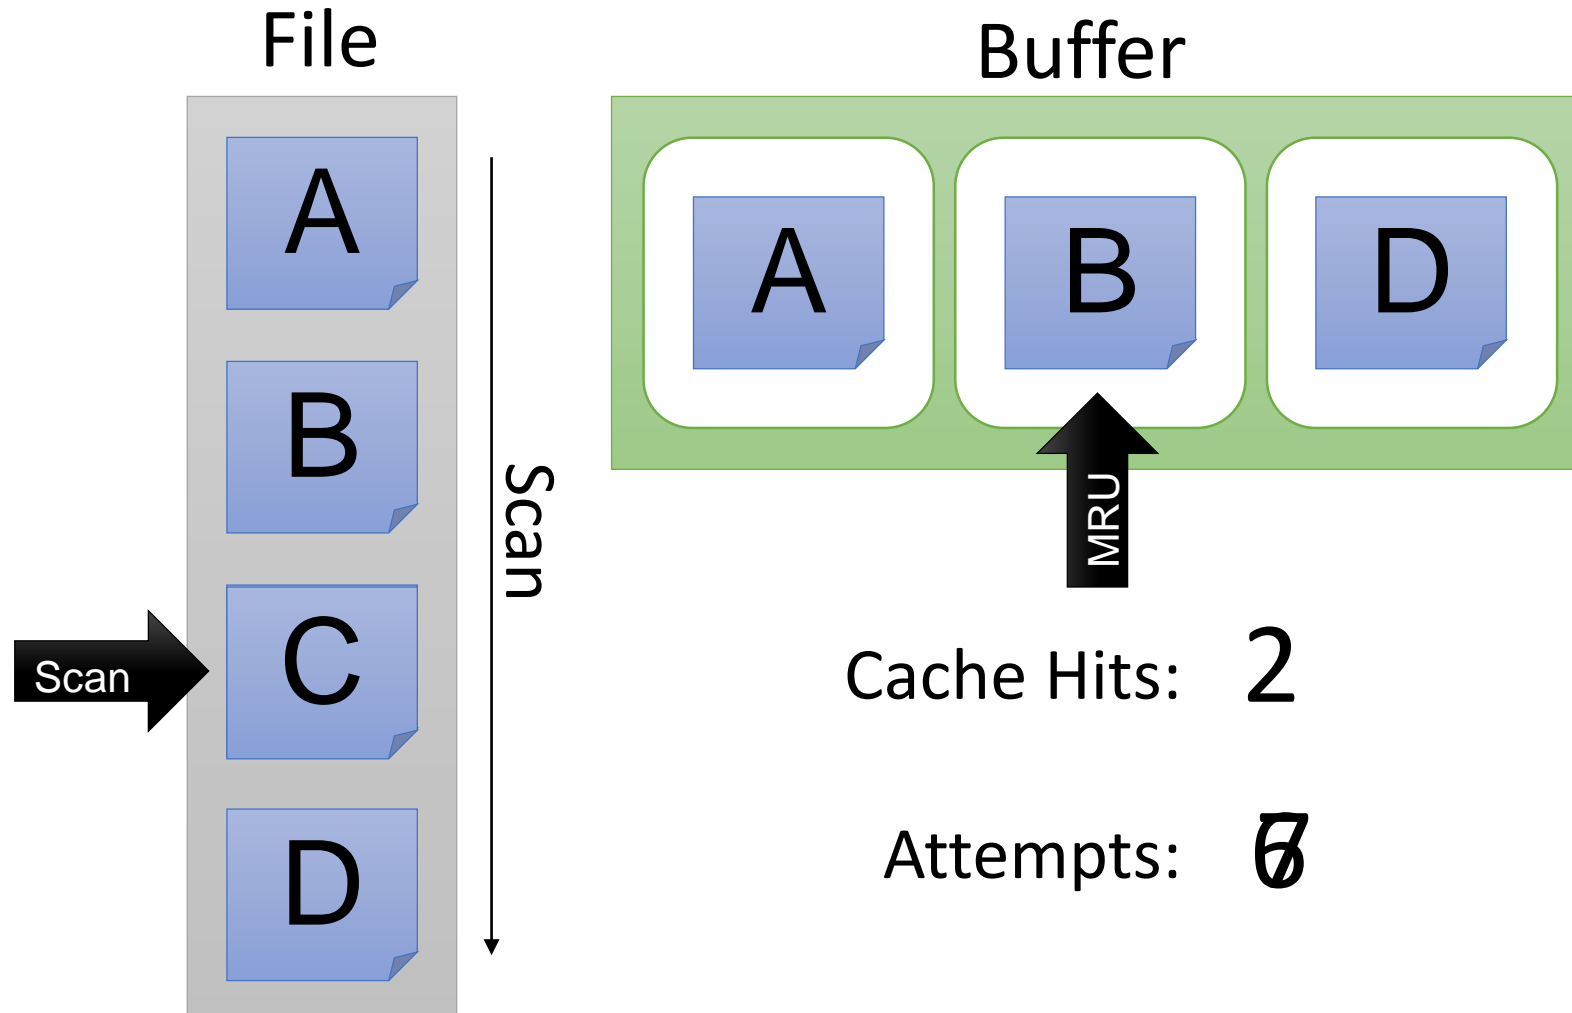




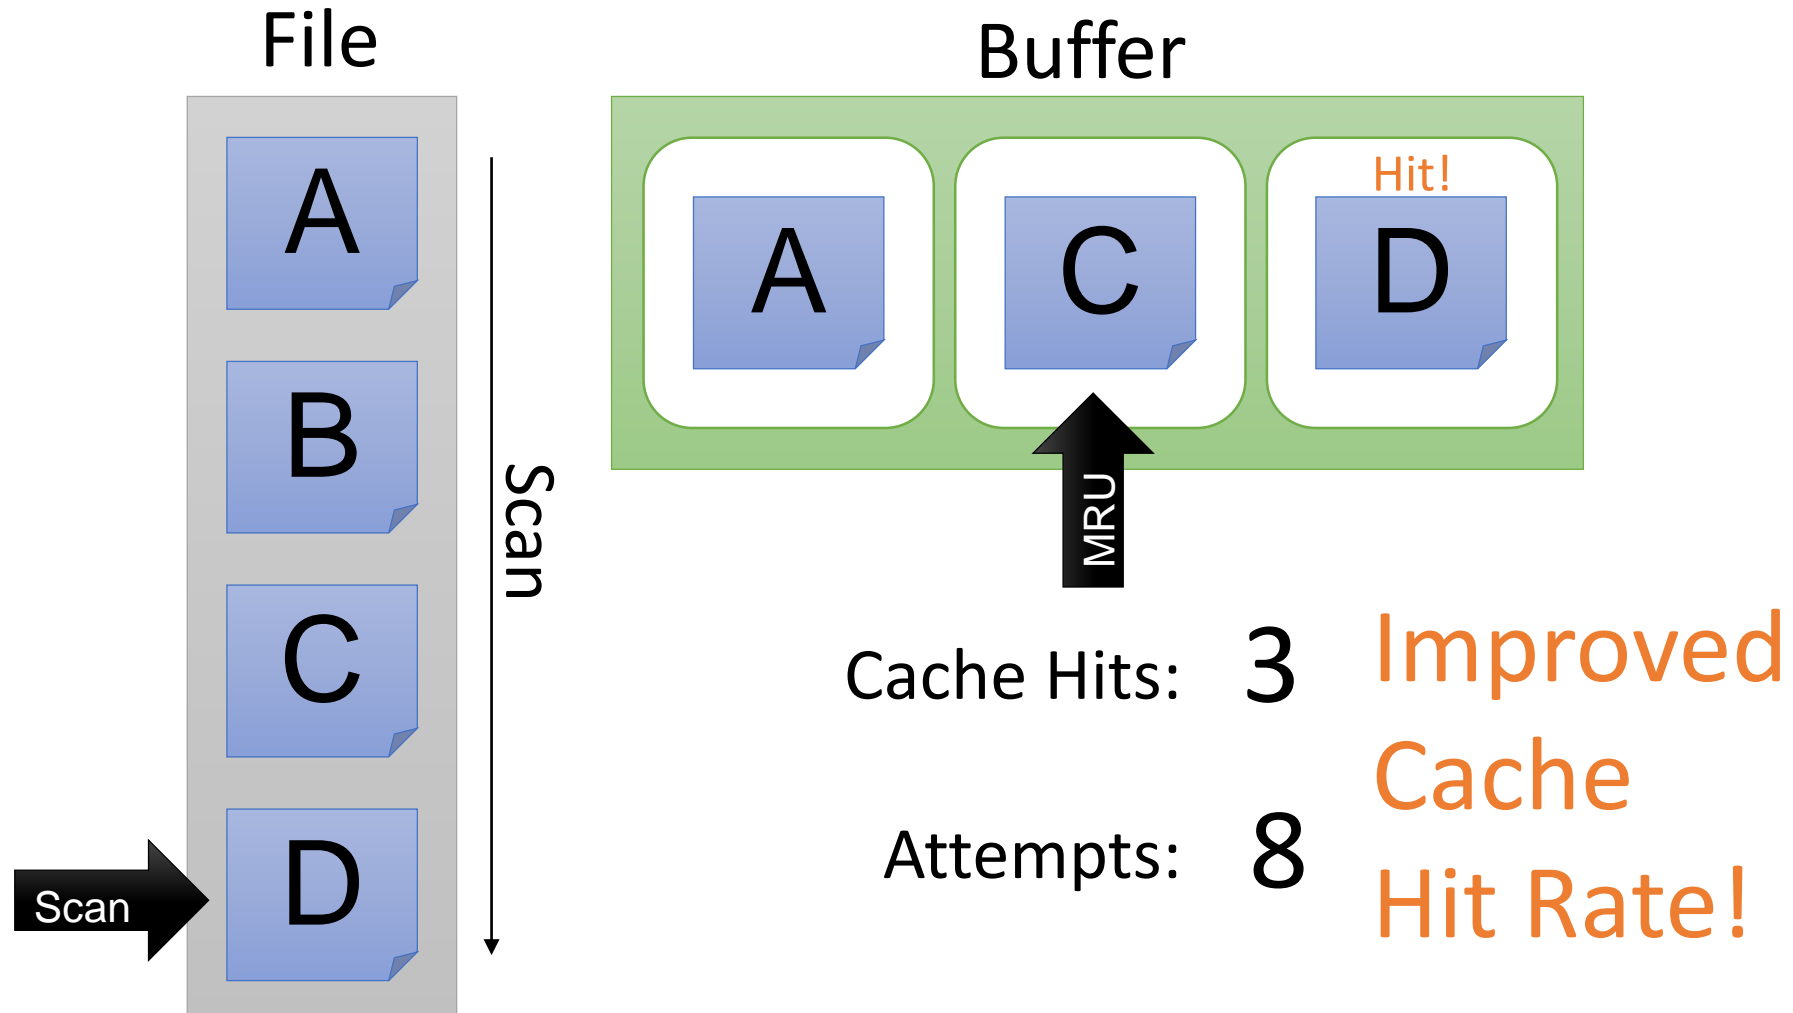
# Repeated Scan of Big File (MRU)




# Repeated Scan of Big File (MRU)



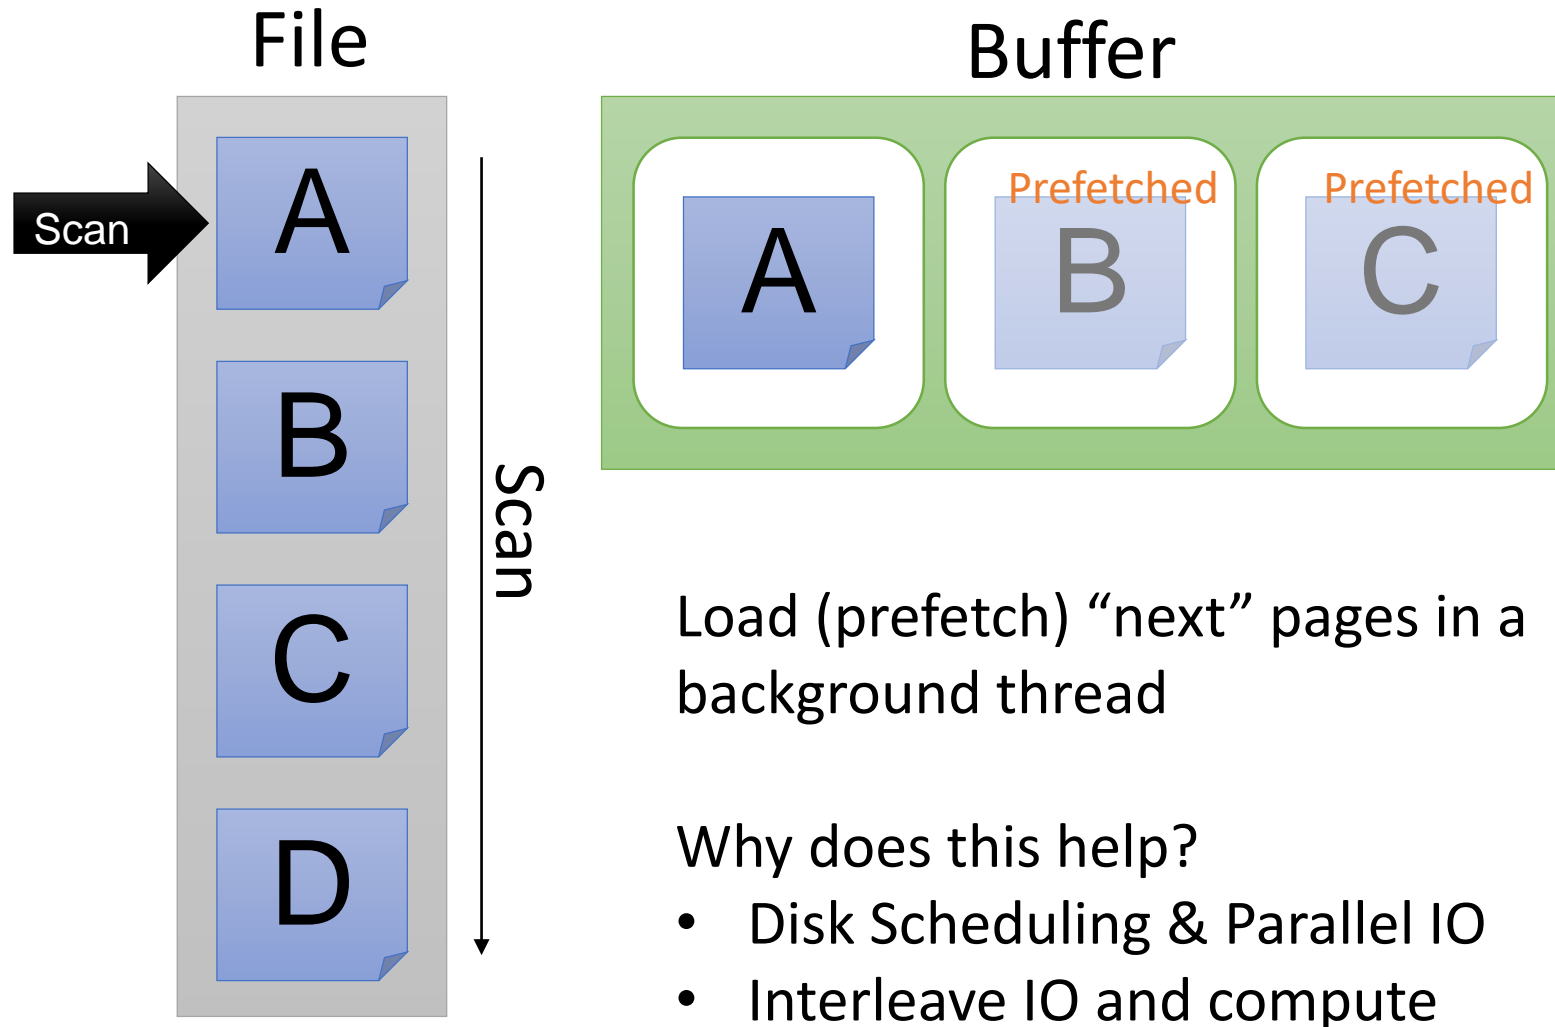
# Repeated Scan of Big File (MRU)



# Repeated Scan of Big File (MRU)



# Background Prefetching



# The Buffer Manager

- A **buffer manager** handles supporting operations for the buffer:
  - Primarily, handles & executes the “replacement policy”
    - i.e. finds a page in buffer to flush/release if buffer is full and a new page needs to be read in
- DBMSs typically implement their own buffer management routines