

CS150: Database & Datamining

Lecture 28: NoSQL I

Xuming He
Spring 2019

Acknowledgement: Slides are adopted from the Berkeley course CS186 by Joey Gonzalez and Joe Hellerstein, Stanford CS145 by Peter Bailis, IIT Course 236363.

SQL Review

- SQL stands for the query language
- But commonly refers to the traditional RDBMS:
 - Relational storage of data
 - Each tuple is stored consecutively
 - Joins as first-class citizens
 - In fact, normal forms prefer joins to maintenance
 - Strong guarantees on transaction management
 - No consistency worries when many transactions operate simultaneously on common data
- Focus on *scaling up*
 - That is, make a single machine do more, faster

Trends Drive Common Requirements

Social media + mobile
computing



- Explosion in data, always available, constantly read and updated
- High load of simple requests of a common nature
- Some consistency can be compromised (e.g., 🗑)


Cloud computing + open
source



- Affordable resources for management / analysis of data
- People of various skills / budgets need software solutions for distributed analysis of massive data

Database solutions need to *scale out*
(utilize distribution, “scale horizontally”)

Compromises Required



What is needed for effective distributed, data- and user-intensive applications?

1. Use data models and storage that allow to avoid joins of big objects
2. Relax the guarantees on consistency

Transaction in distributed systems

- Consider a parallel or distributed database
 - One that handles updates
 - Unlike, say, Spark
- Data is *partitioned (sharded)* across nodes
 - This is how we scale up data volume!
 - Assume only one copy of each record (for now)
- If a transaction touches one machine
 - Life is good
- If a transaction touches multiple machines
 - ACID becomes extremely expensive!

Consistency

- Assume each transaction is assigned to a node that serves as its “coordinator”
- When a multi-node txn finishes, the DBMS needs to ask all of the nodes involved whether it is safe to commit
 - All nodes must agree on the outcome
- Need **two-phase commit**

Two-Phase Commit: Motivation

Coordinator



Subordinate 1



Each subordinate
holds fraction of
database

Subordinate 2

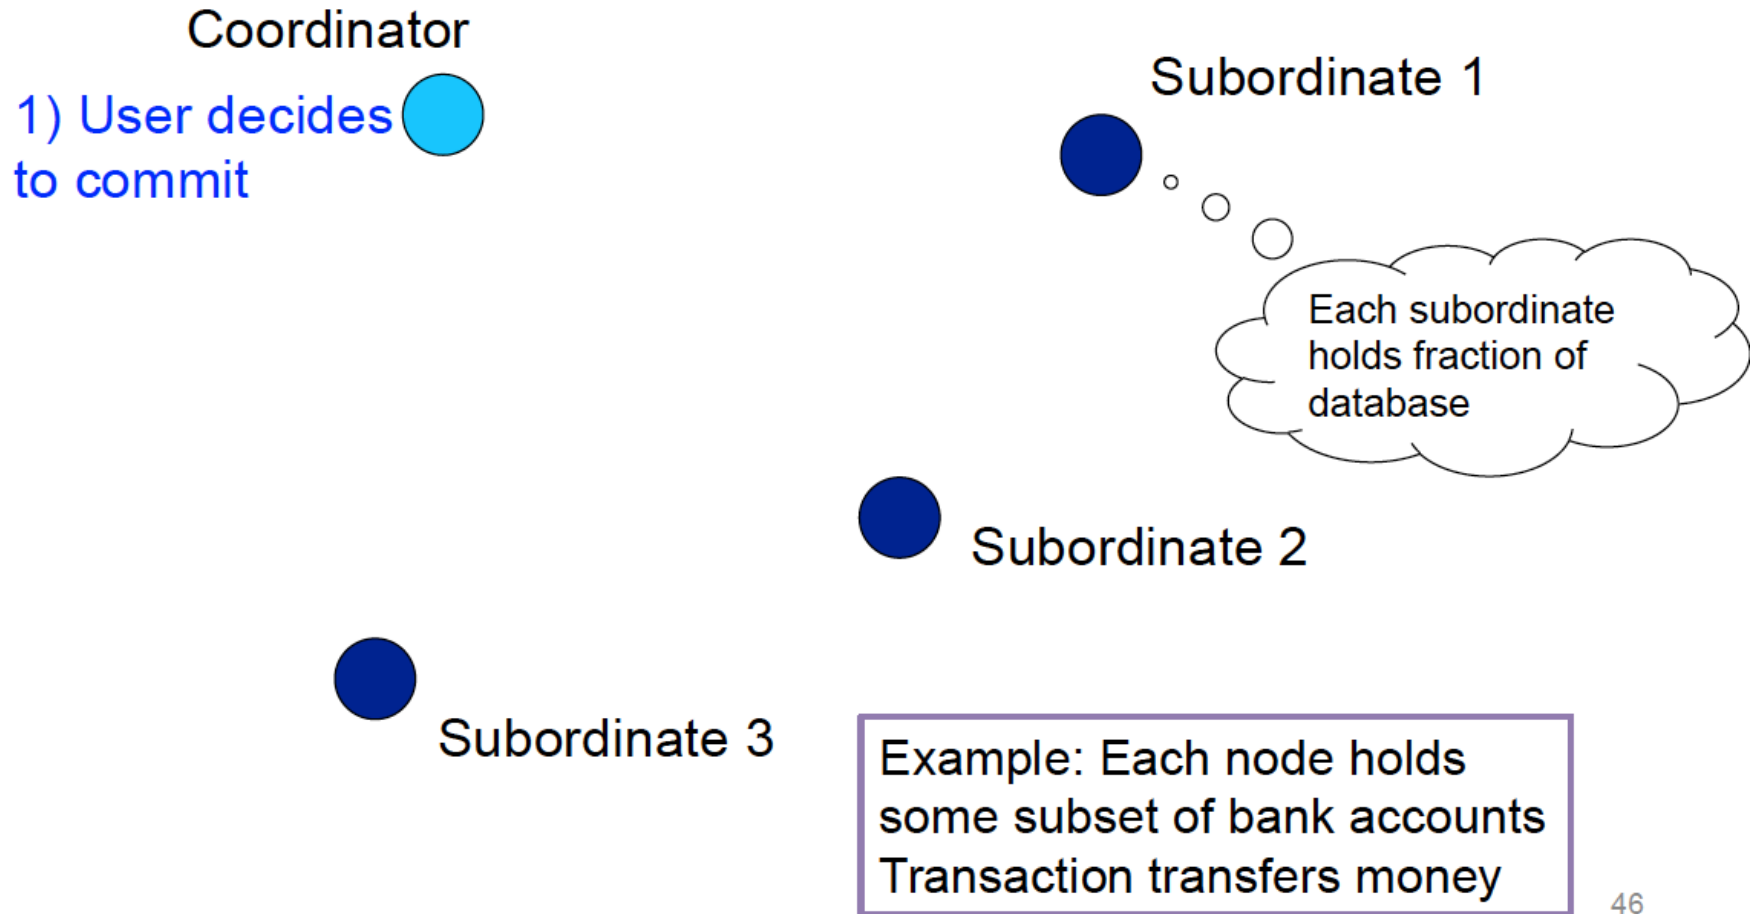


Subordinate 3

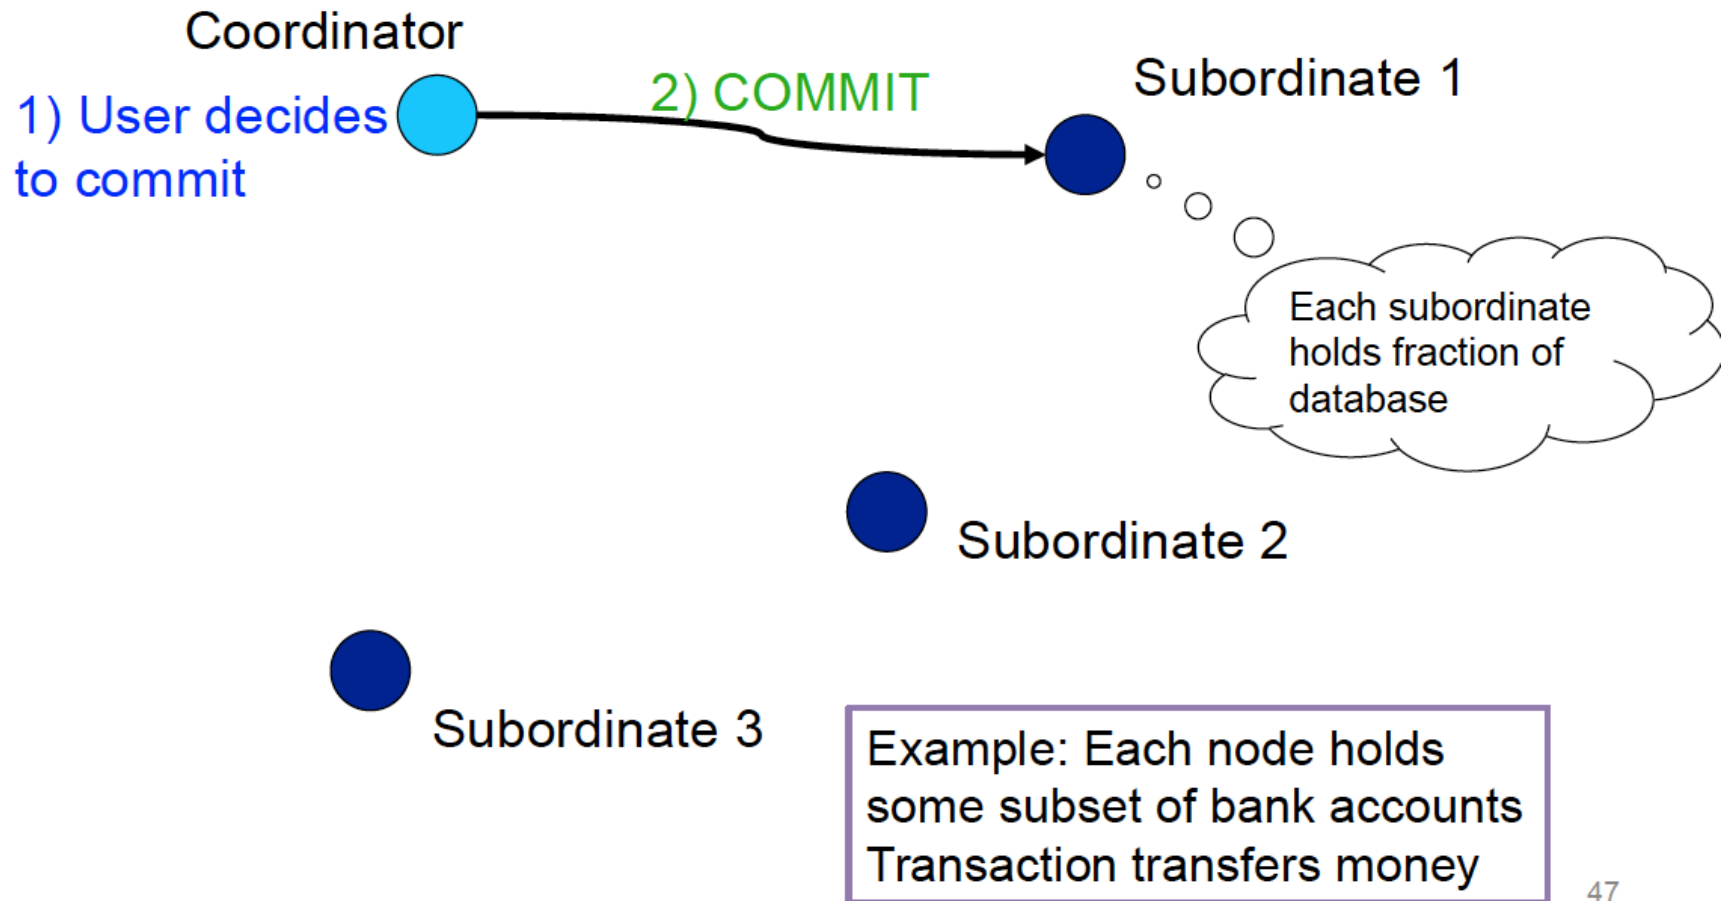


Example: Each node holds
some subset of bank accounts
Transaction transfers money

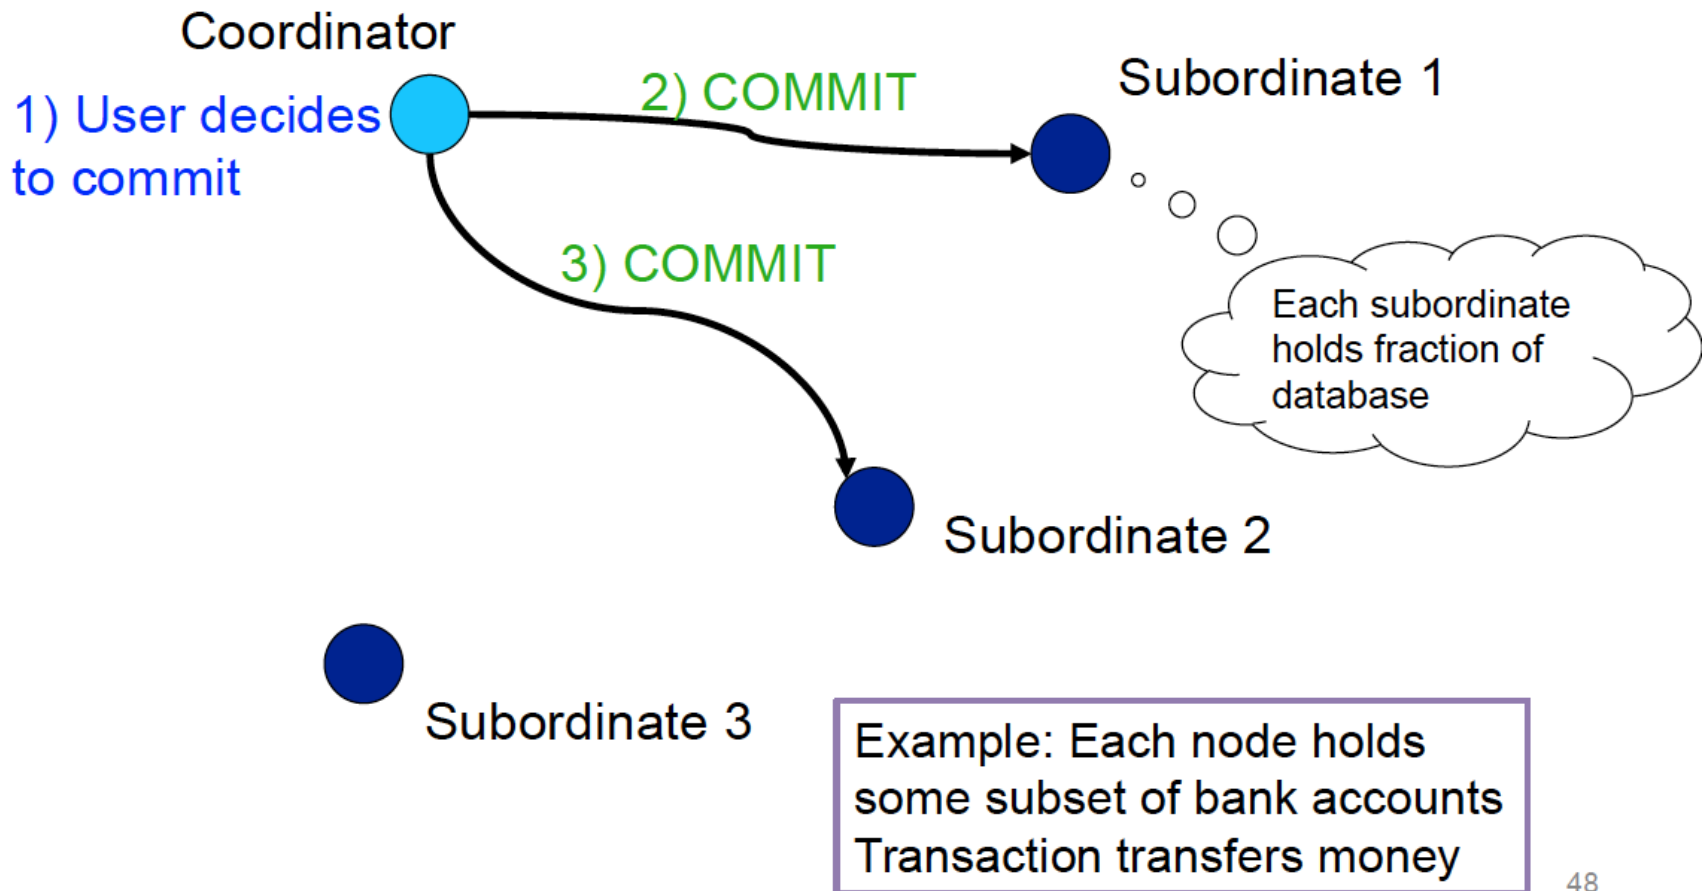
Two-Phase Commit: Motivation



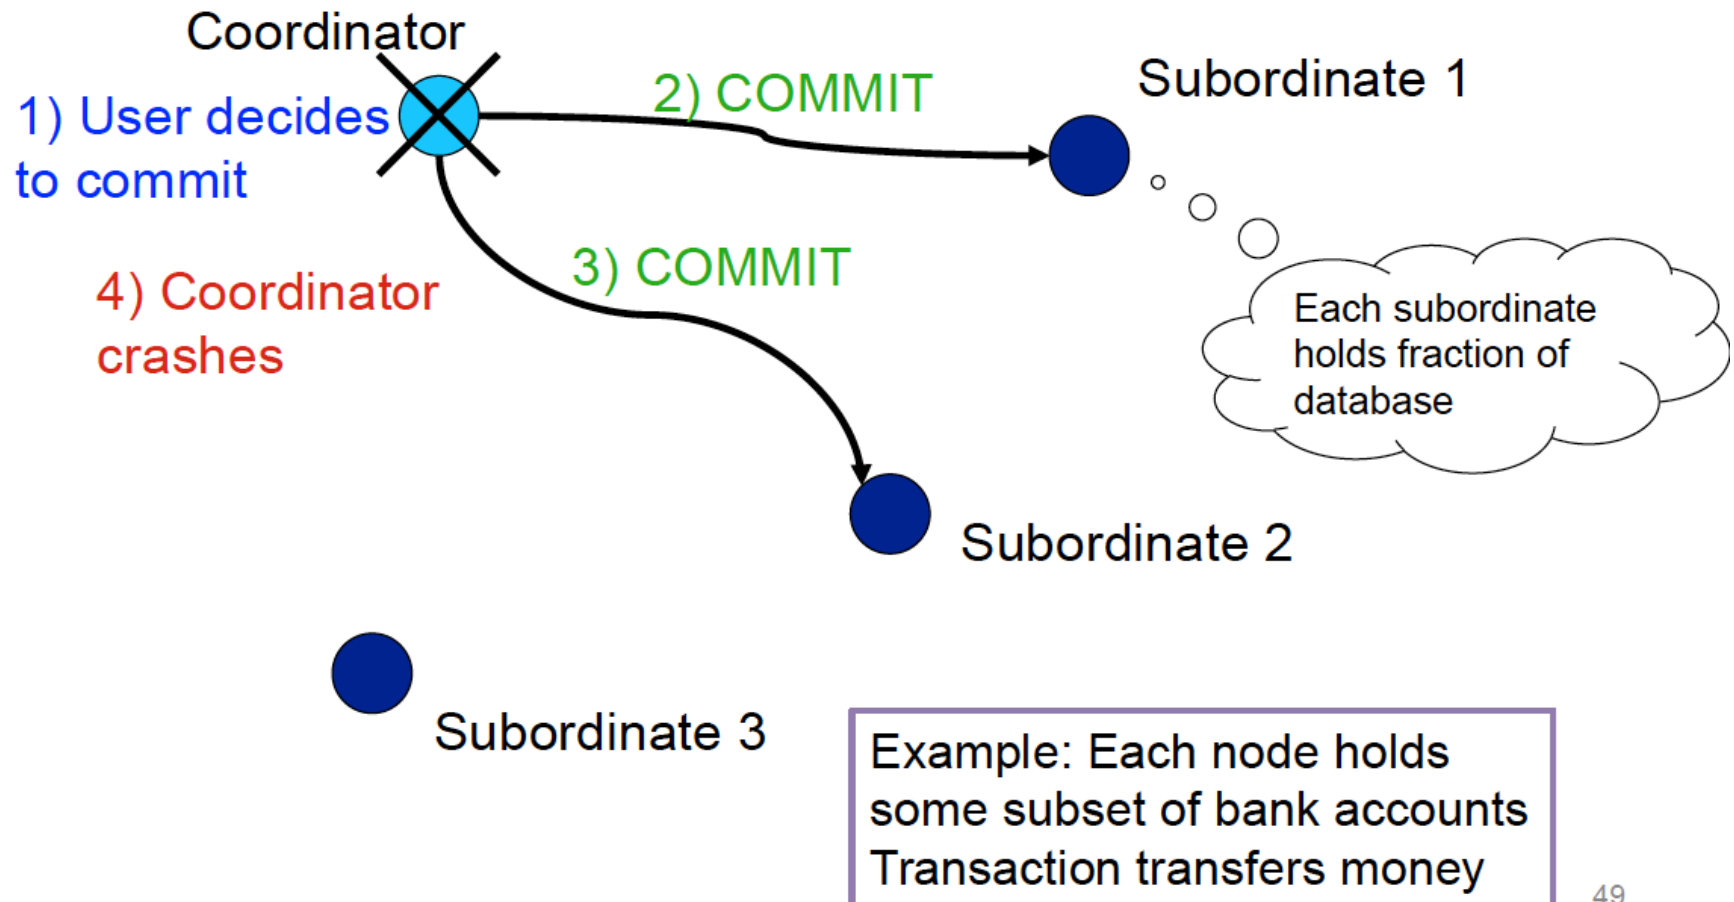
Two-Phase Commit: Motivation



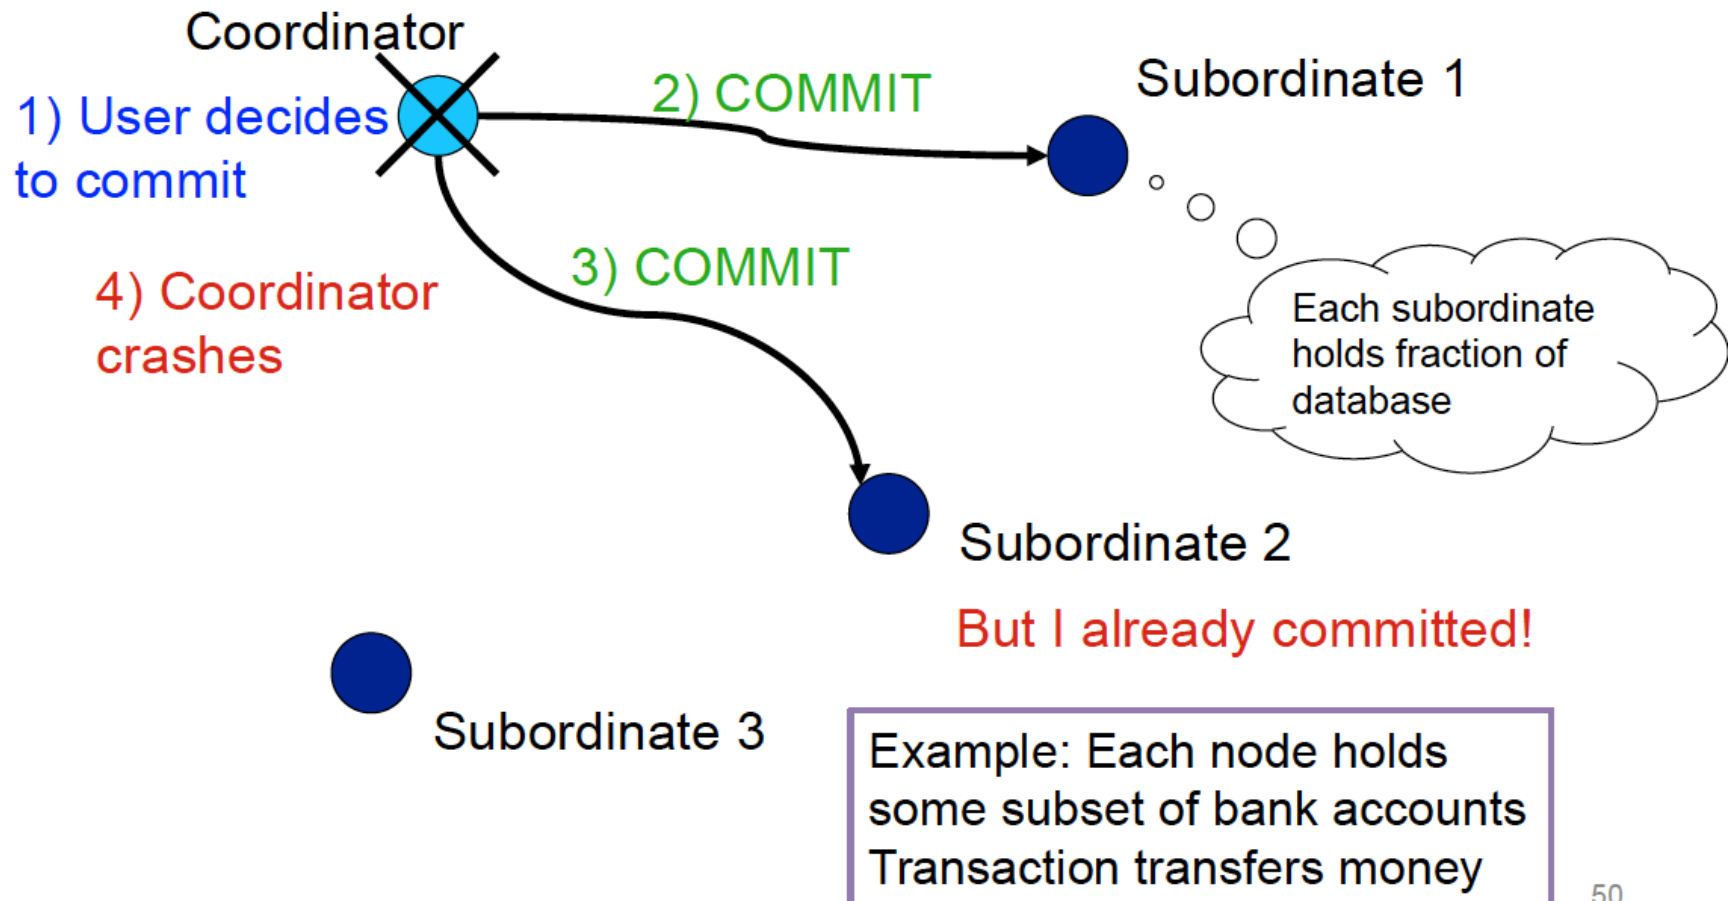
Two-Phase Commit: Motivation



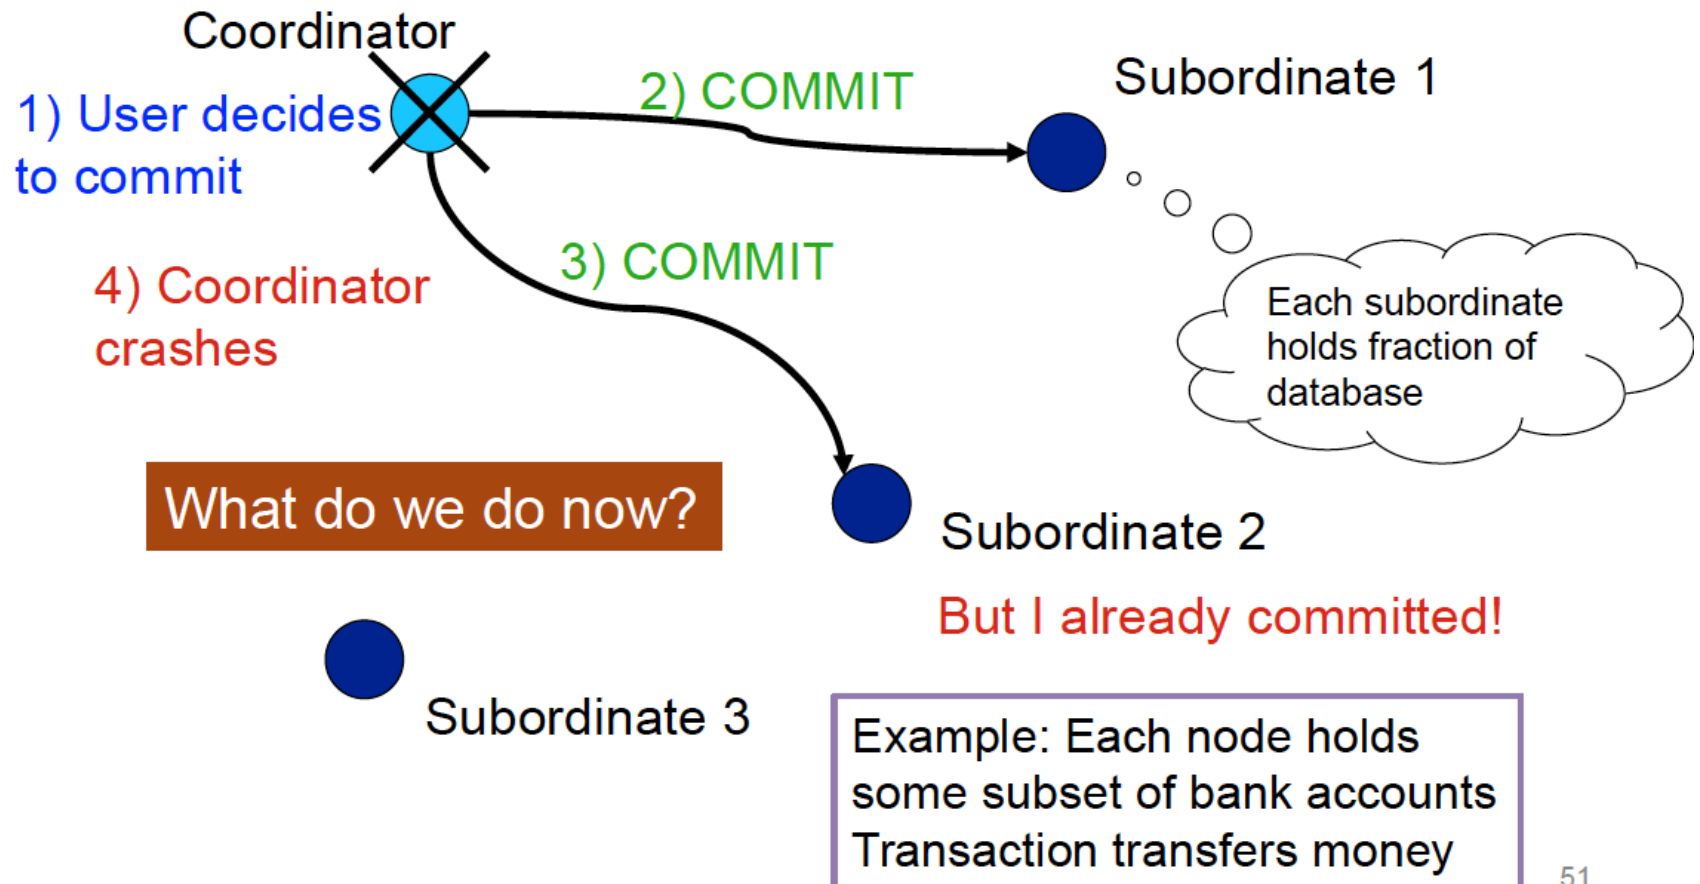
Two-Phase Commit: Motivation



Two-Phase Commit: Motivation



Two-Phase Commit: Motivation



2-Phase Commit

➤ Phase 1:

- Coordinator tells participants to “prepare”
- Participants respond with yes/no votes
 - *Unanimity* required for yes!

➤ Phase 2:

- Coordinator disseminates result of the vote

➤ Need to do some logging for failure handling!

2PC: Phase 1 illustrated

Coordinator



Subordinate 1



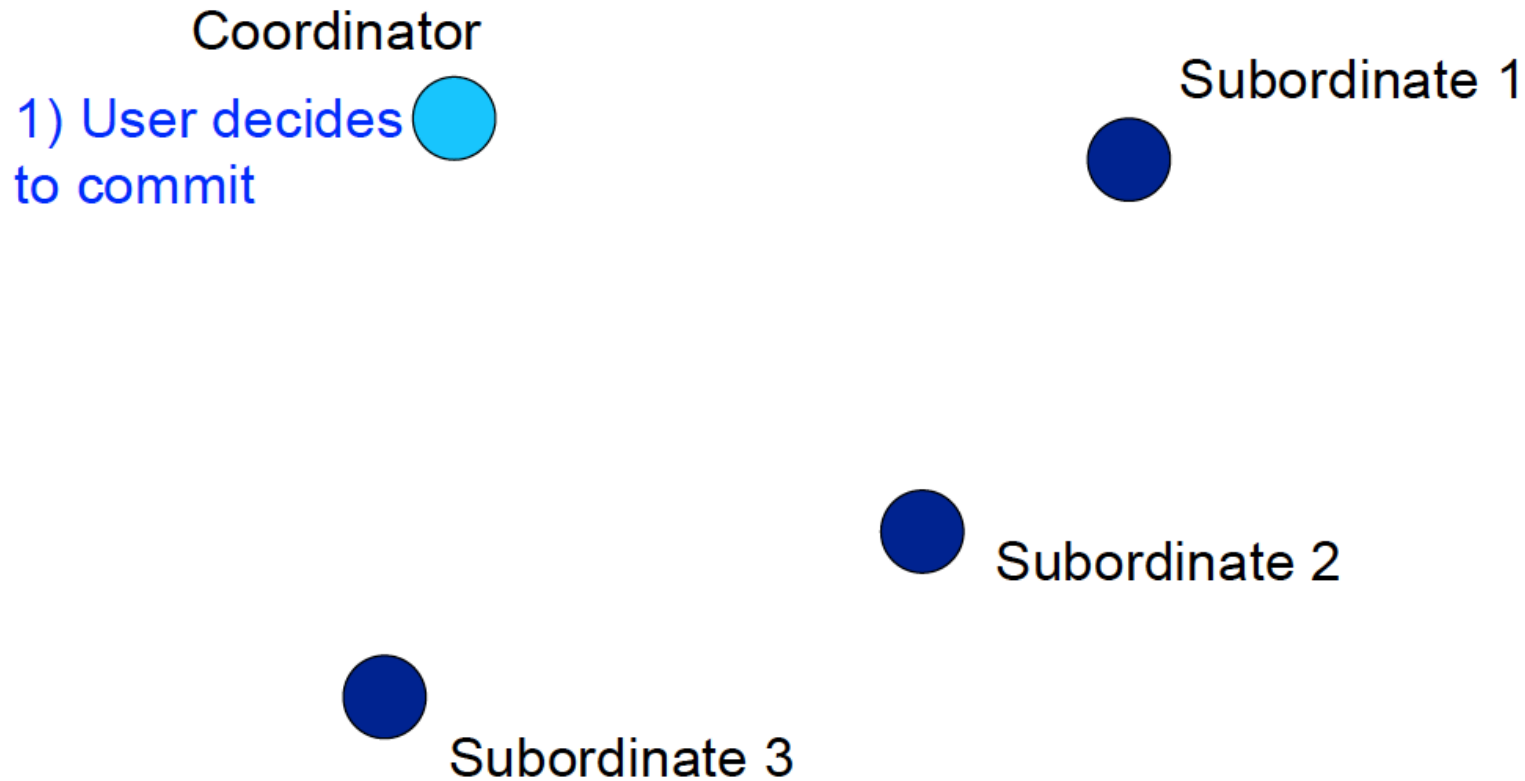
Subordinate 2



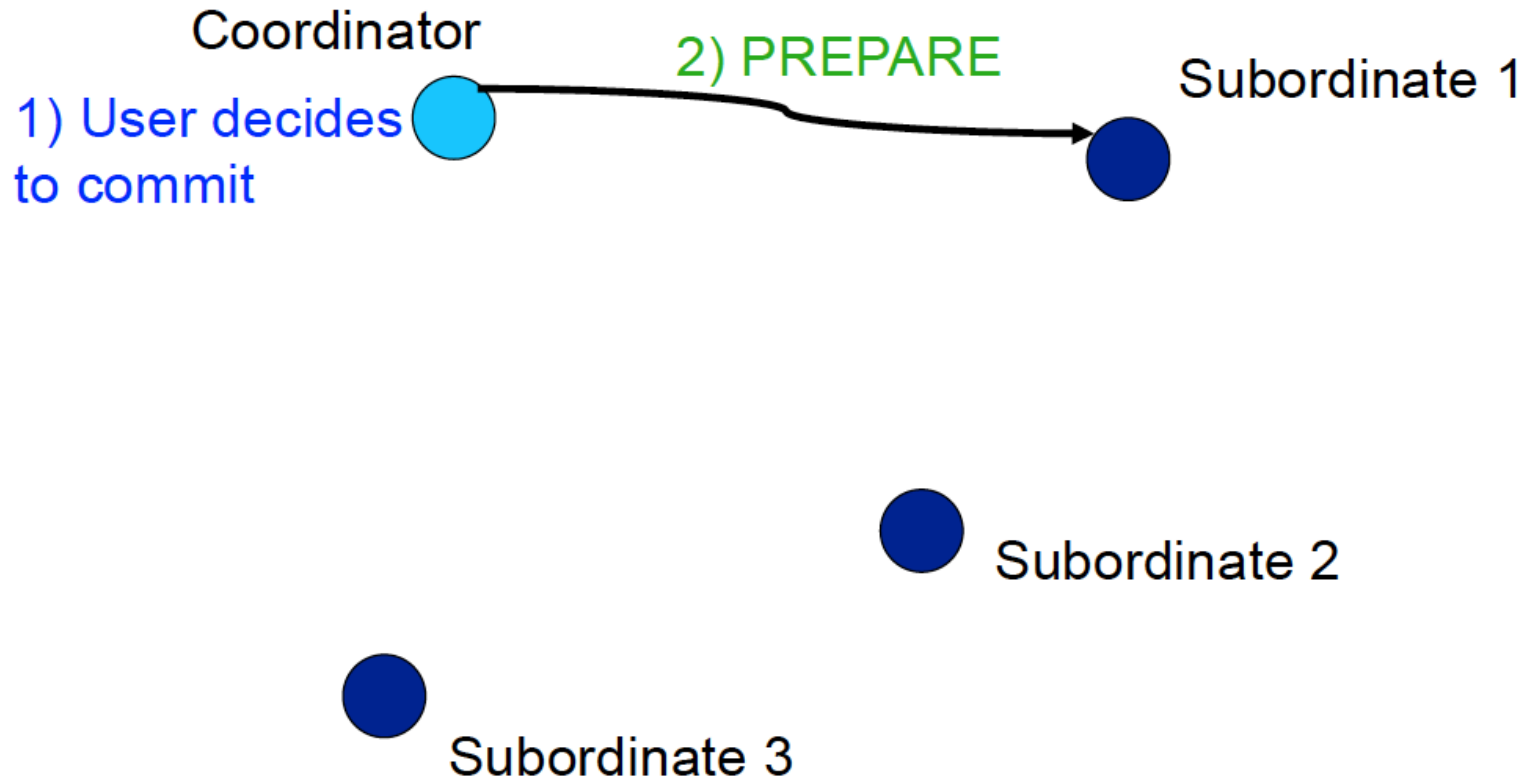
Subordinate 3



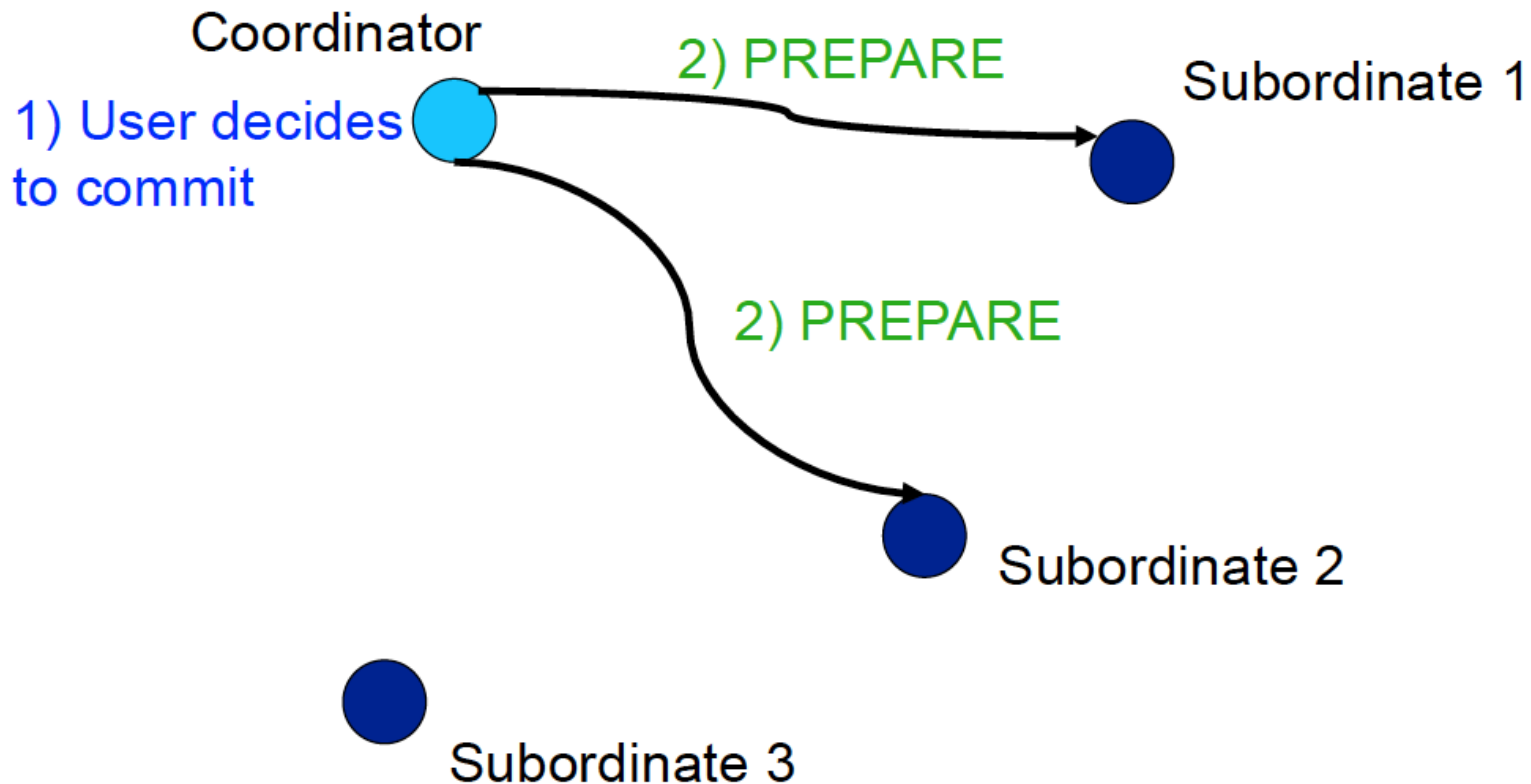
2PC: Phase 1 illustrated



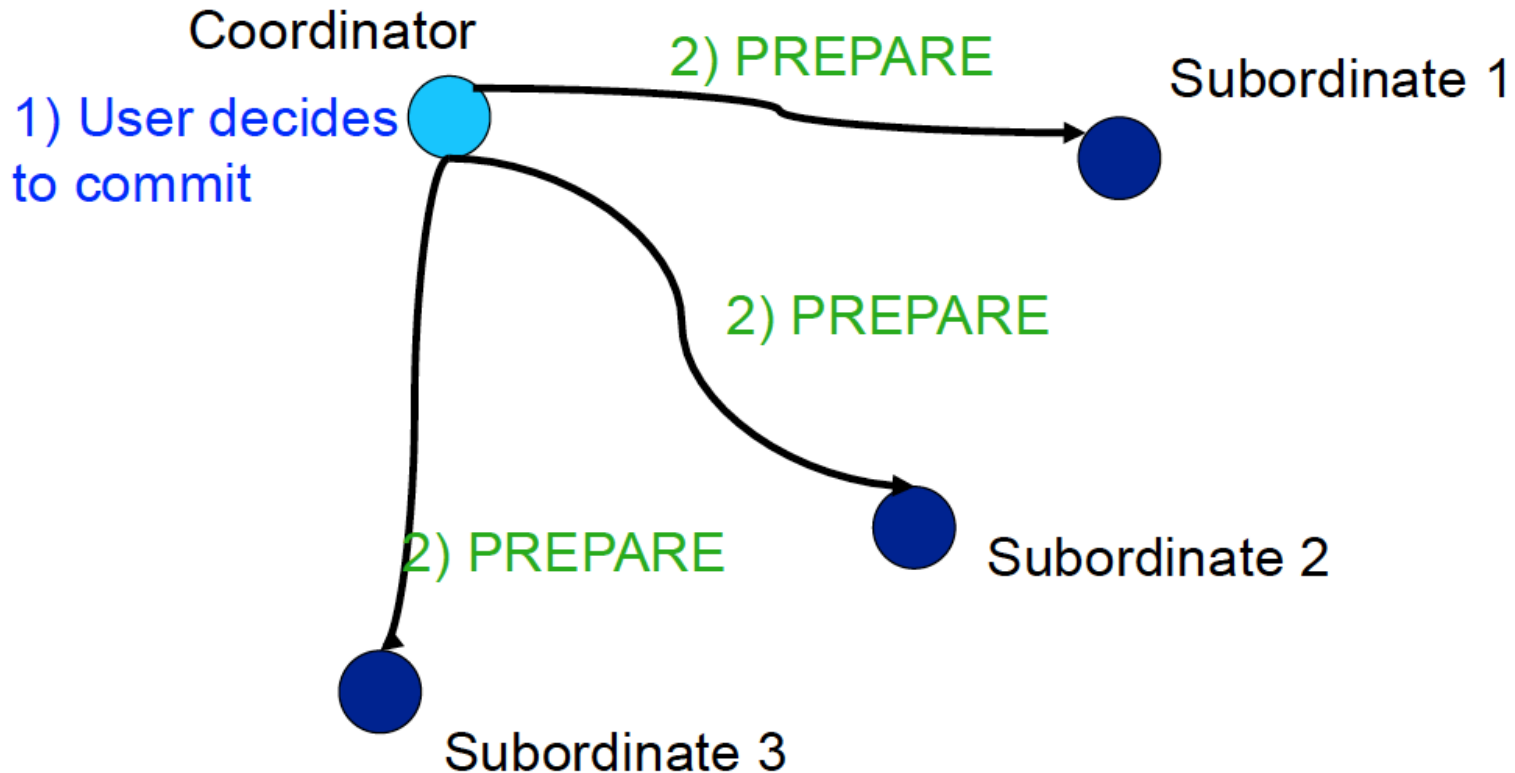
2PC: Phase 1 illustrated



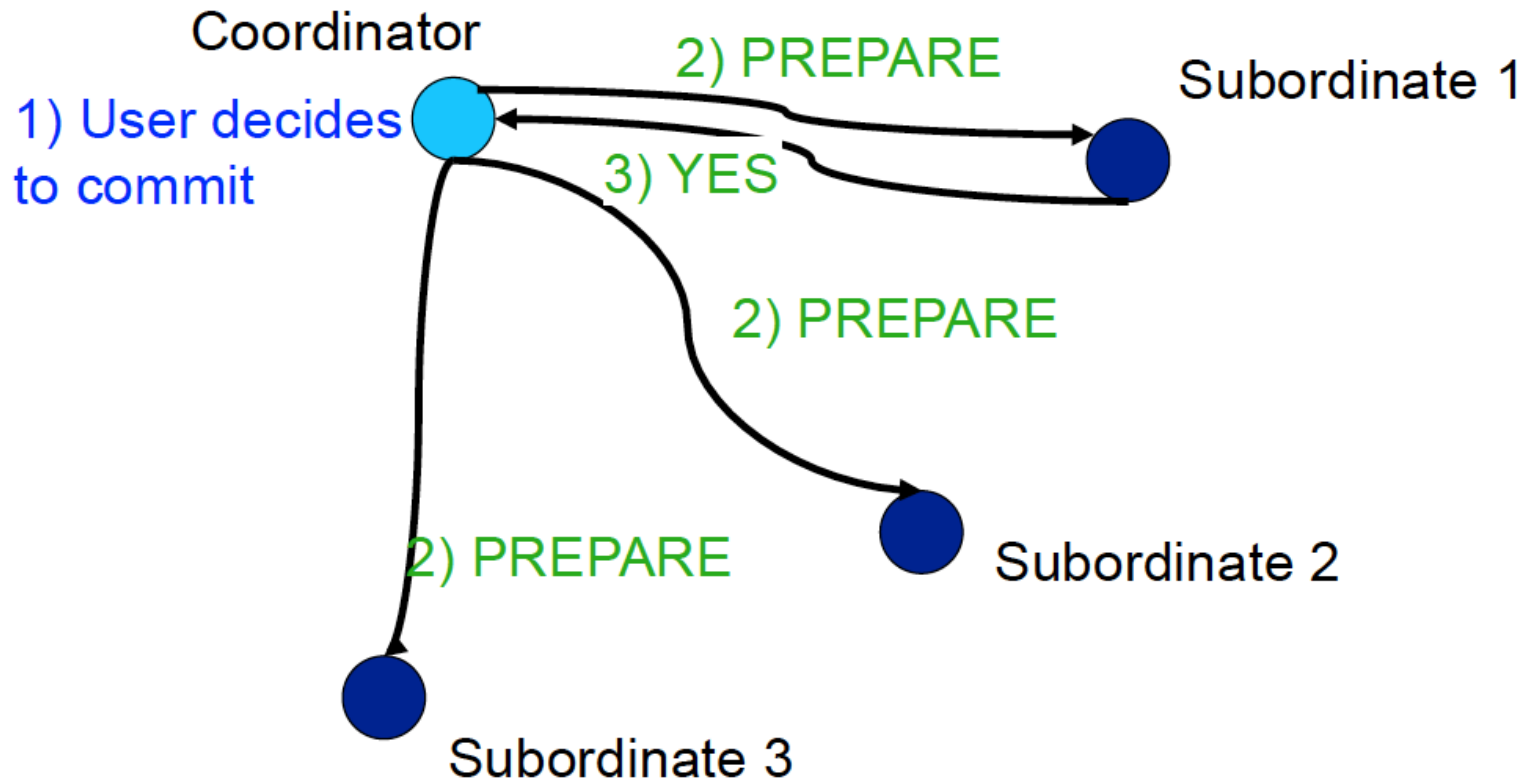
2PC: Phase 1 illustrated



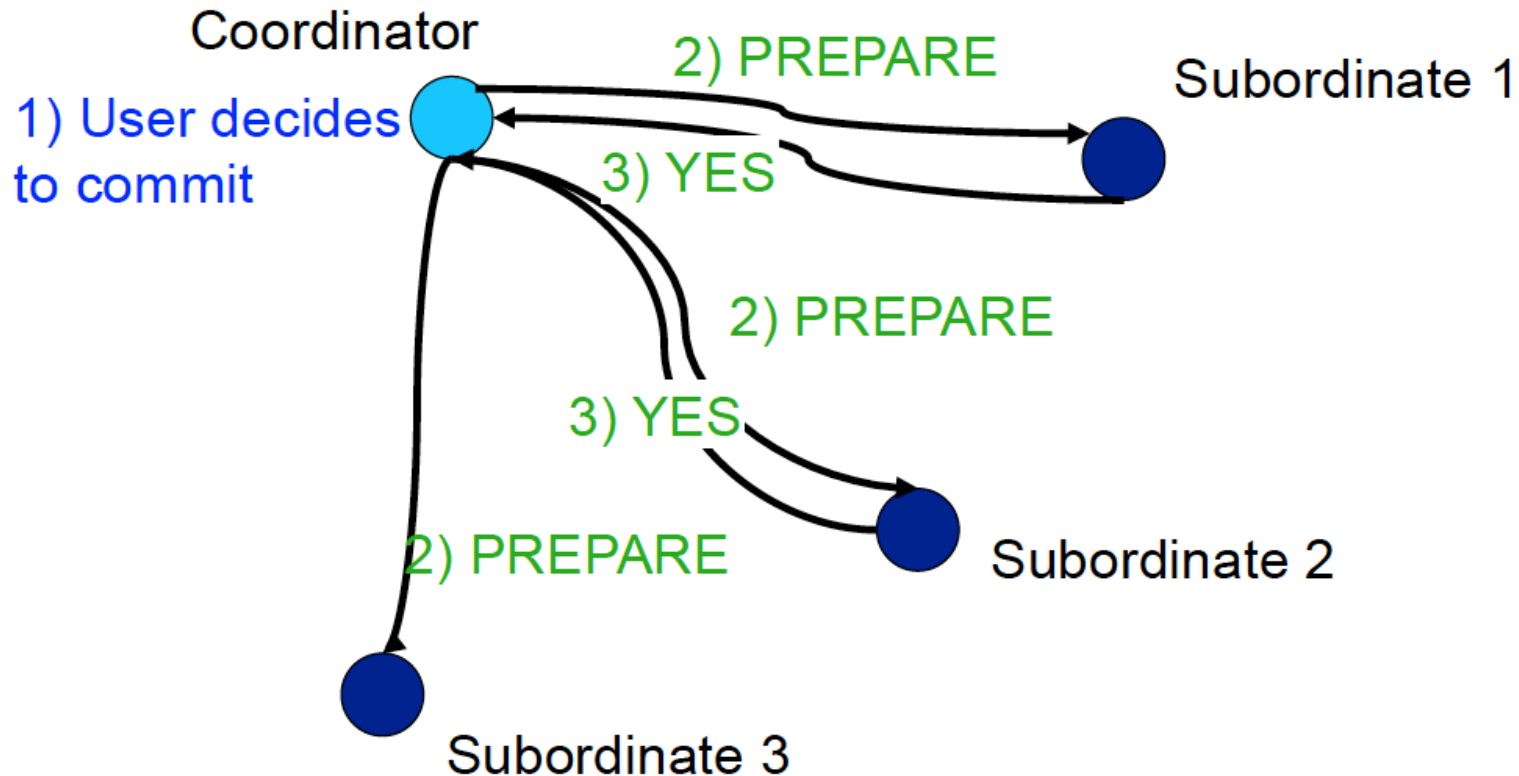
2PC: Phase 1 illustrated



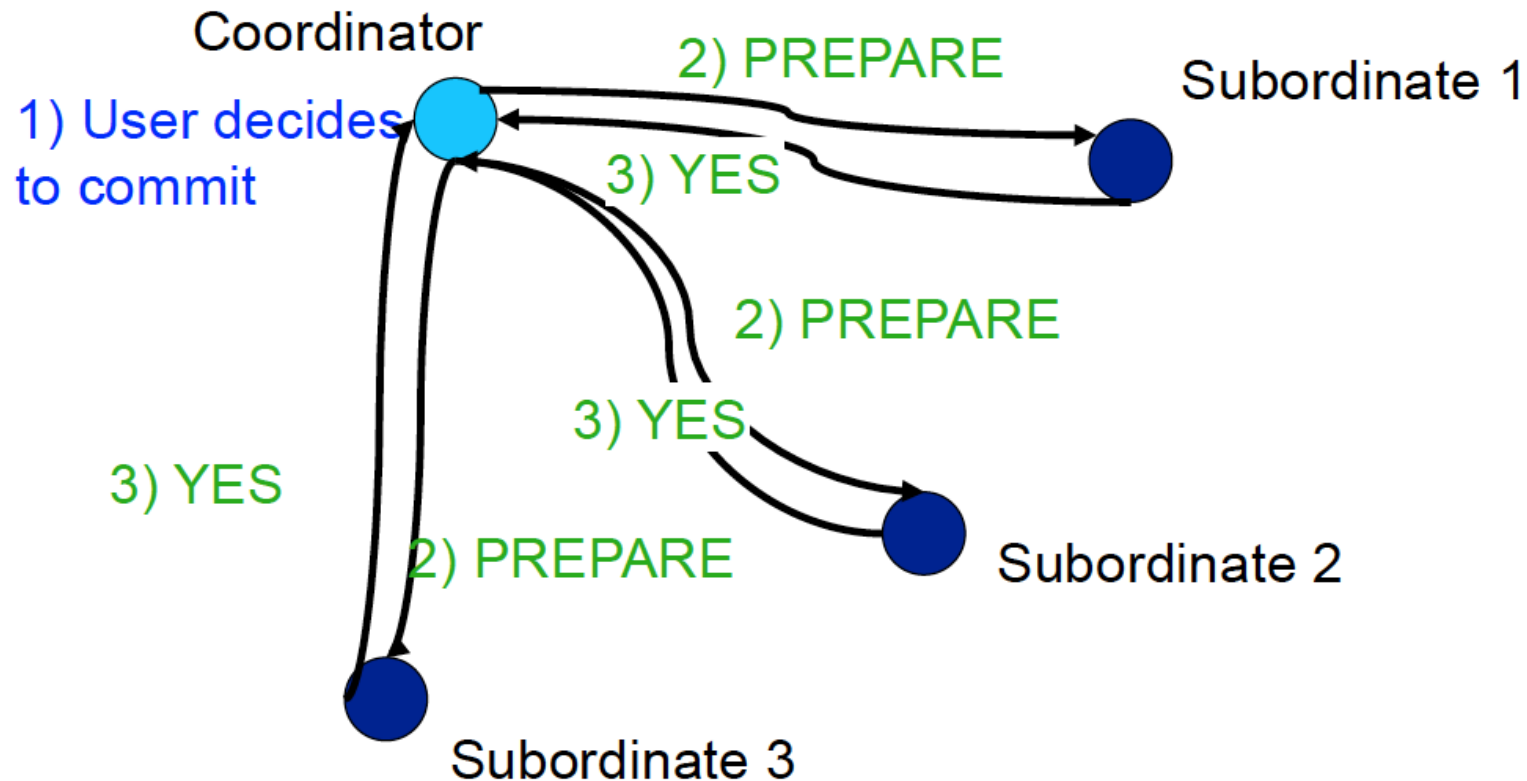
2PC: Phase 1 illustrated




2PC: Phase 1 illustrated




2PC: Phase 1 illustrated




2PC: Phase 2 illustrated




2PC: Phase 2 illustrated




2PC: Phase 2 illustrated



2PC: Phase 2 illustrated



2PC: Phase 2 illustrated



2-Phase Commit

- Multiple servers run parts of the same transaction
- They all must commit, or none should commit
- Two-phase commit is a complicated protocol that ensures that
- 2PC can also be used for WRITE with replication: commit the write at all replicas before declaring success

2-Phase Commit

Assumptions:

- Each site logs actions at that site, but there is no global log
- There is a special site, called the coordinator, which plays a special role
- 2PC involves sending certain messages: as each message is sent, it is logged at the sending site, to aid in case of recovery

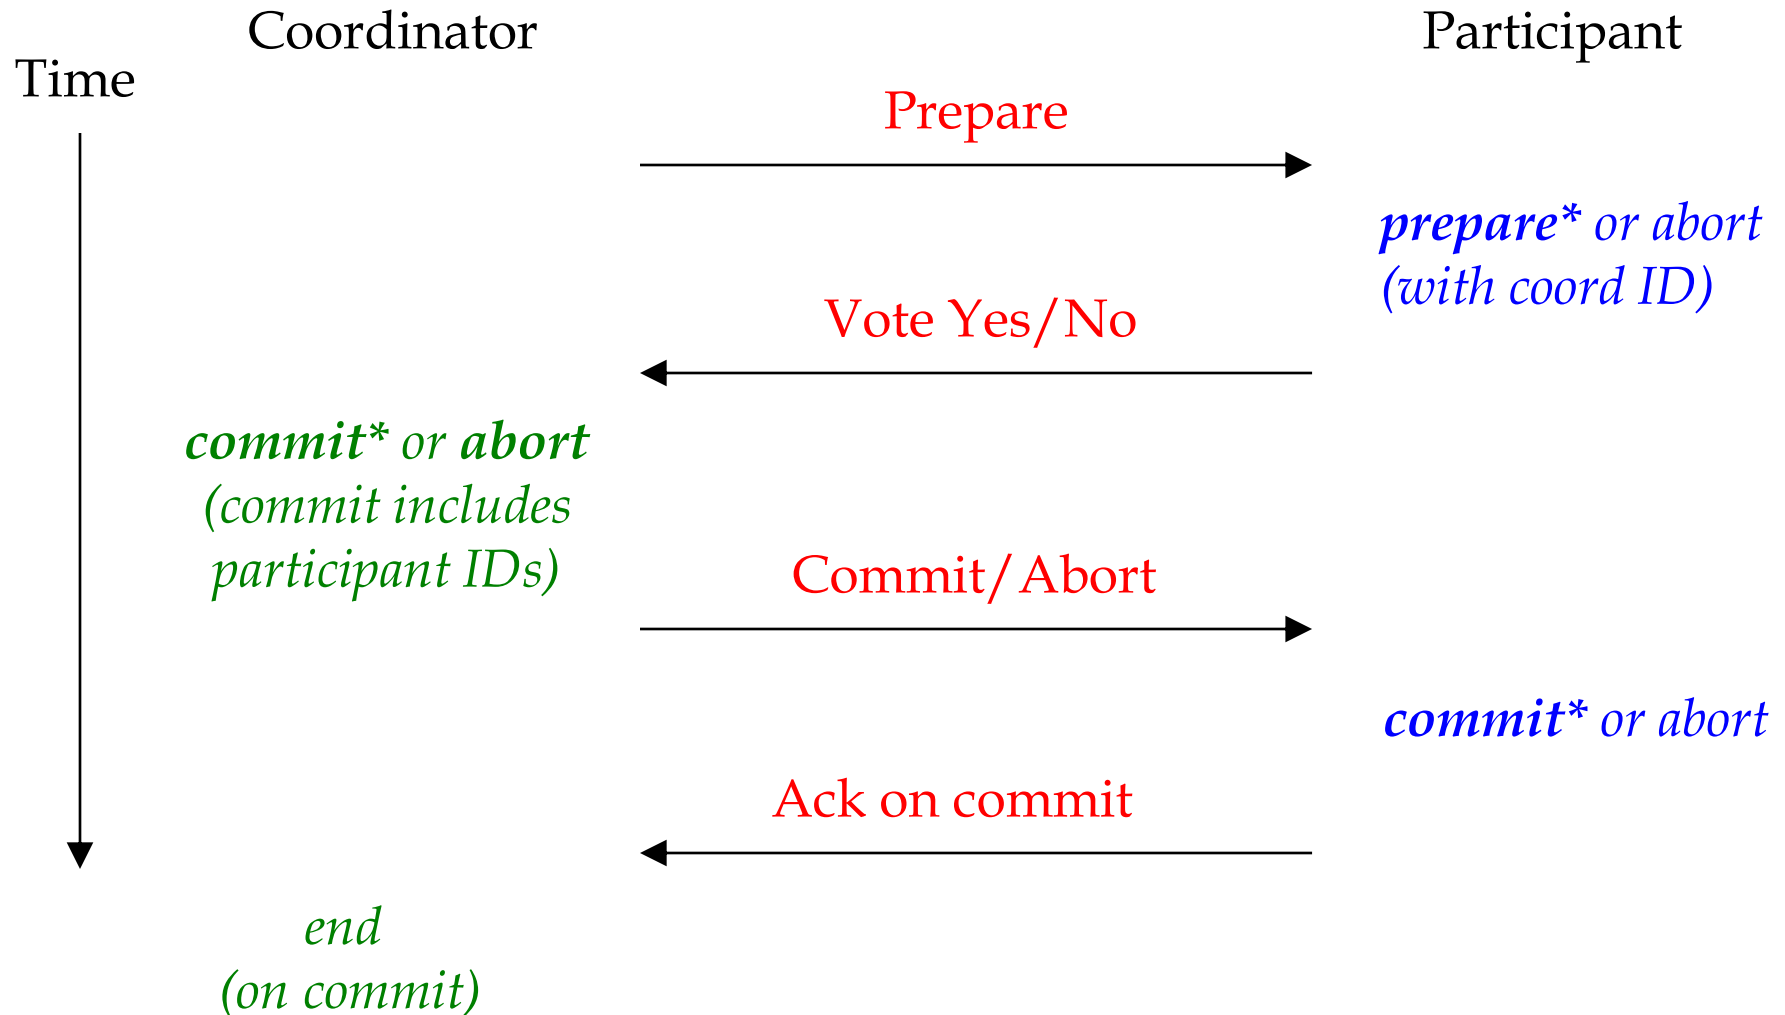
2-Phase Commit and Recovery

1. Coordinator sends prepare message
2. Subordinates receive prepare statement; force-write **<prepare>** log entry; answers yes or no
3. If coordinator receives only yes, force write **<commit>**, sends commit messages;
If at least one no, or timeout, force write **<abort>**, sends abort messages
4. If subordinate receives abort, force-write **<abort>**, sends ack message and aborts; if receives commit, force-write **<commit>**, sends ack, commits.
5. When coordinator receives all ack, writes **<end log>**

2PC: Messages

NOTE

*asterisk**: wait for log flush
before sending next msg



2-Phase Commit and Recovery

Restart after failure: each server recovers locally

1. If it finds a **<commit>** or **<abort>** log entry, then: redo or undo; if the server is coordinator, then re-request all *ack* messages, then write **<end log>**
2. If it finds a **<prepare>** entry, then re-contact the coordinator to ask for commit/abort
3. If no **<prepare>** , **<commit>** or **<abort>**, presume abort

2-Phase Commit: Summary

- Transaction: ACID properties, but expensive
- Relies on central coordinator: both performance bottleneck, and single-point-of-failure
- Solution: Paxos = distributed protocol
 - Complex: will not discuss at all

NoSQL Databases

NoSQL

➤ Not Only SQL

- Not the other thing!
- Term introduced by Carlo Strozzi in 1998 to describe an alternative database model
- Became **the name of a movement** following Eric Evans's reuse for a distributed-database event

➤ Seminal papers:

- Google's BigTable
 - Chang, Dean, Ghemawat, Hsieh, Wallach, Burrows, Chandra, Fikes, Gruber: Bigtable: A Distributed Storage System for Structured Data. OSDI 2006: 205-218
- Amazon's DynamoDB
 - DeCandia, Hastorun, Jampani, Kakulapati, Lakshman, Pilchin, Sivasubramanian, Vosshall, Vogels: Dynamo: amazon's highly available key-value store. SOSP 2007: 205-220

NoSQL from nosql-database.org

“

- Next Generation Databases mostly addressing some of the points: being *non-relational*, *distributed*, *open-source* and *horizontally scalable*.
- The original intention has been modern web-scale databases. The movement began early 2009 and is growing rapidly. Often more characteristics apply such as: *schema-free*, *easy replication support*, *simple API*, eventually consistent / *BASE (not ACID)*, a huge amount of data and more.
- So the misleading term “nosql” (the community now translates it mostly with “not only sql”) should be seen as an alias to something like the definition above.

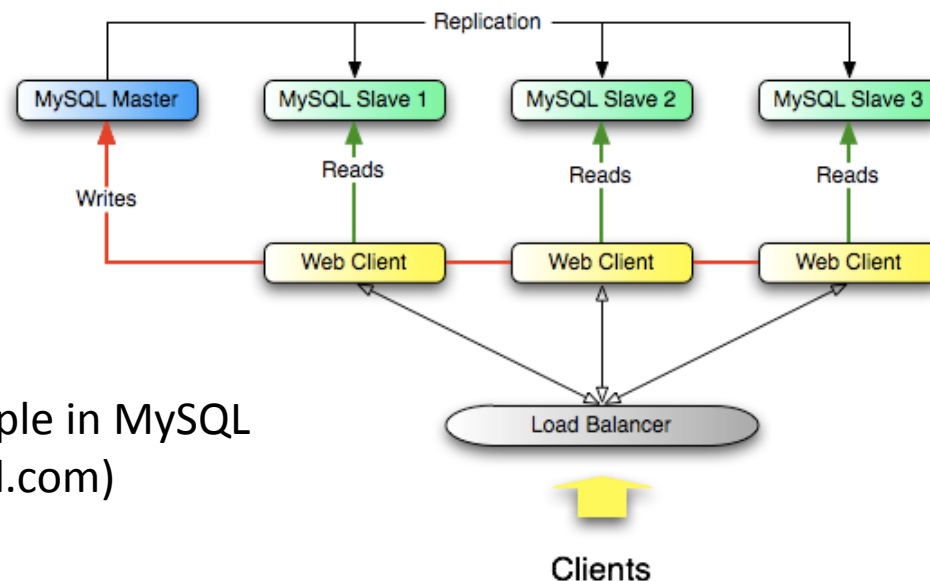
”

Common NoSQL Features

- Non-relational data models
- Flexible structure
 - No need to fix a **schema**, attributes can be added and replaced on the fly
- Massive read/write performance; availability via **horizontal scaling**
 - **Replication** and **sharding** (data partitioning)
 - Potentially thousands of machines worldwide
- Open source (very often)
- APIs to impose **locality**

Database Replication

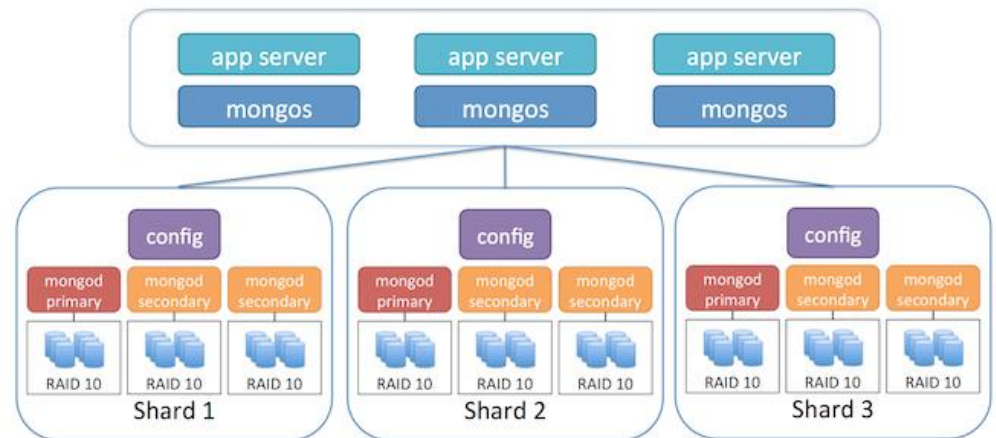

- Data replication: storing the same data on several machines (“nodes”)
- Useful for:
 - **Availability** (parallel requests are made against replicas)
 - **Reliability** (data can survive hardware faults)
 - **Fault tolerance** (system stays alive when nodes/network fail)
- Typical architecture: master-slave



Replication example in MySQL
(dev.mysql.com)

Database Sharding

- Simply partitioning data across multiple nodes
- Useful for
 - **Scaling** (more data)
 - **Availability**



Replication + sharding example in MongoDB
(mongodb-documentation.readthedocs.org)

True and False Conceptions

➤ True:

- SQL does not effectively handle common Web needs of massive (datacenter) data
- SQL has guarantees that can sometimes be compromised for the sake of scaling
- Joins are not for free, sometimes undoable

➤ False:

- NoSQL says NO to SQL
- Nowadays NoSQL is the only way to go
- Joins can always be avoided by structure redesign

Highlighted Database Features

➤ Data model

- What data is being stored?

➤ CRUD interface


- API for Create, Read, Update, Delete
- Sometimes preceding S for Search

➤ Transaction consistency guarantees


➤ Replication and sharding model

- What's automated and what's manual?


We Will Look at 4 Data Models




Key/Value Store



Column-Family Store



Document Store



Graph Databases