Database Management Systems
Chapter 4: Logical Database Design

UNIVERSITY OF
MARYLAND

ROBERT H. SMITH
SCHOOL OF BUSINESS

DR. ADAM LEE

Objectives

- Define terms.
- List five properties of relations.
- State two properties of candidate keys.
- Define first, second, and third normal form.
- Describe problems from merging relations.
- Transform ER and EER diagrams to relations.
- Create tables with entity and relational integrity constraints.
- Use normalization to convert anomalous tables to well-structured relations.

UNIVERSITY OF MARYLAND

Components of Relational Model

- Data structure
  - Tables (relations)
  - Rows
  - Columns
- Data manipulation:
  - Powerful SQL operations for retrieving and modifying data.
- Data integrity:
  - Mechanisms for implementing business rules that maintain integrity of manipulated data.

Relation

- A relation is a named, two-dimensional table of data.
- A table consists of rows (records) and columns (attributes or fields).
- Requirements for a table to qualify as a relation:
  - It must have a unique name.
  - Every attribute value must be atomic (not multivalued, not composite).
  - Every row must be unique.
  - Attributes (columns) in tables must have unique names.
  - The orders of columns and rows must be irrelevant.
- NOTE: All relations are in 1st Normal form (1NF).

Correspondence with ER Model

- Relations (tables) correspond with:
  - Entity types
  - Many-to-many relationship types
- Rows correspond with:
  - Entity instances
  - Many-to-many relationship instances
- Columns correspond with:
  - Attributes

- NOTE: The word relation (in relational database) is NOT the same as the word relationship (in ER model).

Key Fields

- Keys are special fields that serve two main purposes:
  - **Primary keys** are unique identifiers of the relation. Examples include employee numbers, social security numbers, etc. This guarantees that all rows are unique.
  - **Foreign keys** are identifiers that enable a dependent relation (on the many side of a relationship) to refer to its parent relation (on the one side of the relationship).
- Keys can be simple (a single attribute) or composite (a set of attributes).
- Keys usually are used as indexes to speed up the response to user queries (more on this in Chapter 5).

Figure 4-3: Schema for Four Relations

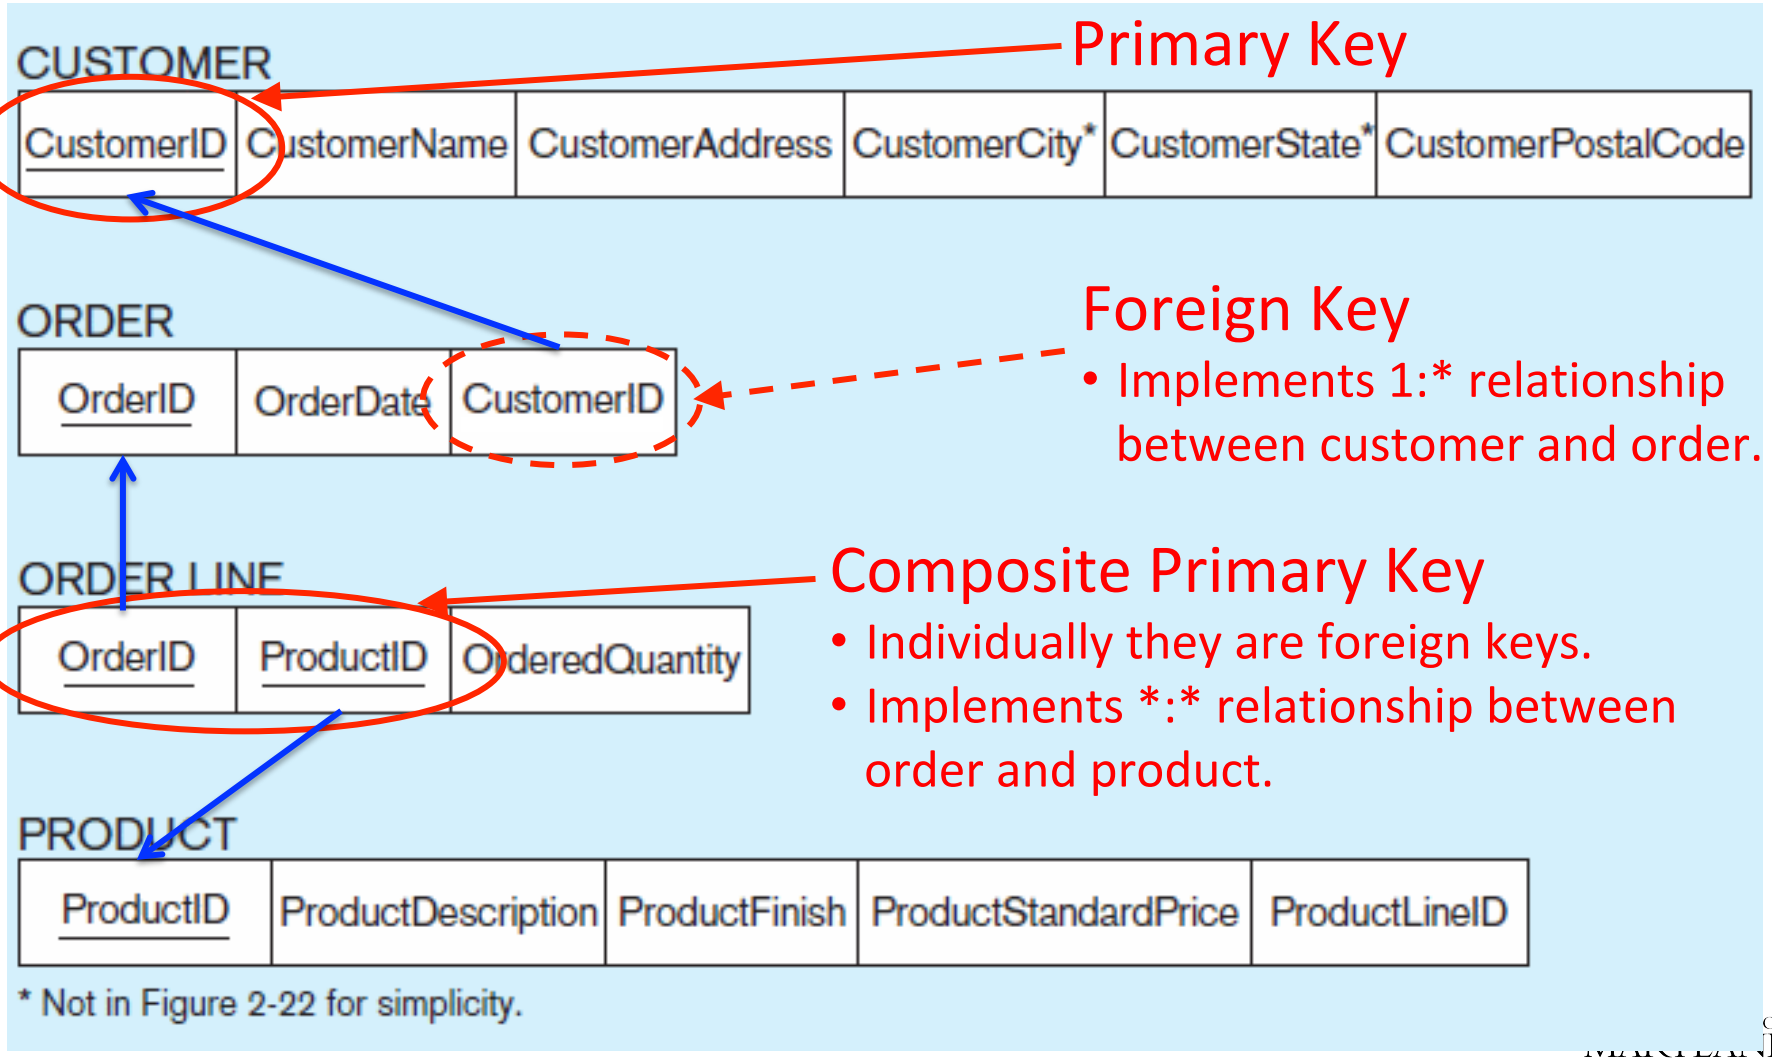

Integrity Constraints

- Atomic Constraint.

- **Domain Constraints**:
  - Allowable format and size for an attribute.
  - Allowable values for an attribute.

- Entity Integrity.
- Referential Integrity.

Table 4-1: Domain Definitions Enforce Domain Integrity Constraints

TABLE 4-1 Domain Definitions for INVOICE Attributes

Attribute	Domain Name	Description	Domain
CustomerID	Customer IDs	Set of all possible customer IDs	character: size 5
CustomerName	Customer Names	Set of all possible customer names	character: size 25
CustomerAddress	Customer Addresses	Set of all possible customer addresses	character: size 30
CustomerCity	Cities	Set of all possible cities	character: size 20
CustomerState	States	Set of all possible states	character: size 2
CustomerPostalCode	Postal Codes	Set of all possible postal zip codes	character: size 10
OrderID	Order IDs	Set of all possible order IDs	character: size 5
OrderDate	Order Dates	Set of all possible order dates	date: format mm/dd/yy
ProductID	Product IDs	Set of all possible product IDs	character: size 5
ProductDescription	Product Descriptions	Set of all possible product descriptions	character: size 25
ProductFinish	Product Finishes	Set of all possible product finishes	character: size 15
ProductStandardPrice	Unit Prices	Set of all possible unit prices	monetary: 6 digits
ProductLineID	Product Line IDs	Set of all possible product line IDs	integer: 3 digits
OrderedQuantity	Quantities	Set of all possible ordered quantities	integer: 3 digits

Integrity Constraints

- Atomic Constraint.
- Domain Constraints.

- **Entity Integrity**:
  - Every relation has a primary key.
  - Every primary key attribute is NON-null.
  - All primary key attributes MUST have valid data.

- Referential Integrity.

Integrity Constraints

- Atomic Constraint.
- Domain Constraints.
- Entity Integrity.

- **Referential Integrity**:
  - To maintain consistency among the rows of two relations.
  - If there is a foreign key in one relation, either each foreign key value MUST match a primary key value in one relation, or the foreign key value MUST be null.

Referential Integrity

- **ON DELETE** rules
  - **RESTRICT | NO ACTION** – do not allow delete of "primary" side if related rows exist in "foreign" side
  - **CASCADE** – automatically delete "foreign" side rows that correspond with the "primary" side row to be deleted
  - **SET NULL** – set the foreign key in the "foreign" side to null if deleting from the "primary" side → not allowed for weak entities

- **ON UPDATE** rules
  - **RESTRICT | NO ACTION**
  - **CASCADE**
  - **SET NULL**

Figure 4-5: Referential Integrity Constraints

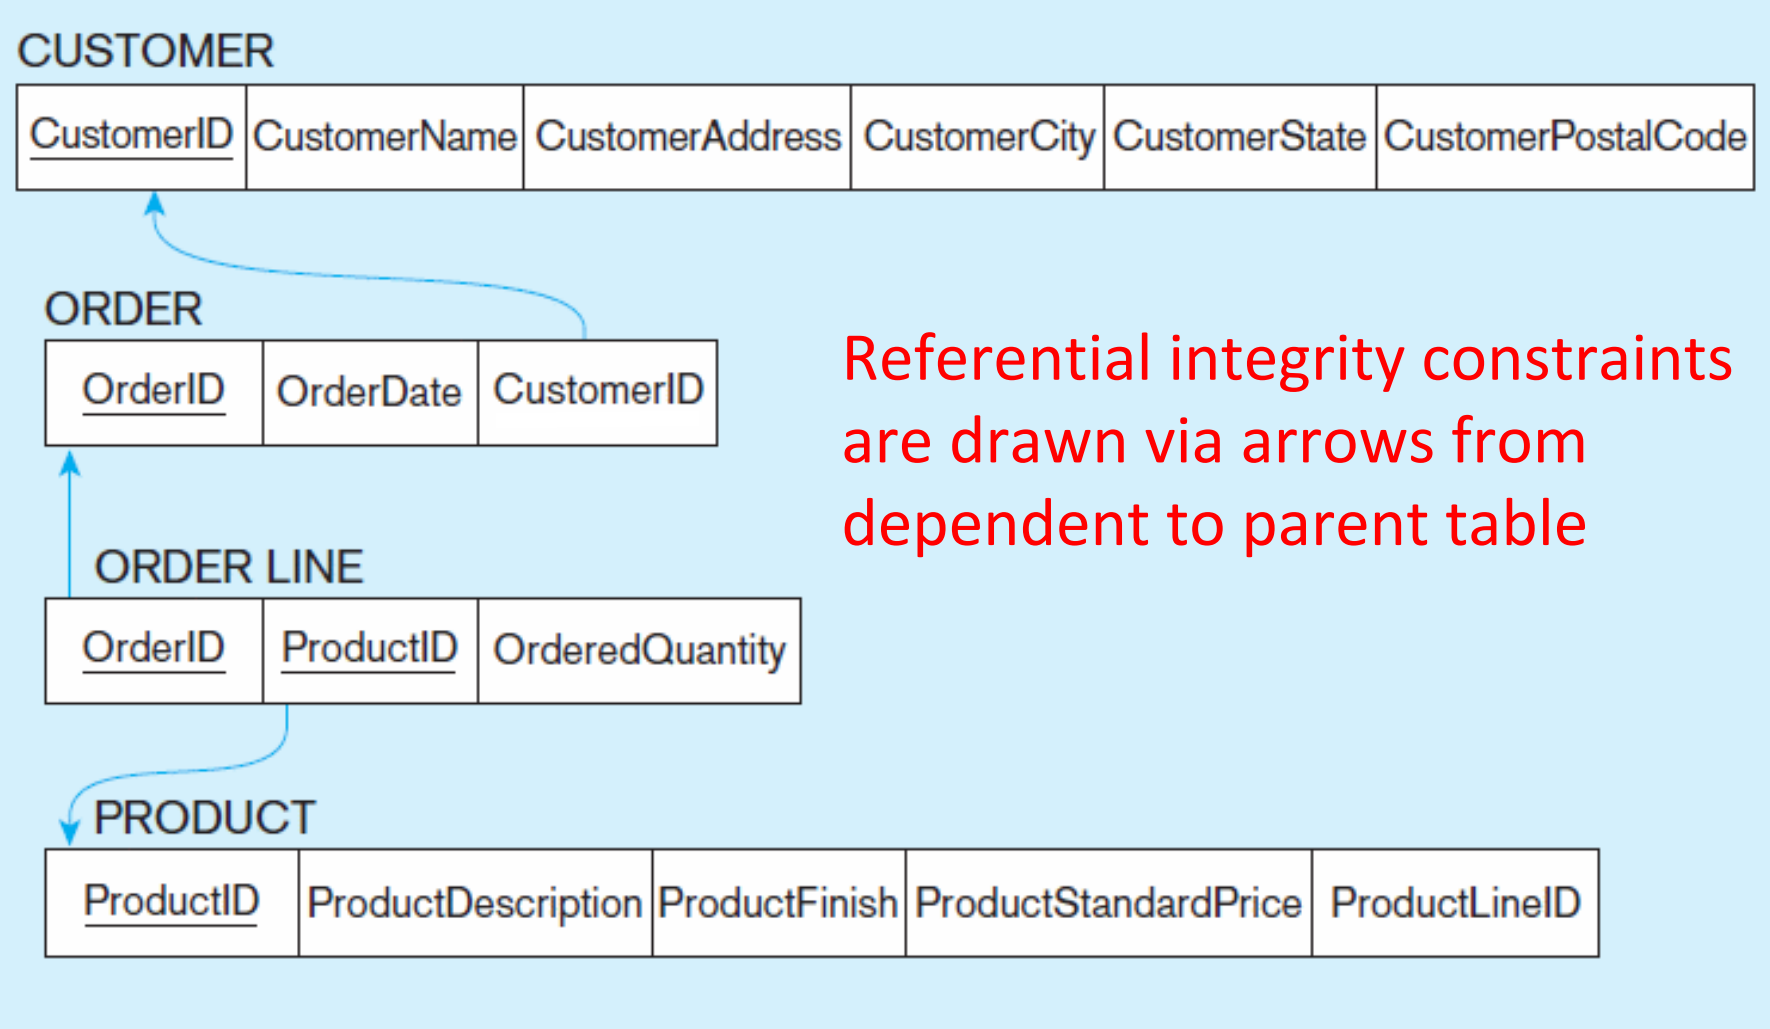

Transforming ER Diagrams into Relations Step 1

- Mapping regular entities to relations:
  - **Simple attributes**: ER attributes map directly onto the relation.
  - **Composite attributes**: Use only their simple, component attributes.
  - **Multivalued attribute**: Becomes a separate relation with a foreign key taken from the superior entity.

Figure 4-8 Mapping a Regular Entity

(a) CUSTOMER entity type with simple attributes

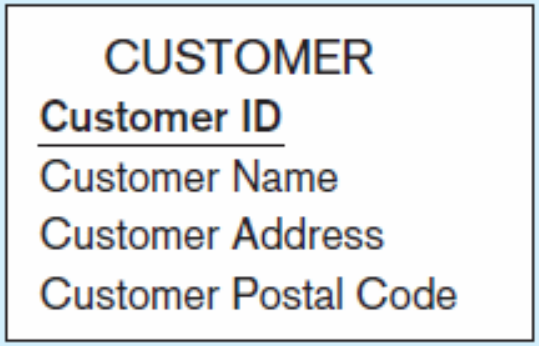

(b) CUSTOMER relation

CUSTOMER

CustomerID	CustomerName	CustomerAddress	CustomerPostalCode

Figure 4-9 Mapping a Composite Attribute

(a) CUSTOMER entity type with composite attribute

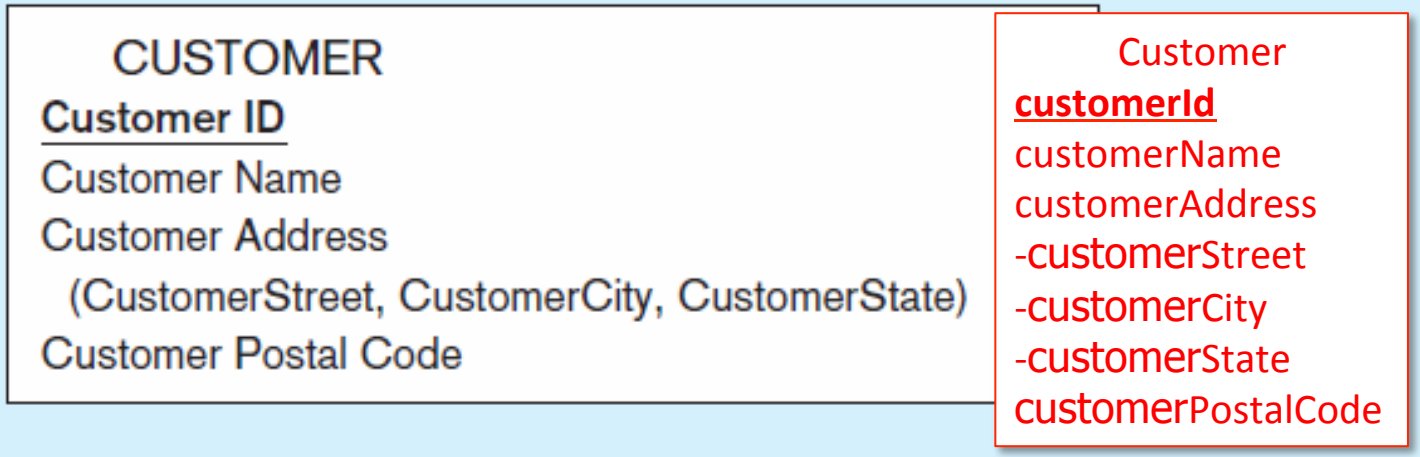

(b) CUSTOMER relation with address detail

CUSTOMER

CustomerID	CustomerName	CustomerStreet	CustomerCity	CustomerState	CustomerPostalCode

Figure 4-10 Mapping a Multivalued Attribute

(a) EMPLOYEE entity type with multivalued attribute

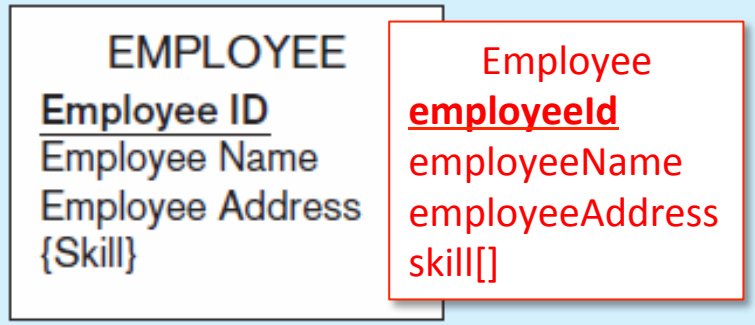

(b) Multivalued attribute becomes a separate relation with foreign key

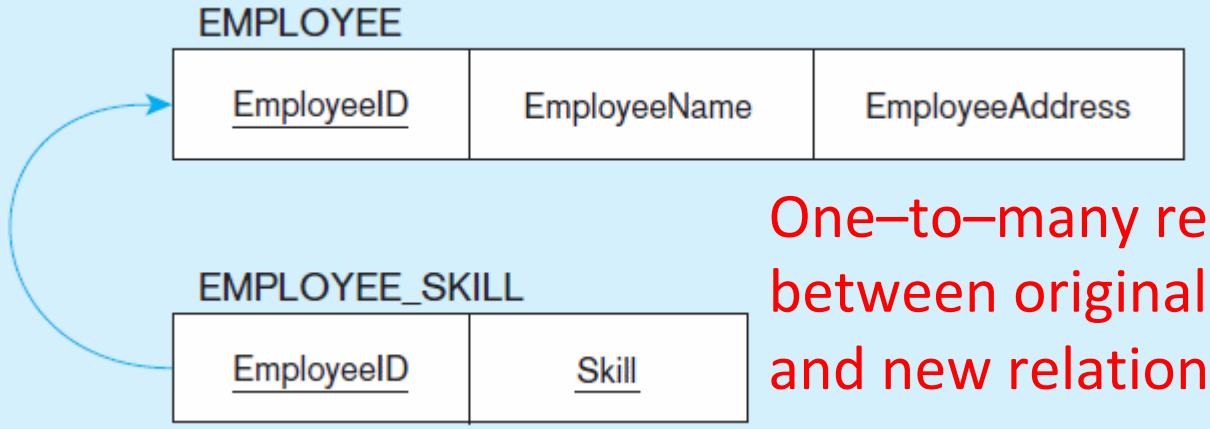

One–to–many relationship between original entity and new relation.

Transforming ER Diagrams into Relations Step 2

- Mapping weak entities:
  - Becomes a separate relation with a foreign key taken from the superior entity.
  - Primary key composed of (i) partial identifier of weak entity, and (ii) primary key of identifying relation (i.e. strong entity).

Figure 4-11 Example of Mapping a Weak Entity

(a) Weak entity DEPENDENT

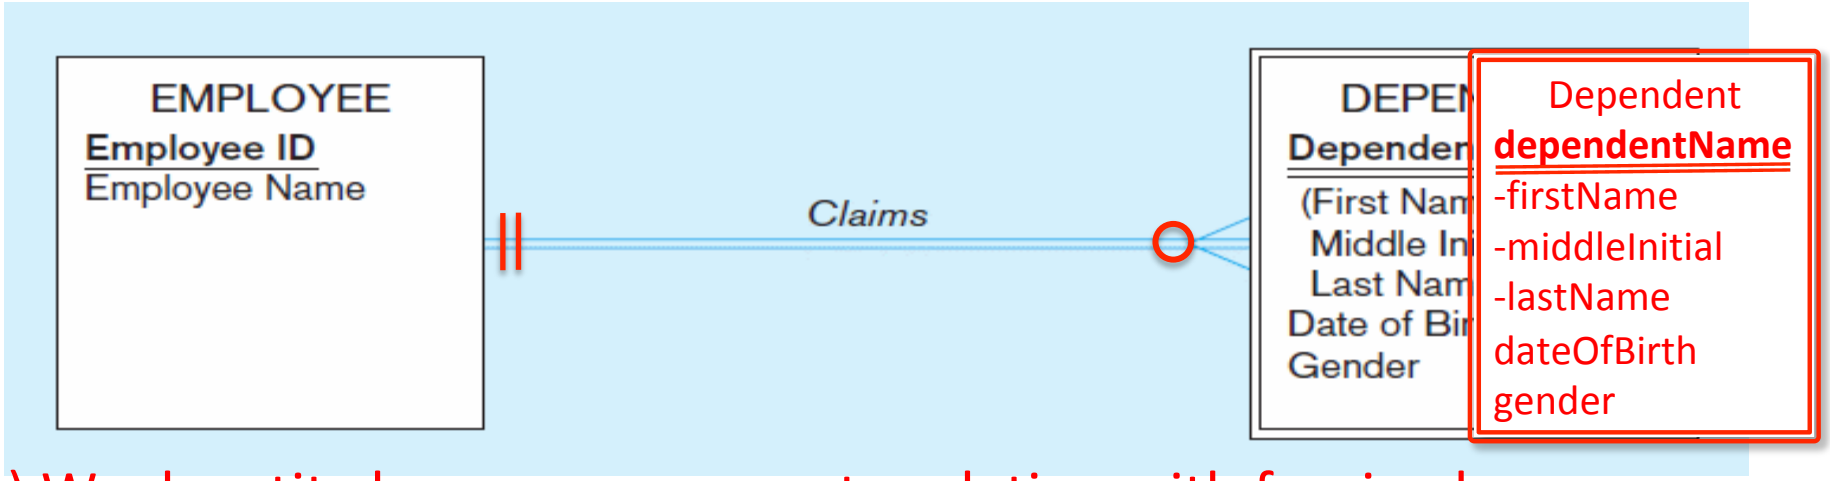

(b) Weak entity becomes a separate relation with foreign key

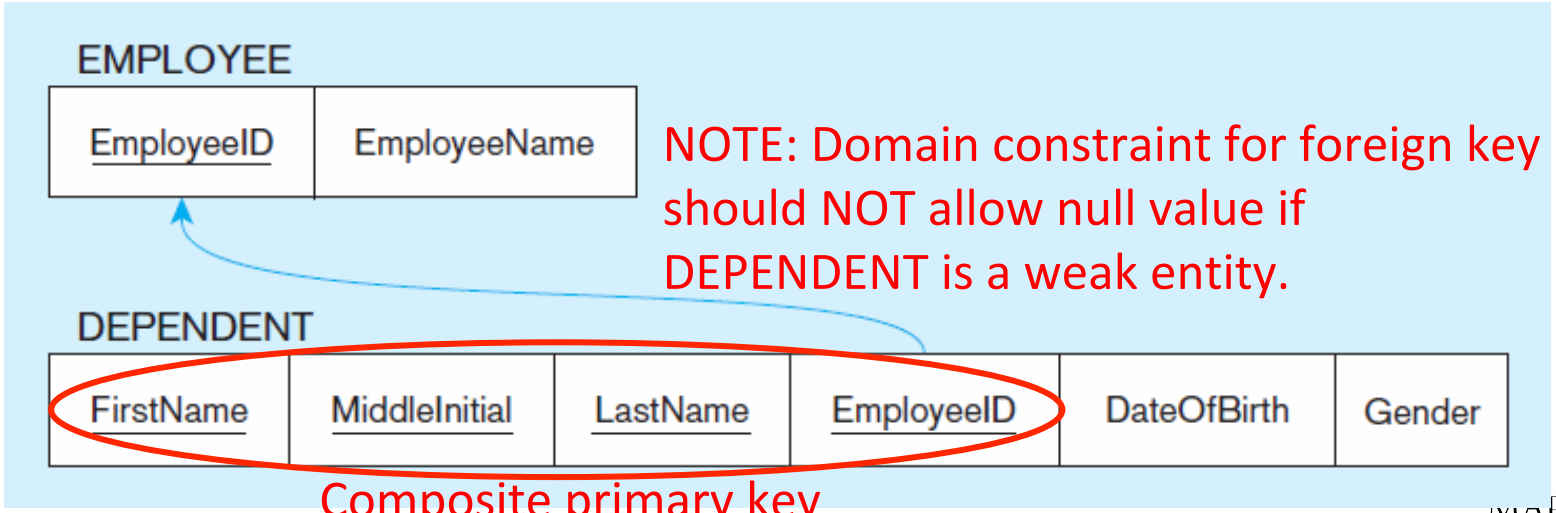

Transforming ER Diagrams into Relations Step 3

- Mapping binary relationships:
  - **One-to-One** – Primary key on mandatory side becomes a foreign key on optional side.
  - **One-to-Many** – Primary key on the one side becomes a foreign key on the many side.
  - **Many-to-Many** – Create a new relation with the primary keys of the two entities as its composite primary key.

Figure 4-14: Example of Mapping a Binary 1:1 Relationship

(a) *In charge* relationship type with attribute Date Assigned

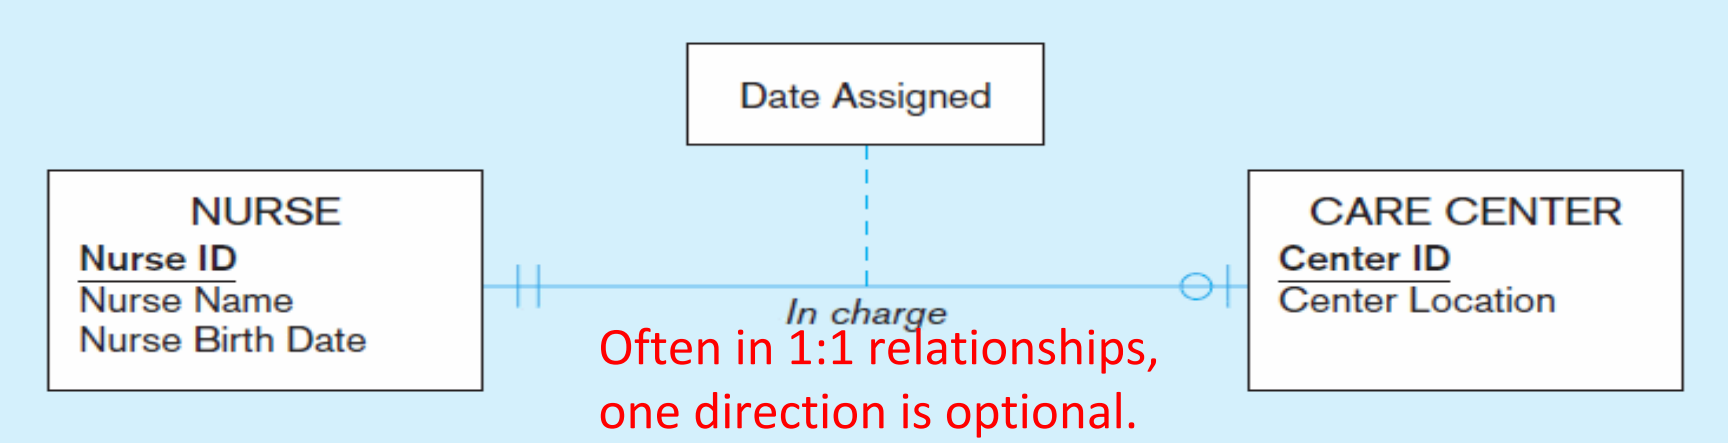

(b) Add foreign key and attribute to the optional entity type

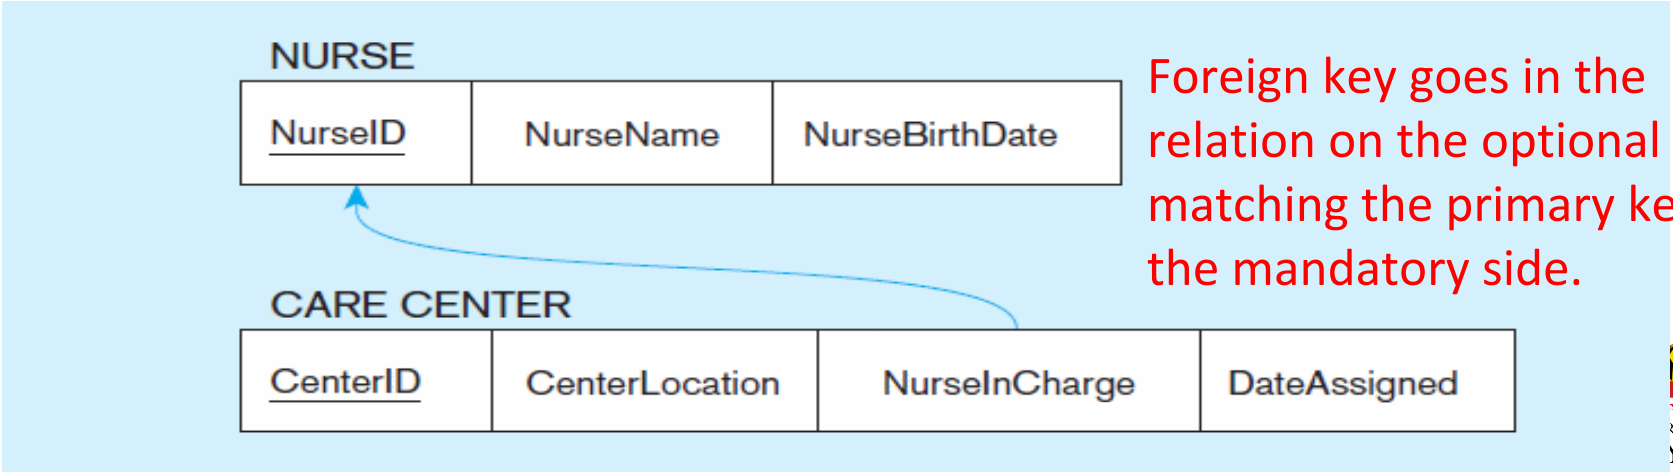

Figure 4-12: Example of Mapping a Binary 1:M Relationship

(a) *Submits* relationship type between CUSTOMERS and ORDERS

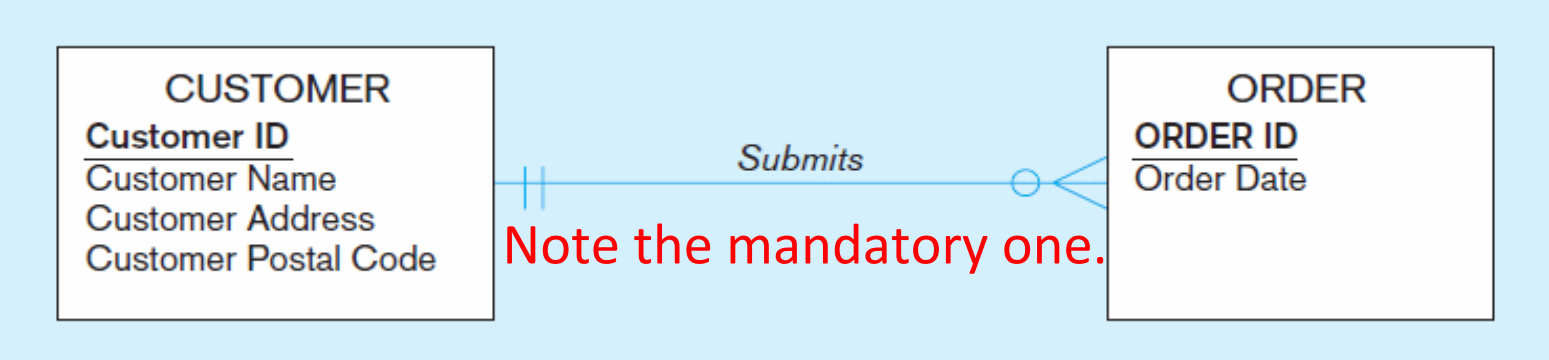

(b) Add foreign key to the optional entity type

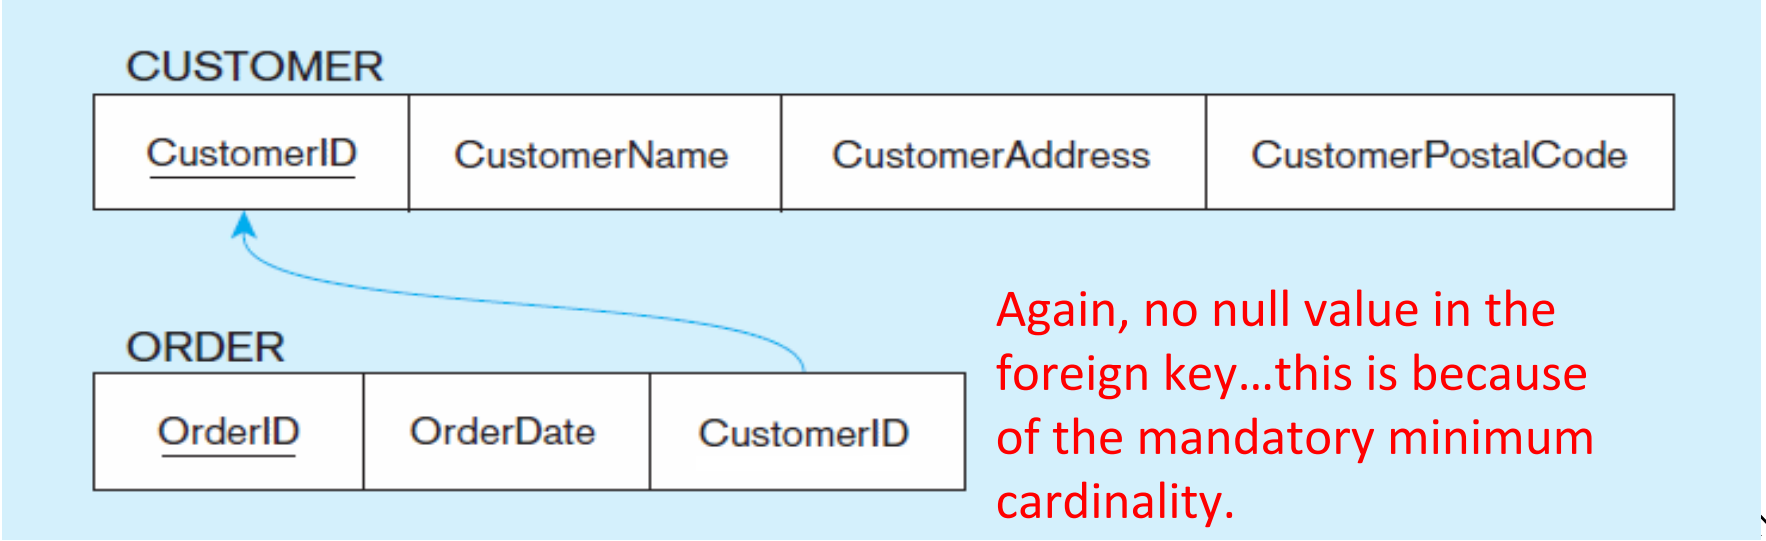

Figure 4-13: Example of Mapping a Binary M:N Relationship

(a) *Completes* relationship type between EMPLOYEE and COURSE

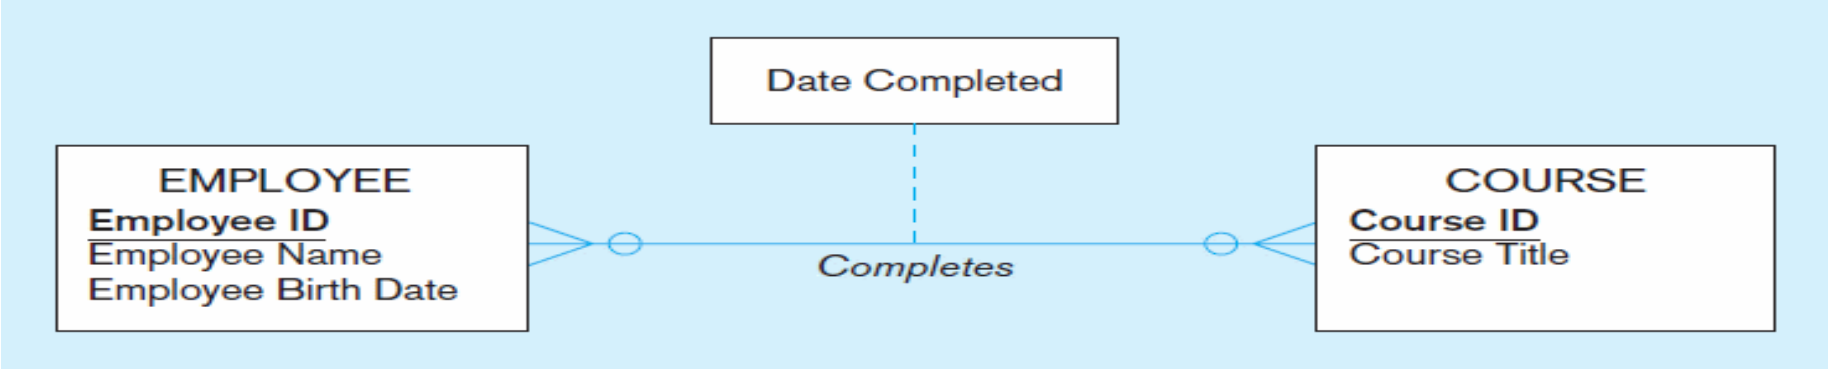

(b) Create new relation with two foreign keys and attribute

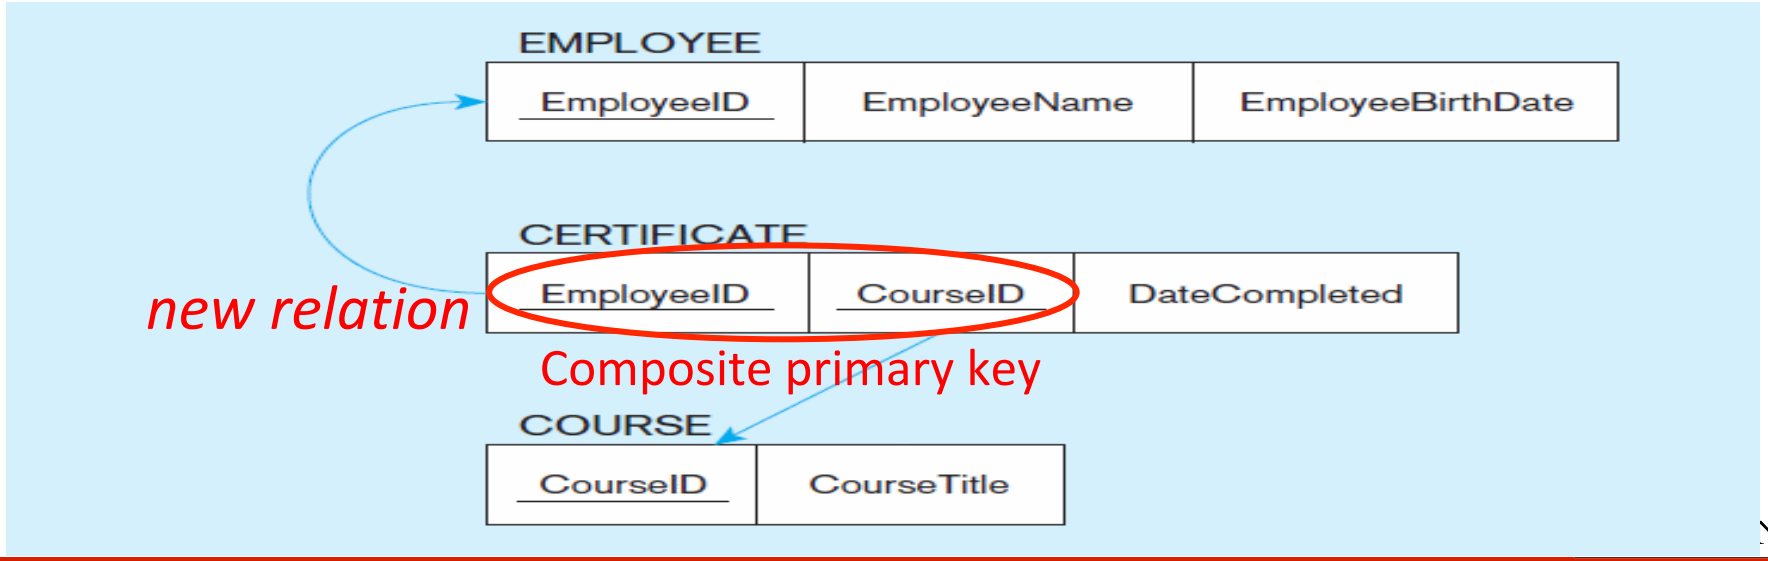

Transforming ER Diagrams into Relations Step 4

- Mapping associative entities:
  - Identifier Not Assigned:
    - Default primary key for the association relation is composed of the primary keys of the two entities (as in *:* relationship).
  - Identifier Assigned:
    - It is natural and familiar to end-users.
    - Default identifier may not be unique.

Figure 4-15: Example of Mapping a Associative Entity

(a) Two strong entities and an associative entity

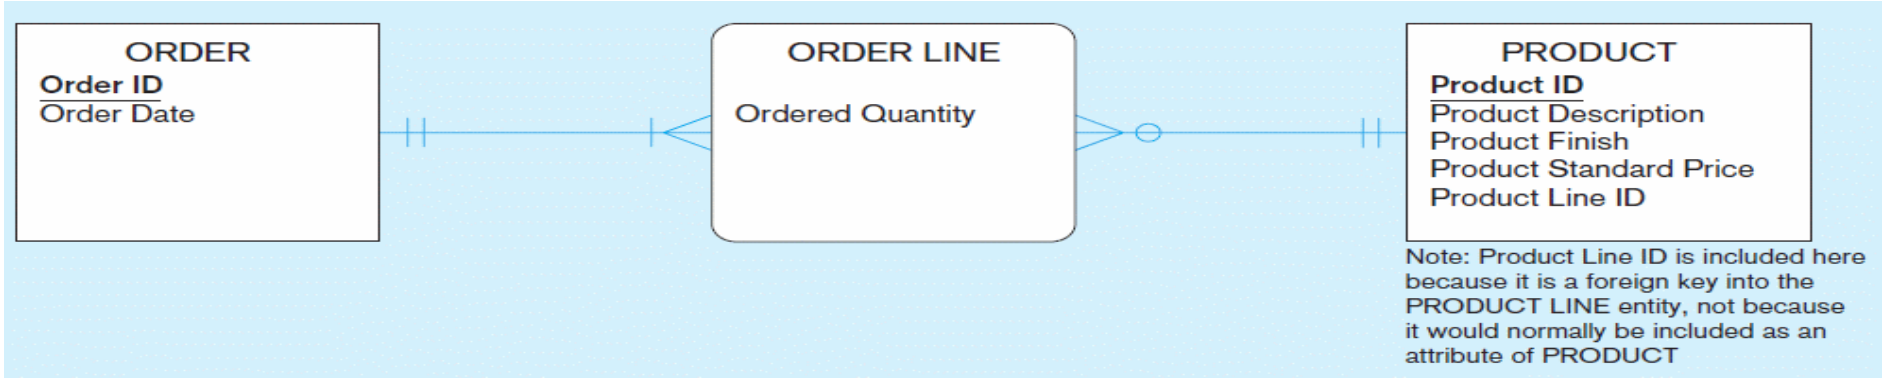

(b) Three resulting relations

Figure 4-16: Example of Mapping a Associative Entity with an Identifier

(a) Two strong entities and an associative entity

(b) Three resulting relations

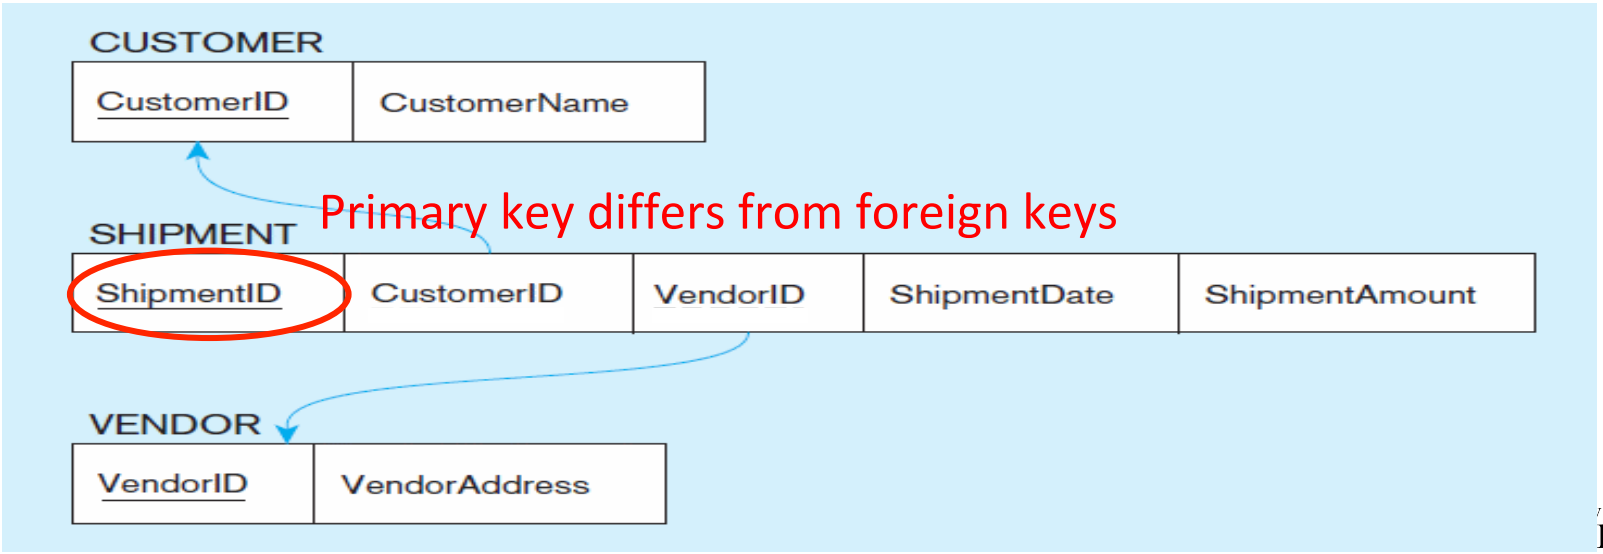

Transforming ER Diagrams into Relations Step 5

- Mapping unary relationships:
  - **One-to-Many** – Recursive foreign key in the same relation.
  - **Many-to-Many** – Two relations:
    - One for the entity type.
    - One for an associative relation in which the primary key has two attributes, both taken from the primary key of the entity.

Figure 4-17: Mapping a Unary 1:N Relationship

(a) EMPLOYEE entity with unary relationship

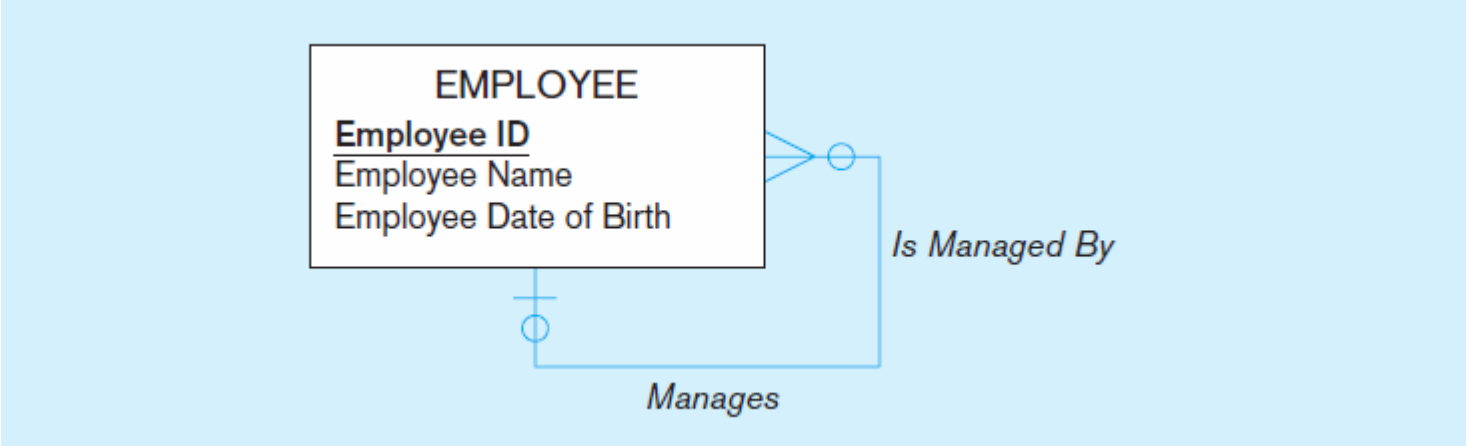

(b) EMPLOYEE relation with recursive foreign key

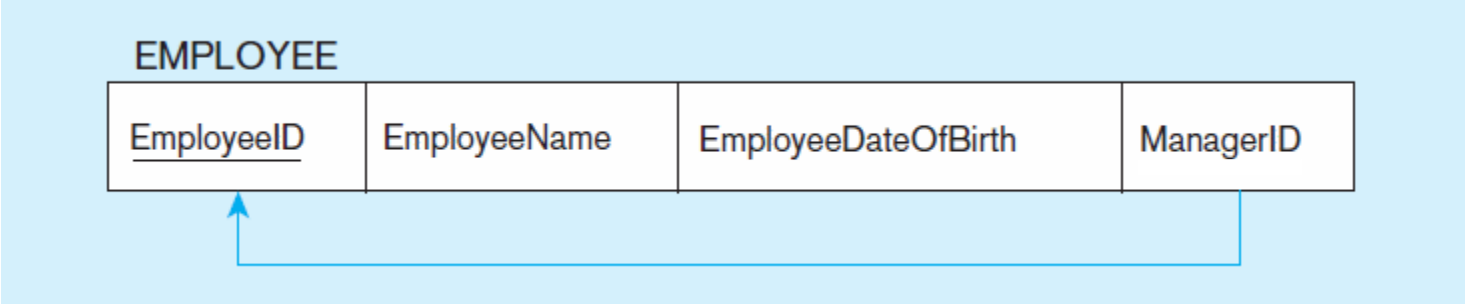

Figure 4-18: Mapping a Unary M:N Relationship

(a) Bill-of-materials relationships (M:N)

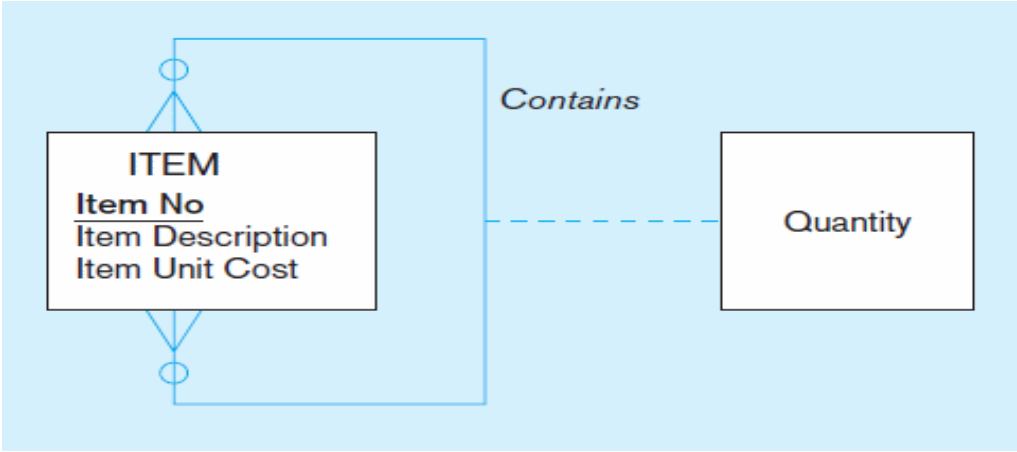

(b) ITEM and COMPONENT relations

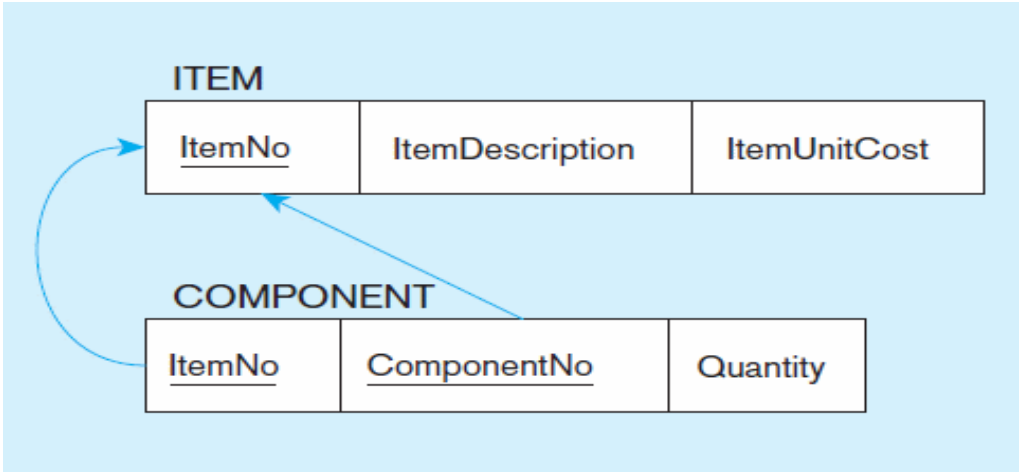

Transforming ER Diagrams into Relations Step 6

- Mapping ternary (and *n*-ary) relationships:
  - One relation for each entity and one for the associative entity.
  - Associative entity has foreign keys to each entity in the relationship.

Figure 4-19: Mapping a Ternary Relationship

(a) PATIENT TREATMENT ternary relationship with associative entity

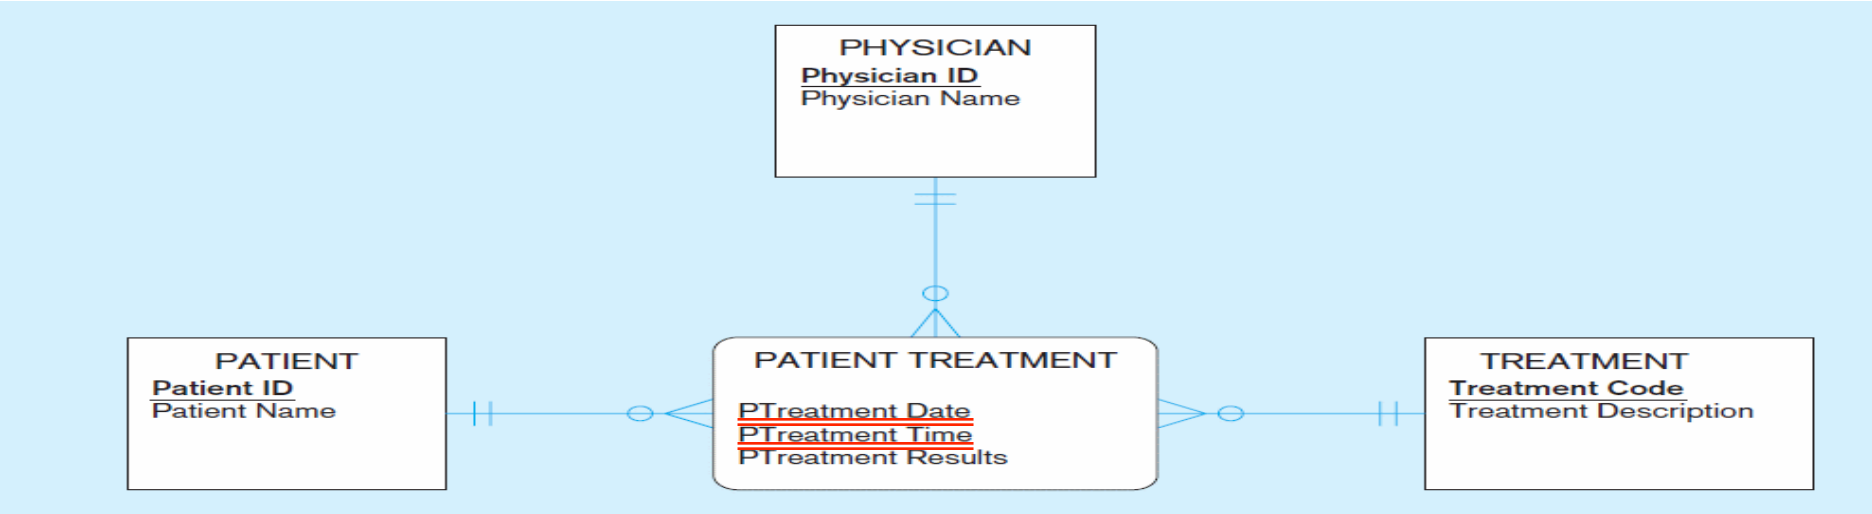

(b) PATIENT TREATMENT relation with three foreign keys

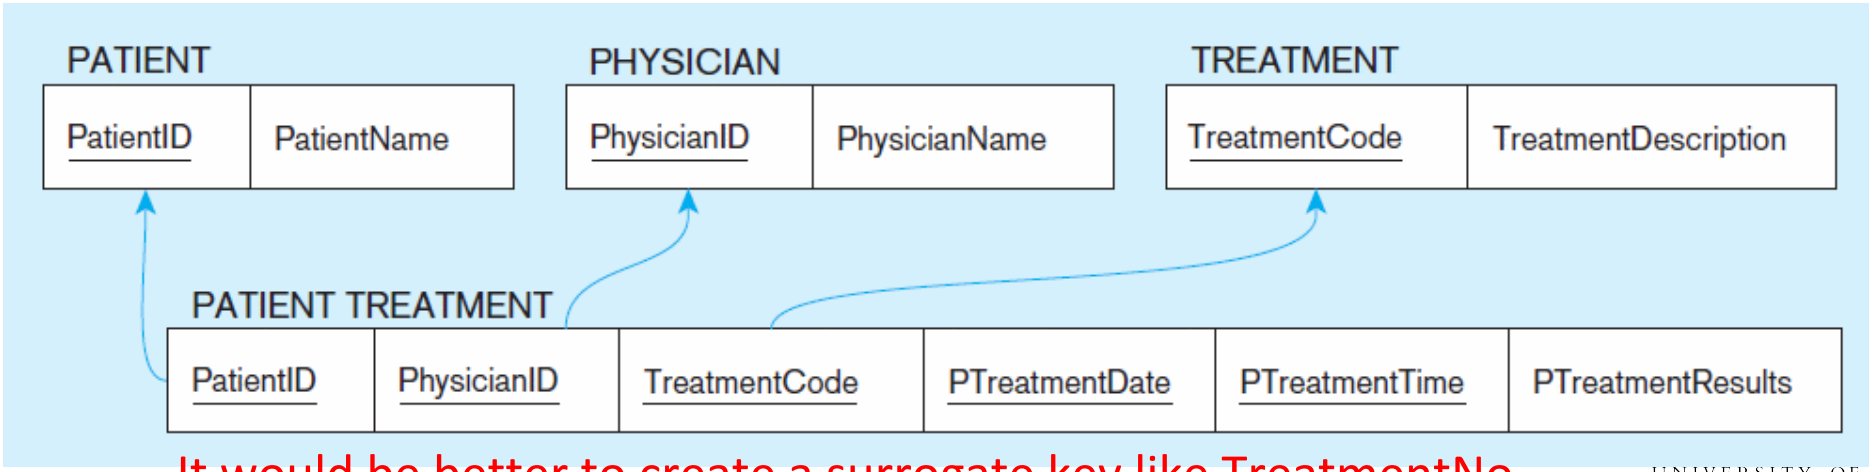

It would be better to create a surrogate key like TreatmentNo.

Figure 4-19: Mapping a Ternary Relationship

(a) PATIENT TREATMENT ternary relationship with associative entity

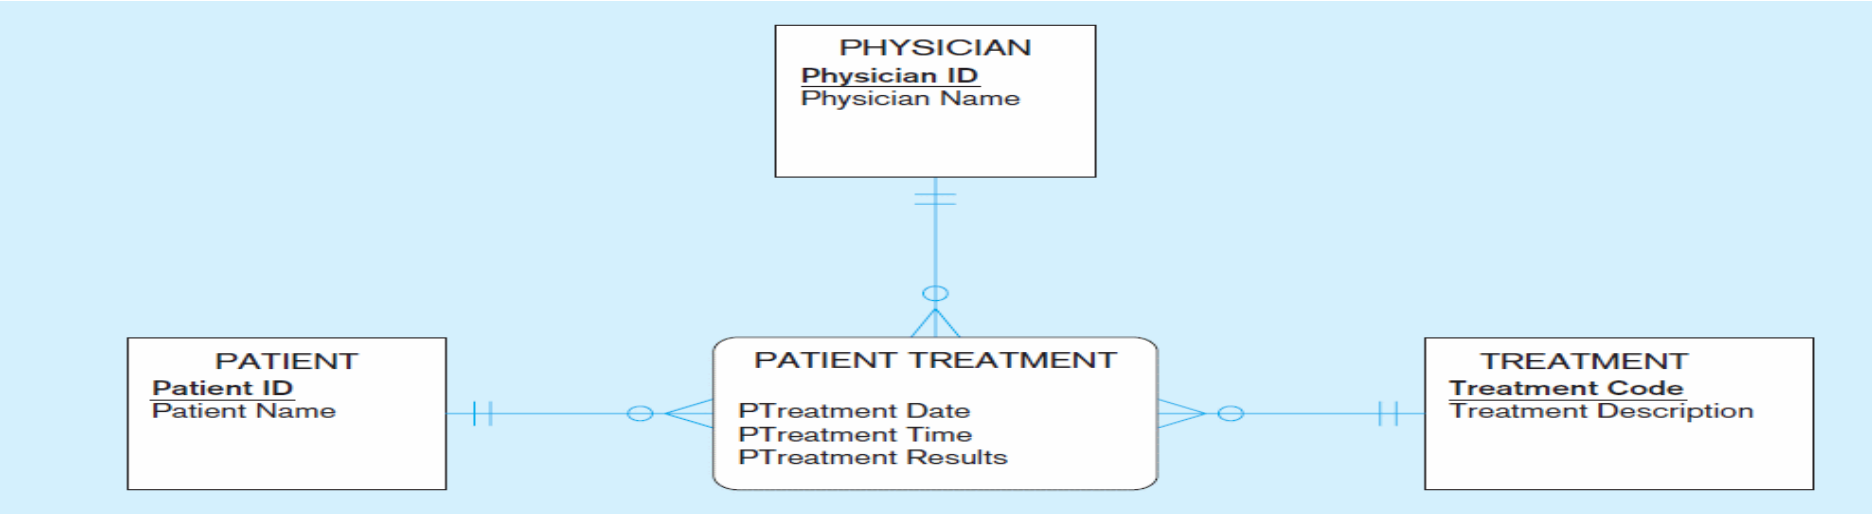

(b) PATIENT TREATMENT relation with three foreign keys

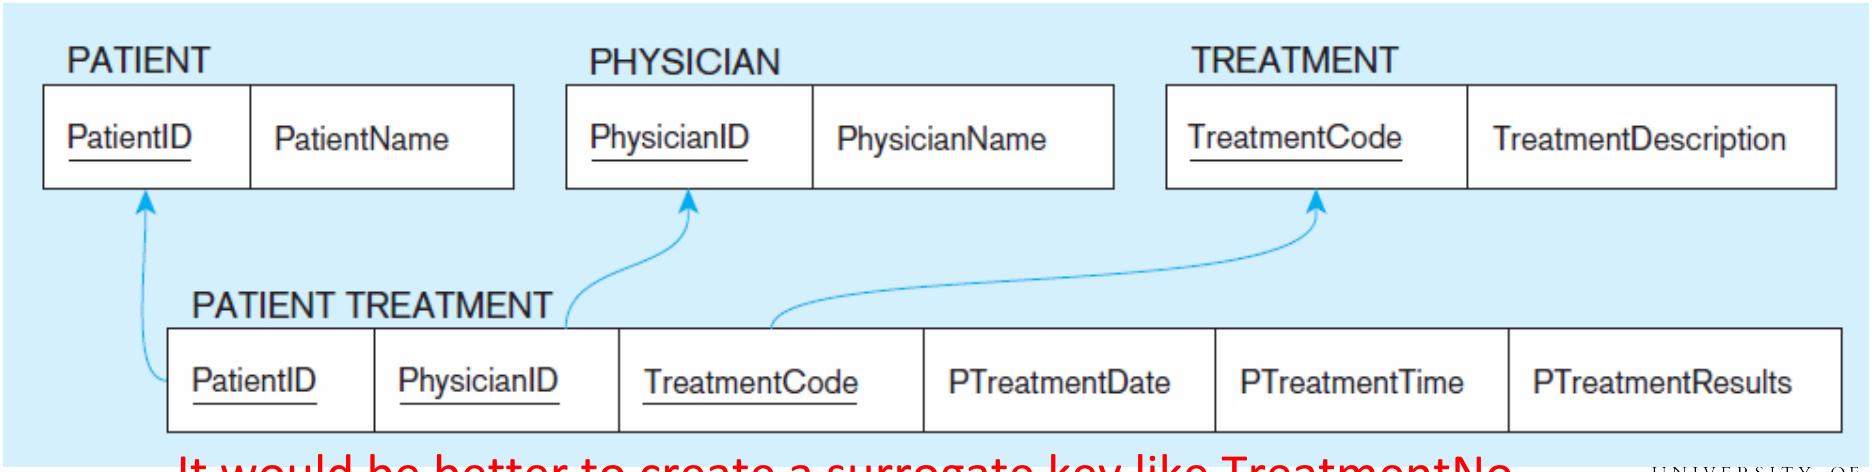

It would be better to create a surrogate key like TreatmentNo.

Transforming ER Diagrams into Relations Step 7

- Mapping supertype/subtype relationships:
  - One relation for supertype and for each subtype.
  - Supertype attributes (including identifier and subtype discriminator) go into supertype relation.
  - Subtype attributes go into each subtype; primary key of supertype relation also becomes primary key of subtype relation.
  - 1:1 relationship established between supertype and each subtype, with supertype as primary table.

Figure 4-20: Supertype/Subtype Relationships

(a) EMPLOYEE with three subtypes

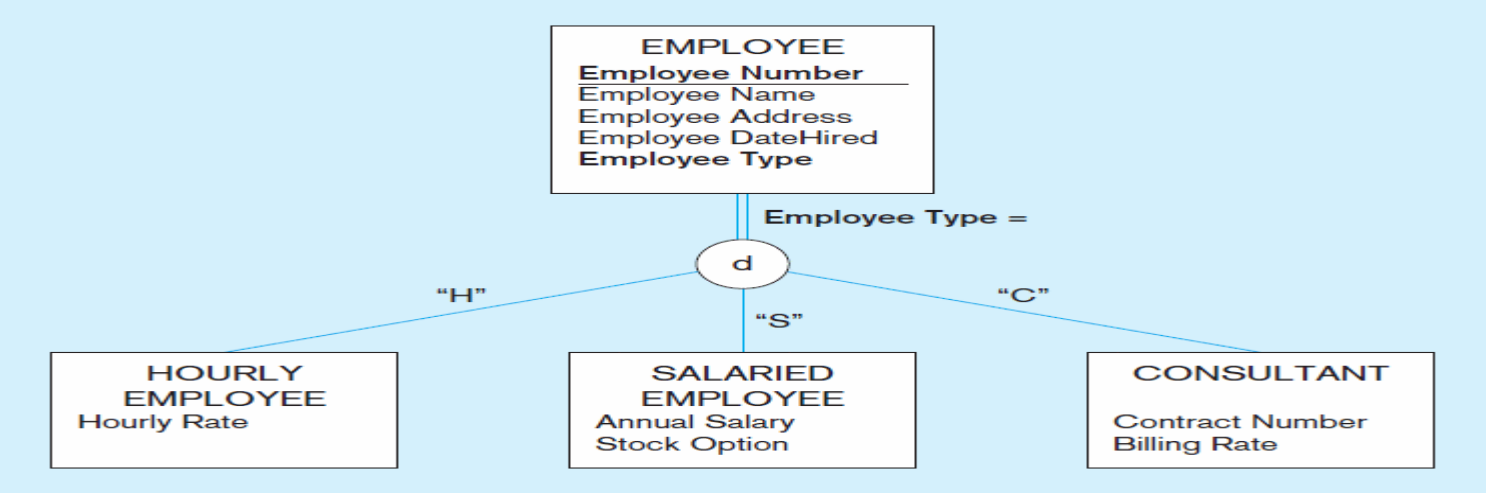

(b) These are implemented as one-to-one relationships

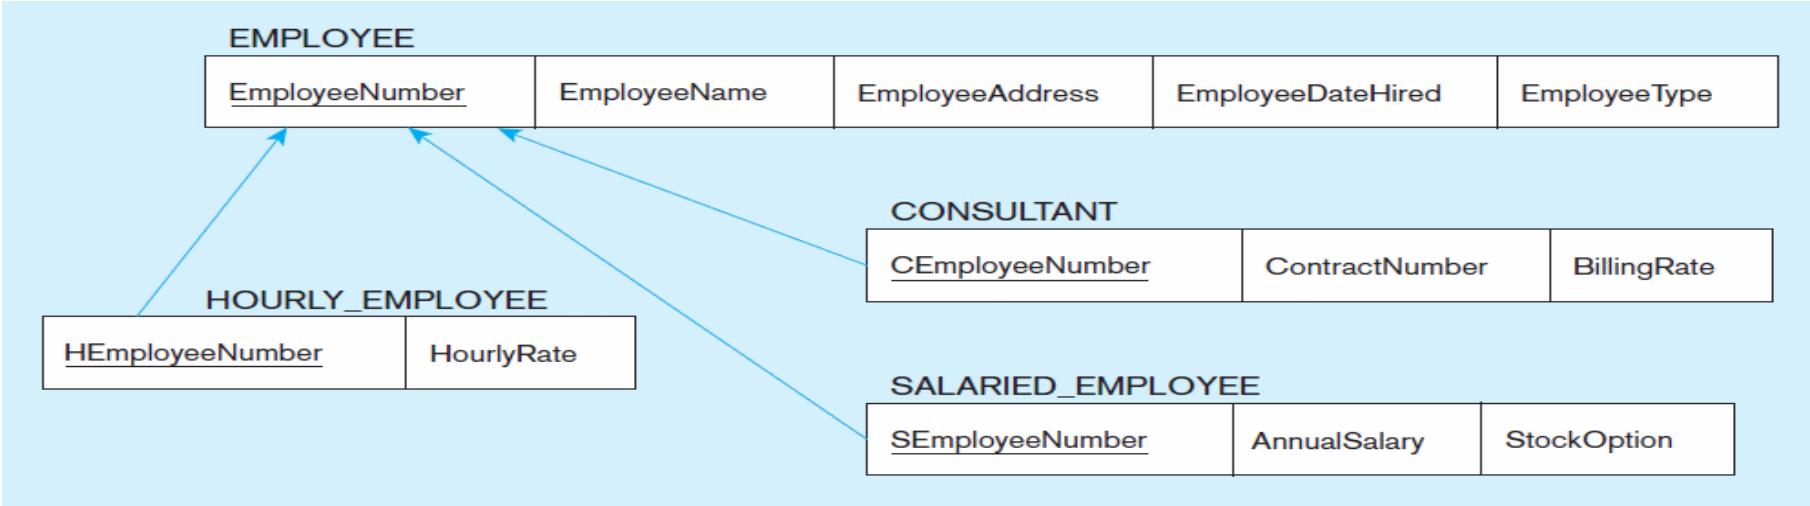

Data Normalization

- Primarily a tool to validate and improve a logical design so that it satisfies certain constraints that avoid unnecessary duplication of data.
- The process of decomposing relations with anomalies to produce smaller, well-structured relations.

Well-Structured Relations

- A relation that contains minimal data redundancy and allows users to insert, delete, and update rows without causing data inconsistencies.
- Goal is to avoid anomalies:
  - **Insertion Anomaly** – adding new rows forces user to create duplicate data.
  - **Deletion Anomaly** – deleting rows may cause a loss of data that would be needed for other future rows.
  - **Modification Anomaly** – changing data in a row forces changes to other rows because of duplication.
- **General rule of thumb**:
  - A table should not pertain to more than one entity type.

Figure 4-2: EMPLOYEE2

EMPLOYEE2

EmpID	Name	DeptName	Salary	CourseTitle	DateCompleted
100	Margaret Simpson	Marketing	48,000	SPSS	6/19/201X
				Surveys	10/7/201X
140	Alan Beeton	Accounting	52,000	Tax Acc	12/8/201X
110	Chris Lucero	Info Systems	43,000	Visual Basic	1/12/201X
				C++	4/22/201X
15 0	Susan Martin	Marketing	42,000	SPSS	6/19/201X
				Java	8/12/201X

- **Question** – Is this a relation?
- **Answer** – Yes, unique rows and no multivalued attributes.

- **Question** – What's the primary key?
- **Answer** – Composite: EmpID, CourseTitle?

Anomalies in This Table

- **Insertion** – Cannot enter a new employee without having the employee take a class (or at least empty fields of class information).
- **Deletion** – If we remove employee 140, we lose information about the existence of a Tax Acc class.
- **Modification** – Giving a salary increase to employee 100 forces us to update multiple records.

- Why do these anomalies exist?
- Because there are two themes (entity types) in this one relation. This results in data duplication and an unnecessary dependency between the entities.

UNIVERSITY OF MARYLAND

Figure 4.22: Steps in Normalization

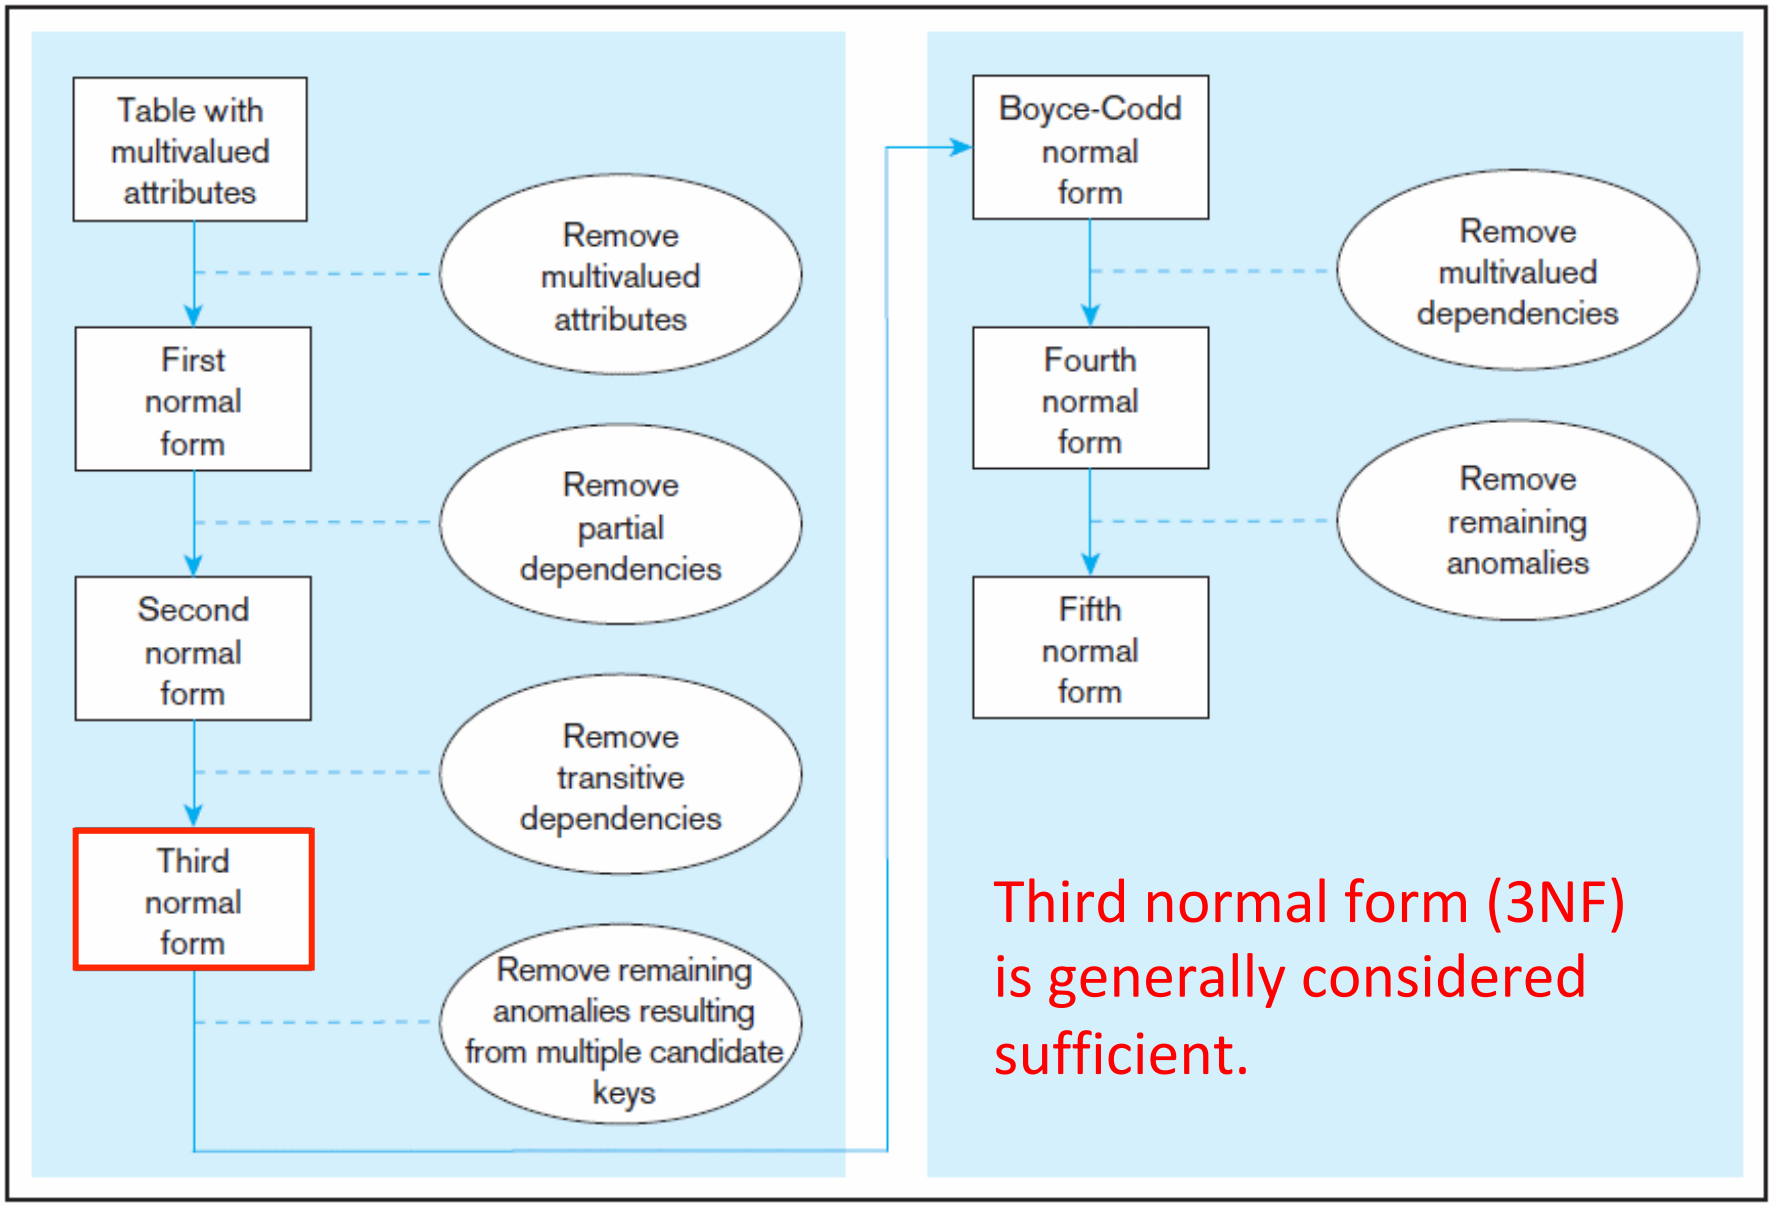

Functional Dependencies and Keys

- Functional Dependency:
  - The value of one attribute (the determinant) determines the value of another attribute.

- Candidate Key:
  - A minimal unique identifier. One of the candidate keys will become the primary key.
  - E.g., perhaps there is both credit card number and SS# in a table… in this case both are candidate keys.
  - Each non-key field is functionally dependent on every candidate key.

UNIVERSITY OF MARYLAND

First Normal Form (1NF)

- No multivalued attributes.
- Every attribute value is atomic.
- Fig. 4-25 is not in 1NF (multivalued attributes).
  - i.e. It is not a relation.
- Fig. 4-26 is in 1NF.
- All relations are in 1NF.

Figure 4-25: Table with Multivalued Attributes, not in 1NF

OrderID	Order Date	Customer ID	Customer Name	Customer Address	ProductID	Product Description	Product Finish	Product StandardPrice	Ordered Quantity
1006	10/24/2010	2	Value Furniture	Plano, TX	7	Dining Table	Natural Ash	800.00	2
					5	Writer's Desk	Cherry	325.00	2
					4	Entertainment Center	Natural Maple	650.00	1
1007	10/25/2010	6	Furniture Gallery	Boulder, CO	11	4–Dr Dresser	Oak	500.00	4
					4	Entertainment Center	Natural Maple	650.00	3

FIGURE 4-25 INVOICE data (Pine Valley Furniture Company)

NOTE: This is NOT a relation.

Figure 4-26: Table with No Multivalued Attributes and Unique Rows, in 1NF

OrderID	Order Date	Customer ID	Customer Name	Customer Address	ProductID	Product Description	Product Finish	Product StandardPrice	Ordered Quantity
1006	10/24/2010	2	Value Furniture	Plano, TX	7	Dining Table	Natural Ash	800.00	2
1006	10/24/2010	2	Value Furniture	Plano, TX	5	Writer's Desk	Cherry	325.00	2
1006	10/24/2010	2	Value Furniture	Plano, TX	4	Entertainment Center	Natural Maple	650.00	1
1007	10/25/2010	6	Furniture Gallery	Boulder, CO	11	4–Dr Dresser	Oak	500.00	4
1007	10/25/2010	6	Furniture Gallery	Boulder, CO	4	Entertainment Center	Natural Maple	650.00	3

FIGURE 4-26 INVOICE relation (1NF) (Pine Valley Furniture Company)

NOTE: This is a relation, but not a well-structured one.

Anomalies in This Table

- **Insertion** – If new product is ordered for order 1007 of existing customer, customer data must be re-entered, causing duplication.
- **Deletion** – If we delete the Dining Table from Order 1006, we lose information concerning this item's finish and price.
- **Modification** – Changing the price of product ID 4 requires update in multiple records.

- Why do these anomalies exist?
- Because there are multiple themes (entity types) in one relation. This results in duplication and an unnecessary dependency between the entities.

UNIVERSITY OF MARYLAND

Second Normal Form (2NF)

- 1NF PLUS every non-key attribute is fully functionally dependent on the ENTIRE primary key.
- Every non-key attribute must be defined by the entire key, not by only part of the key.
- No partial functional dependencies.

Figure 4-27: Functional Dependency Diagram

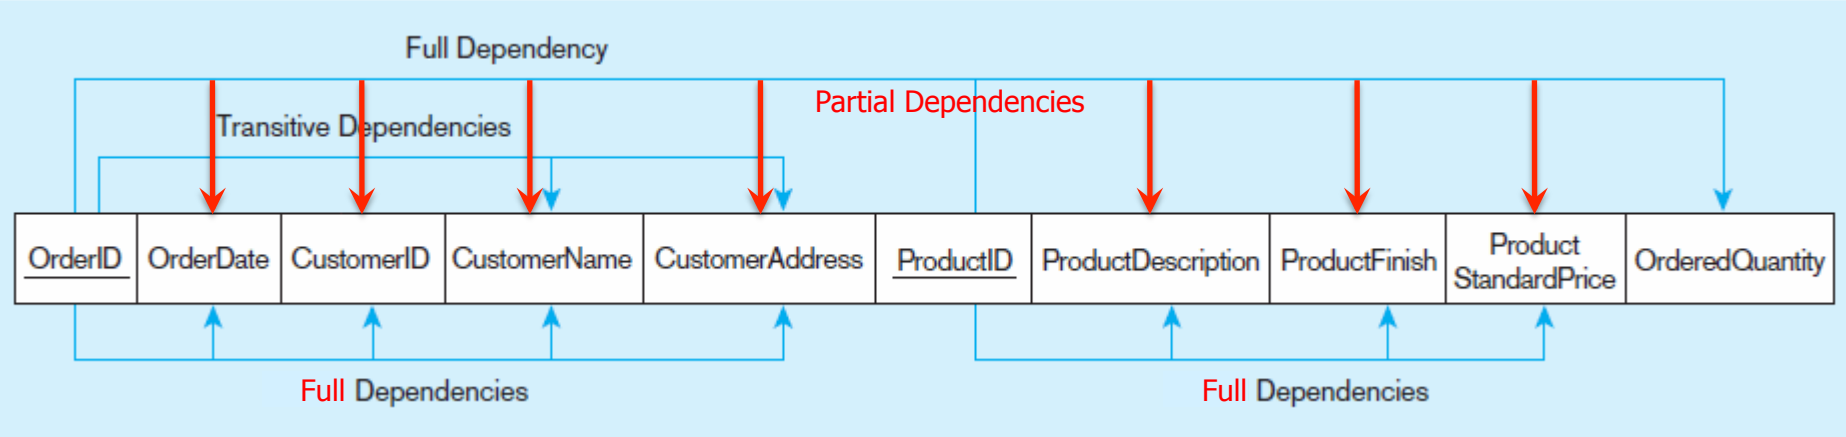

- OrderID → OrderDate, CustomerID, CustomerName, CustomerAddress
- CustomerID → CustomerName, CustomerAddress
- ProductID → ProductDescription, ProductFinish, ProductStandardPrice
- OrderID, ProductID → OrderQuantity, OrderDate, CustomerID, CustomerName, CustomerAddress, ProductDescription, ProductFinish, ProductStandardPrice

Therefore, NOT in 2NF

Figure 4-28: Removing Partial Dependencies

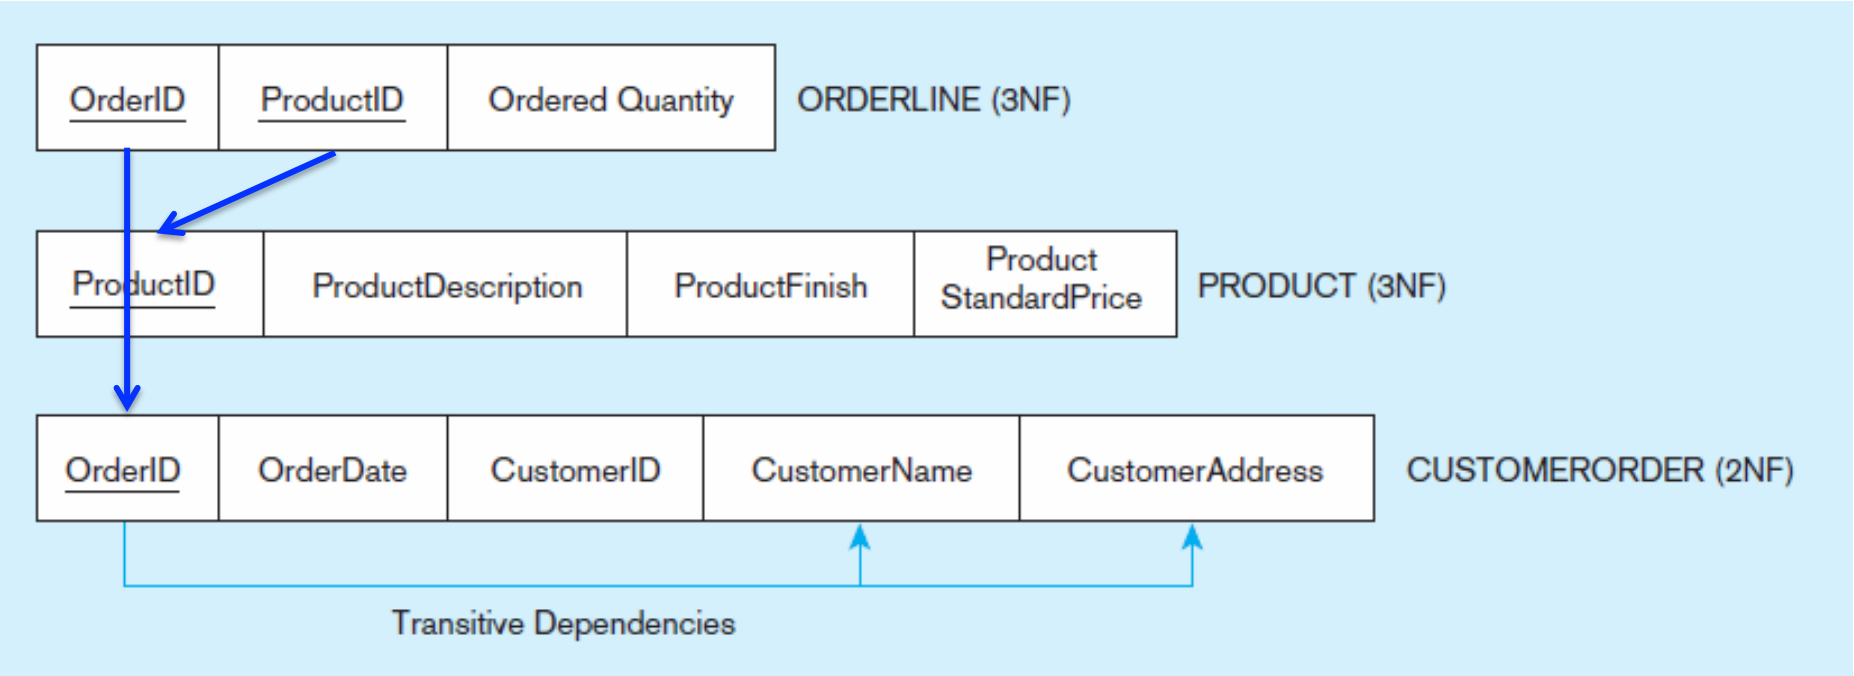

- OrderID, ProductID → OrderQuantity
- ProductID → ProductDescription, ProductFinish, ProductStandardPrice
- OrderID → OrderDate, CustomerID, CustomerName, CustomerAddress But there are still transitive dependencies.
- CustomerID → CustomerName, CustomerAddress

Third Normal Form (3NF)

- 2NF PLUS no transitive dependencies (functional dependencies on non-primary-key attributes).
- Note: This is called transitive, because the primary key is a determinant for another attribute, which in turn is a determinant for a third.
- Solution: Non-key determinant with transitive dependencies go into a new table; non-key determinant becomes primary key in the new table and stays as foreign key in the old table.

Figure 4-29: Removing Transitive Dependencies

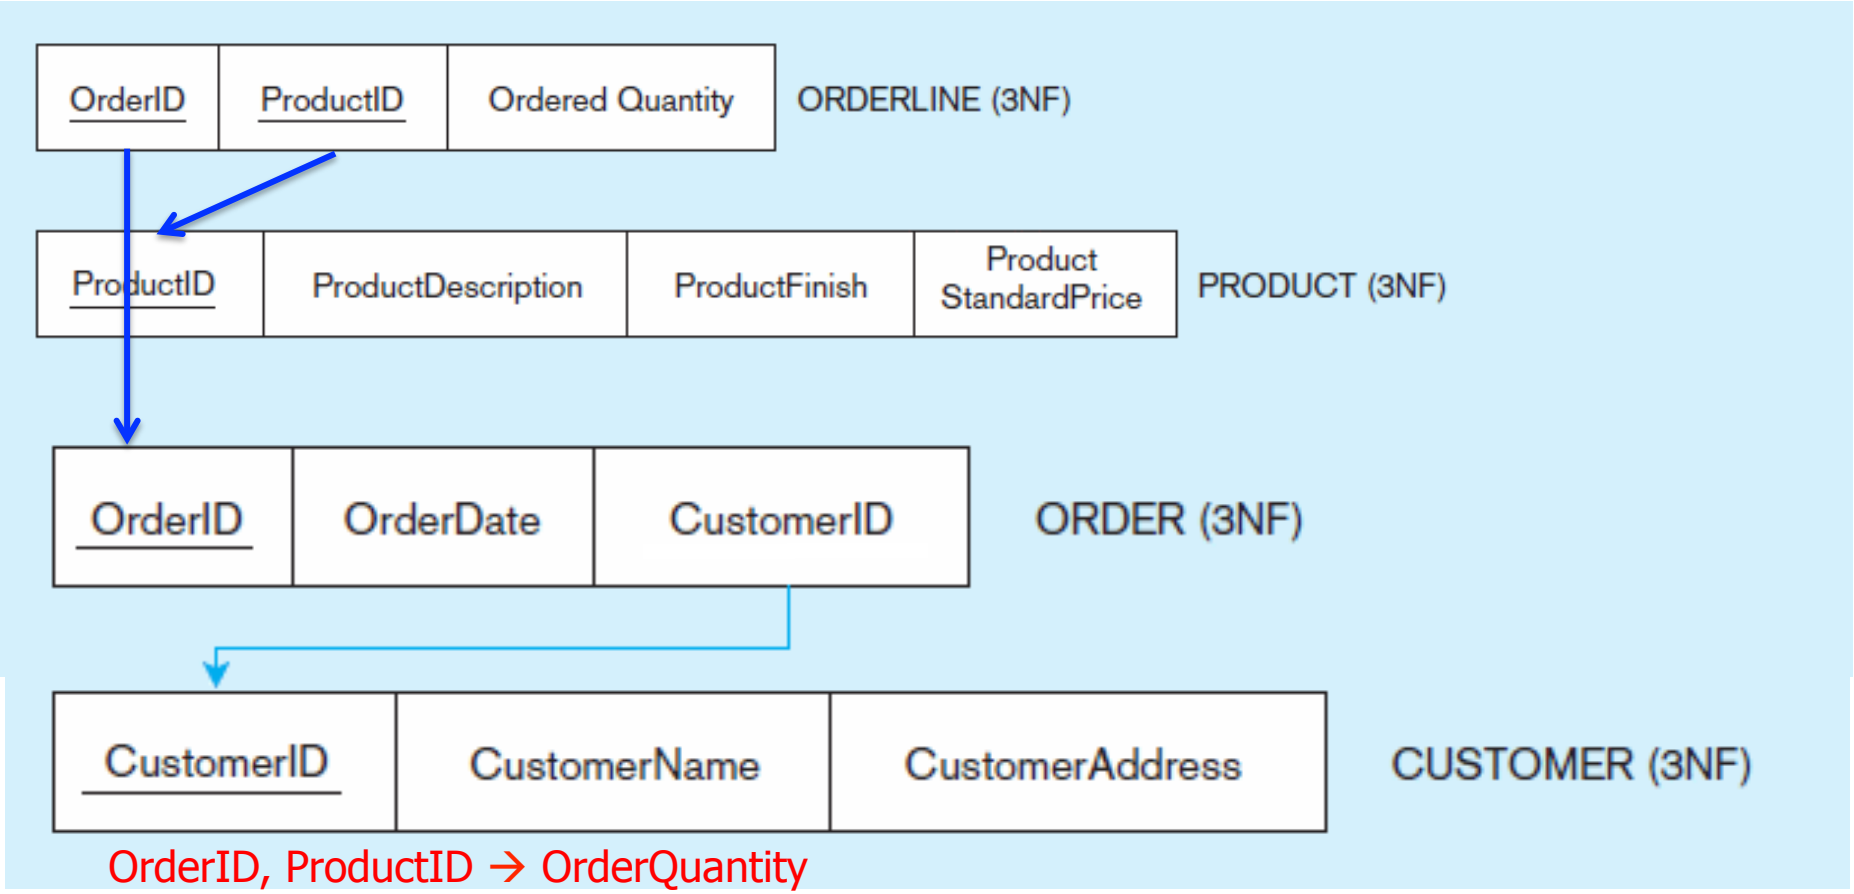

OrderID, ProductID → OrderQuantity

ProductID → ProductDescription, ProductFinish, ProductStandardPrice

OrderID → OrderDate, CustomerID

CustomerID → CustomerName, CustomerAddress

Merging Relations

- **View Integration** – Combining entities from multiple ER models into common relations.
- Issues to watch out for when merging entities from different ER models:
  - **Synonyms** – Two or more attributes with different names but same meaning.
  - **Homonyms** – Attributes with same name but different meanings.
  - **Transitive Dependencies** – Even if relations are in 3NF prior to merging, they may not be after merging.
  - **Supertype/Subtype Relationships** – May be hidden prior to merging.

Enterprise Keys

- Primary keys that are unique in the whole database, not just within a single relation.
- Corresponds with the concept of an object ID in object-oriented systems.

Figure 4-31: Enterprise Keys

(a) Relations with enterprise key

OBJECT (OID, ObjectType)
EMPLOYEE (OID, EmpID, EmpName, DeptName, Salary)
CUSTOMER (OID, CustID, CustName, Address)

(b) Sample data with enterprise key

OBJECT

OID	ObjectType
1	EMPLOYEE
2	CUSTOMER
3	CUSTOMER
4	EMPLOYEE
5	EMPLOYEE
6	CUSTOMER
7	CUSTOMER

EMPLOYEE

OID	EmpID	EmpName	DeptName	Salary
1	100	Jennings, Fred	Marketing	50000
4	101	Hopkins, Dan	Purchasing	45000
5	102	Huber, Ike	Accounting	45000

CUSTOMER

OID	CustID	CustName	Address
2	100	Fred's Warehouse	Greensboro, NC
3	101	Bargain Bonanza	Moscow, ID
6	102	Jasper's	Tallahassee, FL
7	103	Desks 'R Us	Kettering, OH