

Database Management Systems

Chapter 6: Introduction to SQL



UNIVERSITY OF
MARYLAND

ROBERT H. SMITH
SCHOOL OF BUSINESS

DR. ADAM LEE

Objectives

- Define terms.
- Interpret history and role of SQL.
- Define a database using SQL data definition language.
- Write single table queries using SQL.
- Establish referential integrity using SQL.



UNIVERSITY OF
MARYLAND

SQL Overview

- **Structured Query Language**
- The standard for **relational database management systems** (RDBMS).
- **RDBMS**: A database management system that manages data as a collection of tables in which all relationships are represented by common values in related tables.



UNIVERSITY OF
MARYLAND

ROBERT H. SMITH
SCHOOL OF BUSINESS

History of SQL

- **1970** – E. F. Codd developed relational database concept.
- **1974-1979** – System R with Sequel (later SQL) created at IBM Research Lab.
- **1979** – Oracle markets first relational DB with SQL.
- **1981** – SQL/DS first available RDBMS system on DOS/VSE.
- Others followed: INGRES (**1981**), IDM (**1982**), DG/SGL (**1984**), Sybase (**1986**).
- **1986** – ANSI SQL standard released; **1989, 1992, 1999, 2003, 2006, 2008, 2011** – Major ANSI standard updates.
- **Current** – SQL is supported by most major database vendors.



UNIVERSITY OF
MARYLAND

ROBERT H. SMITH
SCHOOL OF BUSINESS

Purpose of SQL Standard

- Specify syntax/semantics for data definition and manipulation.
- Define data structures and basic operations.
- Enable portability of database definition and application modules.
- Specify minimal (level 1) and complete (level 2) standards.
- Allow for later growth/enhancement to standard (referential integrity, transaction management, user-defined functions, extended join operations, national character sets).



UNIVERSITY OF
MARYLAND

ROBERT H. SMITH
SCHOOL OF BUSINESS

Benefits of a Standardized Relational Language

- Reduced training costs.
- Productivity.
- Application portability.
- Application longevity.
- Reduced dependence on a single vendor.
- Cross-system communication.



UNIVERSITY OF
MARYLAND

ROBERT H. SMITH
SCHOOL OF BUSINESS

SQL Environment

- **Catalog:**

- A set of schemas that constitute the description of a database.

- **Schema:**

- The structure that contains descriptions of objects created by a user (base tables, views, constraints).

- **Data Definition Language (DDL):**

- Commands that define a database, including creating, altering, and dropping tables and establishing constraints.

- **Data Manipulation Language (DML):**

- Commands that maintain and query a database.

- **Data Control Language (DCL):**

- Commands that control a database, including administering privileges and committing data.



UNIVERSITY OF
MARYLAND

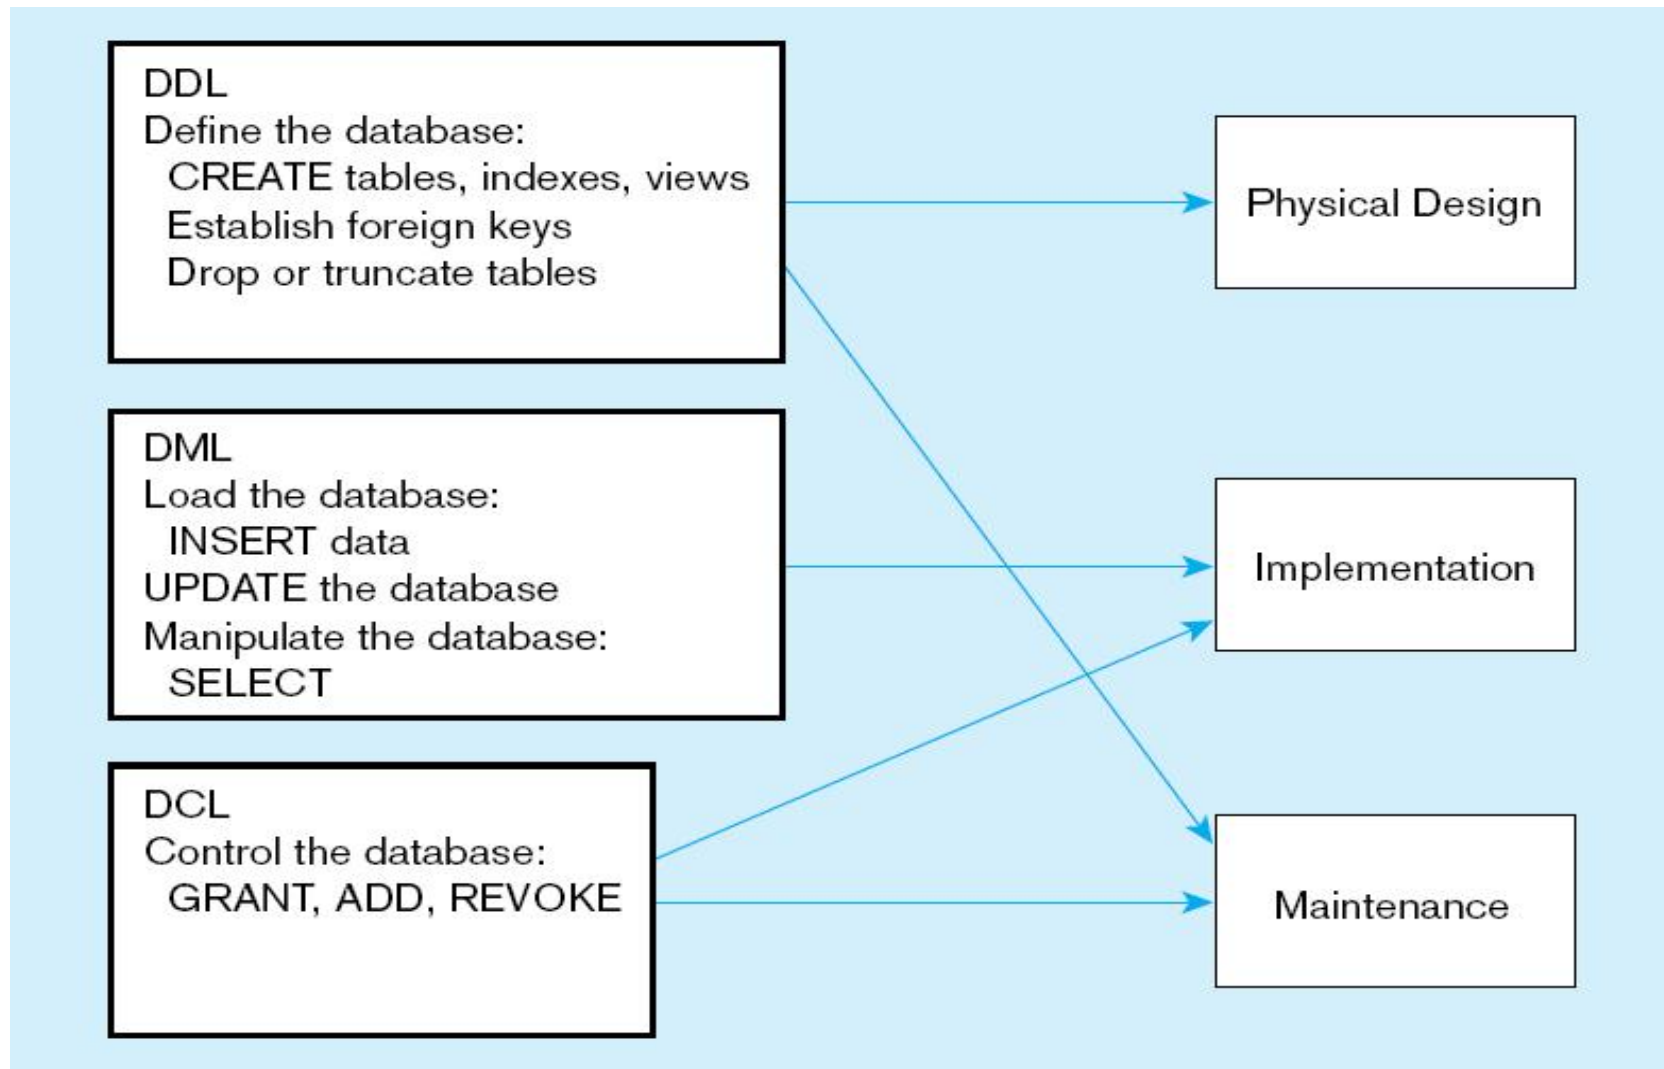
ROBERT H. SMITH
SCHOOL OF BUSINESS

Table 6-2: SQL Data Types

TABLE 6-2 Sample SQL Data Types

String	CHARACTER (CHAR)	Stores string values containing any characters in a character set. CHAR is defined to be a fixed length.
	CHARACTER VARYING (VARCHAR or VARCHAR2)	Stores string values containing any characters in a character set but of definable variable length.
	BINARY LARGE OBJECT (BLOB)	Stores binary string values in hexadecimal format. BLOB is defined to be a variable length. (Oracle also has CLOB and NCLOB, as well as BFILE for storing unstructured data outside the database.)
Number	NUMERIC	Stores exact numbers with a defined precision and scale.
	INTEGER (INT)	Stores exact numbers with a predefined precision and scale of zero.
Temporal	TIMESTAMP TIMESTAMP WITH LOCAL TIME ZONE	Stores a moment an event occurs, using a definable fraction-of-a-second precision. Value adjusted to the user's session time zone (available in Oracle and MySQL)
Boolean	BOOLEAN	Stores truth values: TRUE, FALSE, or UNKNOWN.

Figure 6-4: DDL, DML, DCL, and the Database Development Process



SQL Database Definition

- Data Definition Language (DDL)
- Major CREATE statements:
 - **CREATE SCHEMA** – defines a portion of the database owned by a particular user.
 - **CREATE TABLE** – defines a new table and its columns.
 - **CREATE VIEW** – defines a logical table from one or more tables or views.
- Other CREATE statements:
 - **CREATE CHARACTER SET**
 - **CREATE ASSERTION**
 - **CREATE DOMAIN**



UNIVERSITY OF
MARYLAND

ROBERT H. SMITH
SCHOOL OF BUSINESS

Steps in Table Creation

- Identify **data types** for attributes.
- Identify columns that can and cannot be **null**.
- Identify columns that must be **unique** (candidate keys).
- Identify primary key–**foreign key** mates.
- Determine **default** values.
- Identify **constraints** on columns (domain specifications).
- **Create** the table and associated indexes.



UNIVERSITY OF
MARYLAND

ROBERT H. SMITH
SCHOOL OF BUSINESS

Figure 6-5: General Syntax for CREATE TABLE statement used in DDL

```
CREATE TABLE tablename  
( {column definition [table constraint] } . . .  
[ON COMMIT {DELETE | PRESERVE} ROWS] );
```

where *column definition* ::=

column_name

{*domain name* | *datatype* [(*size*)] }

[*column_constraint_clause* . . .]

[*default value*]

[*collate clause*]

and *table constraint* ::=

[CONSTRAINT *constraint_name*]

Constraint_type [*constraint_attributes*]



UNIVERSITY OF
MARYLAND

Figure 1-3: Project Level Data Models

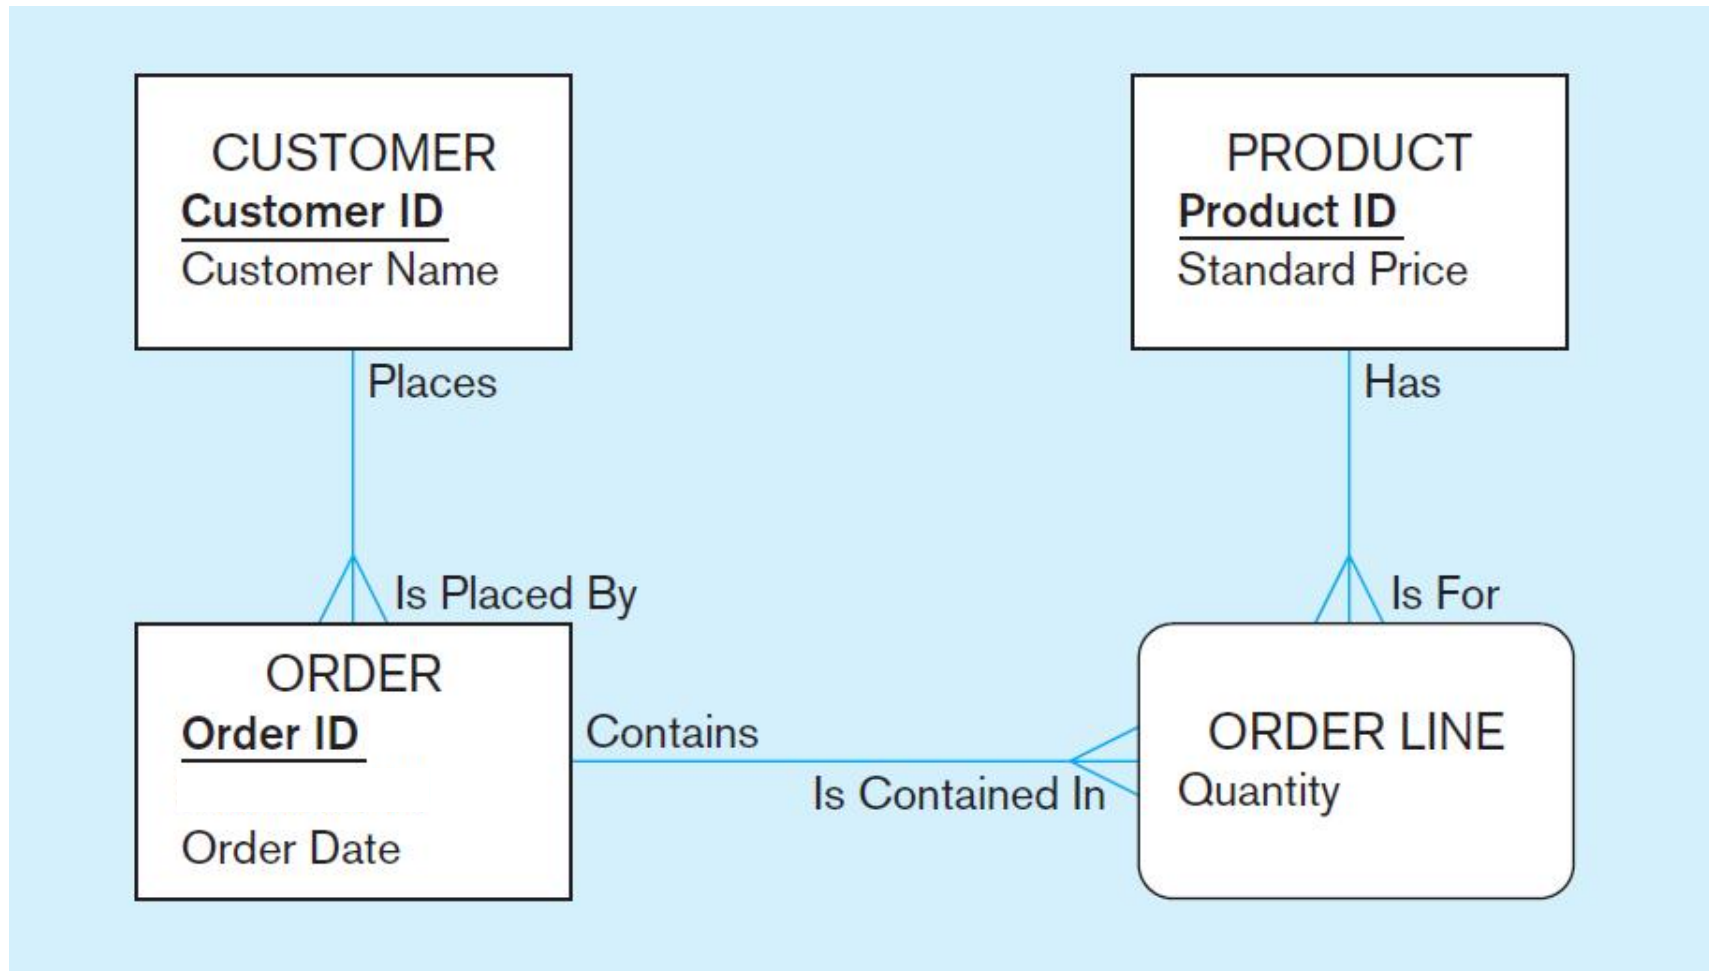


Figure 6-6: SQL Database Definition Commands (Pine Valley Furniture Company)

```
CREATE TABLE Customer_T
    (CustomerID          NUMBER(11,0)    NOT NULL,
     CustomerName        VARCHAR2(25)    NOT NULL,
     CustomerAddress     VARCHAR2(30),
     CustomerCity        VARCHAR2(20),
     CustomerState       CHAR(2),
     CustomerPostalCode  VARCHAR2(9),
 CONSTRAINT Customer_PK PRIMARY KEY (CustomerID));

CREATE TABLE Order_T
    (OrderID             NUMBER(11,0)    NOT NULL,
     OrderDate           DATE DEFAULT SYSDATE,
     CustomerID          NUMBER(11,0),
 CONSTRAINT Order_PK PRIMARY KEY (OrderID),
 CONSTRAINT Order_FK FOREIGN KEY (CustomerID) REFERENCES Customer_T(CustomerID));

CREATE TABLE Product_T
    (ProductID           NUMBER(11,0)    NOT NULL,
     ProductDescription   VARCHAR2(50),
     ProductFinish       VARCHAR2(20)
                        CHECK (ProductFinish IN ('Cherry', 'Natural Ash', 'White Ash',
                                                'Red Oak', 'Natural Oak', 'Walnut')),
     ProductStandardPrice DECIMAL(6,2),
     ProductLineID       INTEGER,
 CONSTRAINT Product_PK PRIMARY KEY (ProductID));

CREATE TABLE OrderLine_T
    (OrderID             NUMBER(11,0)    NOT NULL,
     ProductID           INTEGER         NOT NULL,
     OrderedQuantity     NUMBER(11,0),
 CONSTRAINT OrderLine_PK PRIMARY KEY (OrderID, ProductID),
 CONSTRAINT OrderLine_FK1 FOREIGN KEY (OrderID) REFERENCES Order_T(OrderID),
 CONSTRAINT OrderLine_FK2 FOREIGN KEY (ProductID) REFERENCES Product_T(ProductID));
```

Overall table definitions



Figure 6-6: Attributes and Data Types (Pine Valley Furniture Company)

Attribute (no space)	DATA TYPE	
CREATE TABLE Product_T		
(ProductID	NUMBER(11,0)	NOT NULL,
ProductDescription	VARCHAR2(50),	
ProductFinish	VARCHAR2(20)	
	CHECK (ProductFinish IN ('Cherry', 'Natural Ash', 'White Ash',	
	'Red Oak', 'Natural Oak', 'Walnut')),	
ProductStandardPrice	DECIMAL(6,2),	
ProductLineID	INTEGER,	
CONSTRAINT Product_PK PRIMARY KEY (ProductID));		



UNIVERSITY OF
MARYLAND

Figure 6-6: Required Attribute and Primary Key (Pine Valley Furniture Company)

```
CREATE TABLE Product_T
    (ProductID          NUMBER(11,0) NOT NULL,
     ProductDescription  VARCHAR2(50),
     ProductFinish       VARCHAR2(20)
                          CHECK (ProductFinish IN ('Cherry', 'Natural Ash', 'White Ash',
                                                    'Red Oak', 'Natural Oak', 'Walnut')),
     ProductStandardPrice DECIMAL(6,2),
     ProductLineID       INTEGER,
     CONSTRAINT Product_PK PRIMARY KEY (ProductID));
```

Non-nullable specification

Identifying primary key

Primary keys can never have NULL values – entity integrity.



UNIVERSITY OF
MARYLAND

Figure 6-6: Composite Primary Key (Pine Valley Furniture Company)

CREATE TABLE OrderLine_T		Non-nullable specification
(OrderID	NUMBER(11,0)	NOT NULL,
ProductID	INTEGER	NOT NULL,
OrderedQuantity	NUMBER(11,0),	
CONSTRAINT OrderLine_PK PRIMARY KEY (OrderID, ProductID),		Identifying primary key
CONSTRAINT OrderLine_FK1 FOREIGN KEY (OrderID) REFERENCES Order_T(OrderID),		
CONSTRAINT OrderLine_FK2 FOREIGN KEY (ProductID) REFERENCES Product_T(ProductID));		

Some primary keys are composite – composed of multiple attributes.



UNIVERSITY OF
MARYLAND

Figure 6-6: Default Value and Domain (Pine Valley Furniture Company)

```
CREATE TABLE Order_T
    (OrderID                NUMBER(11,0)    NOT NULL,
     OrderDate              DATE DEFAULT SYSDATE,
     CustomerID             NUMBER(11,0),
 CONSTRAINT Order_PK PRIMARY KEY (OrderID),
 CONSTRAINT Order_FK FOREIGN KEY (CustomerID) REFERENCES Customer_T(CustomerID));

CREATE TABLE Product_T
    (ProductID              NUMBER(11,0)    NOT NULL,
     ProductDescription      VARCHAR2(50),
     ProductFinish           VARCHAR2(20)
     Domain constraint CHECK (ProductFinish IN ('Cherry', 'Natural Ash', 'White Ash',
                                                'Red Oak', 'Natural Oak', 'Walnut')),
     ProductStandardPrice   DECIMAL(6,2),
     ProductLineID           INTEGER,
 CONSTRAINT Product_PK PRIMARY KEY (ProductID));
```

Figure 6-6: Foreign Keys and Relationships (Pine Valley Furniture Company)

```
CREATE TABLE Customer_T
```

(CustomerID	NUMBER(11,0)	NOT NULL,
CustomerName	VARCHAR2(25)	NOT NULL,
CustomerAddress	VARCHAR2(30),	
CustomerCity	VARCHAR2(20),	
CustomerState	CHAR(2),	
CustomerPostalCode	VARCHAR2(9),	

Primary key of
parent table

```
CONSTRAINT Customer_PK PRIMARY KEY (CustomerID));
```

```
CREATE TABLE Order_T
```

(OrderID	NUMBER(11,0)	NOT NULL,
OrderDate	DATE DEFAULT SYSDATE,	
CustomerID	NUMBER(11,0),	

Foreign key of
dependent table

```
CONSTRAINT Order_PK PRIMARY KEY (OrderID),
```

```
CONSTRAINT Order_FK FOREIGN KEY (CustomerID) REFERENCES Customer_T(CustomerID);
```

Data Integrity Controls

- **Referential integrity** – constraint that ensures that foreign key values of a table must match primary key values of a related table in 1:M relationships.
- Restricting:
 - Deletes of primary records.
 - Updates of primary records.
 - Inserts of dependent records.




UNIVERSITY OF
MARYLAND

ROBERT H. SMITH
SCHOOL OF BUSINESS

Figure 6-7: Ensuring Data Integrity through Updates

Relational integrity is enforced via the primary key to foreign key match



Restricted Update: A customer ID can only be deleted if it is not found in ORDER table.

```
CREATE TABLE CustomerT
    (CustomerID          INTEGER DEFAULT '999'    NOT NULL,
     CustomerName        VARCHAR(40)            NOT NULL,
     ...
    CONSTRAINT Customer_PK PRIMARY KEY (CustomerID),
    ON UPDATE RESTRICT);
```

Cascaded Update: Changing a customer ID in the CUSTOMER table will result in that value changing in the ORDER table to match.

```
... ON UPDATE CASCADE);
```

Set Null Update: When a customer ID is changed, any customer ID in the ORDER table that matches the old customer ID is set to NULL.

```
... ON UPDATE SET NULL);
```

Set Default Update: When a customer ID is changed, any customer ID in the ORDER tables that matches the old customer ID is set to a predefined default value.

```
... ON UPDATE SET DEFAULT);
```

Changing Tables

- **ALTER TABLE** statement allows you to change column specifications:

```
ALTER TABLE table_name alter_table_action;
```

- Table Actions:

```
ADD [COLUMN] column_definition  
ALTER [COLUMN] column_name SET DEFAULT default-value  
ALTER [COLUMN] column_name DROP DEFAULT  
DROP [COLUMN] column_name [RESTRICT] [CASCADE]  
ADD table_constraint
```

- Example (adding a new column with a default value):

```
ALTER TABLE CUSTOMER_T  
ADD COLUMN CustomerType VARCHAR2 (2) DEFAULT "Commercial";
```



UNIVERSITY OF
MARYLAND

Removing Tables

- DROP TABLE statement allows you to remove tables from your schema:

DROP TABLE Customer_T



UNIVERSITY OF
MARYLAND

INSERT Statement

- Adds one or more rows to a table.
- Inserting into a table:

```
INSERT INTO Customer_T VALUES  
(001, 'Contemporary Casuals', '1355 S. Himes Blvd.', 'Gainesville', 'FL', 32601);
```

- Inserting a record that has some null attributes requires identifying the fields that actually get data:

```
INSERT INTO Product_T (ProductID,  
ProductDescription, ProductFinish, ProductStandardPrice)  
VALUES (1, 'End Table', 'Cherry', 175, 8);
```

- Inserting from another table:

```
INSERT INTO CaCustomer_T  
SELECT * FROM Customer_T  
WHERE CustomerState = 'CA';
```


Creating Tables with IDENTITY Columns

```
CREATE TABLE Customer_T
(CustomerID INTEGER GENERATED ALWAYS AS IDENTITY
 (START WITH 1
 INCREMENT BY 1
 MINVALUE 1
 MAXVALUE 10000
 NO CYCLE),
CustomerName          VARCHAR2(25) NOT NULL,
CustomerAddress        VARCHAR2(30),
CustomerCity           VARCHAR2(20),
CustomerState          CHAR(2),
CustomerPostalCode     VARCHAR2(9),
CONSTRAINT Customer_PK PRIMARY KEY (CustomerID);
```

Introduced with SQL:2008

- Inserting into a table does not require explicit customer ID entry or field list:

INSERT INTO Customer_T VALUES ('Contemporary Casuals', '1355 S. Himes Blvd.', 'Gainesville', 'FL', 32601);



UNIVERSITY OF
MARYLAND

DELETE Statement

- Removes rows from a table.
- Delete certain rows:
DELETE FROM Customer_T WHERE CustomerState = 'HI';
- Delete all rows:
DELETE FROM Customer_T;



UNIVERSITY OF
MARYLAND

ROBERT H. SMITH
SCHOOL OF BUSINESS

UPDATE Statement

- Modifies data in existing rows:

```
UPDATE Product_T  
SET ProductStandardPrice = 775  
WHERE ProductID = 7;
```



UNIVERSITY OF
MARYLAND

MERGE Statement

```
MERGE INTO Product_T AS PROD
USING
(SELECT ProductID, ProductDescription, ProductFinish,
ProductStandardPrice, ProductLineID FROM Purchases_T) AS PURCH
  ON (PROD.ProductID = PURCH.ProductID)
WHEN MATCHED THEN UPDATE
  PROD.ProductStandardPrice = PURCH.ProductStandardPrice
WHEN NOT MATCHED THEN INSERT
  (ProductID, ProductDescription, ProductFinish, ProductStandardPrice,
  ProductLineID)
  VALUES(PURCH.ProductID, PURCH.ProductDescription,
  PURCH.ProductFinish, PURCH.ProductStandardPrice,
  PURCH.ProductLineID);
```

- Makes it easier to update a table ... allows combination of **INSERT** and **UPDATE** in one statement.
- Useful for updating master tables with new data.



UNIVERSITY OF
MARYLAND

ROBERT H. SMITH
SCHOOL OF BUSINESS

Schema Definition

- Control processing/storage efficiency:
 - Choice of indexes.
 - File organizations for base tables.
 - File organizations for indexes.
 - Data clustering.
 - Statistics maintenance.
- Creating indexes:
 - Speed up random/sequential access to base table data.
 - Example:
CREATE INDEX Name_IDX **ON** Customer_T(CustomerName).
 - This makes an index for the CustomerName field of the Customer_T table.



SELECT Statement

- Used for queries on single or multiple tables.
- Clauses of the **SELECT** statement:
 - **SELECT**: List the columns (and expressions) to be returned from the query
 - **FROM**: Indicate the table(s) or view(s) from which data will be obtained
 - **WHERE**: Indicate the conditions under which a row will be included in the result
 - **GROUP BY**: Indicate categorization of results
 - **HAVING**: Indicate the conditions under which a category (group) will be included
 - **ORDER BY**: Sorts the result according to specified criteria



UNIVERSITY OF
MARYLAND

ROBERT H. SMITH
SCHOOL OF BUSINESS

Figure 6-10: SQL Processing Order

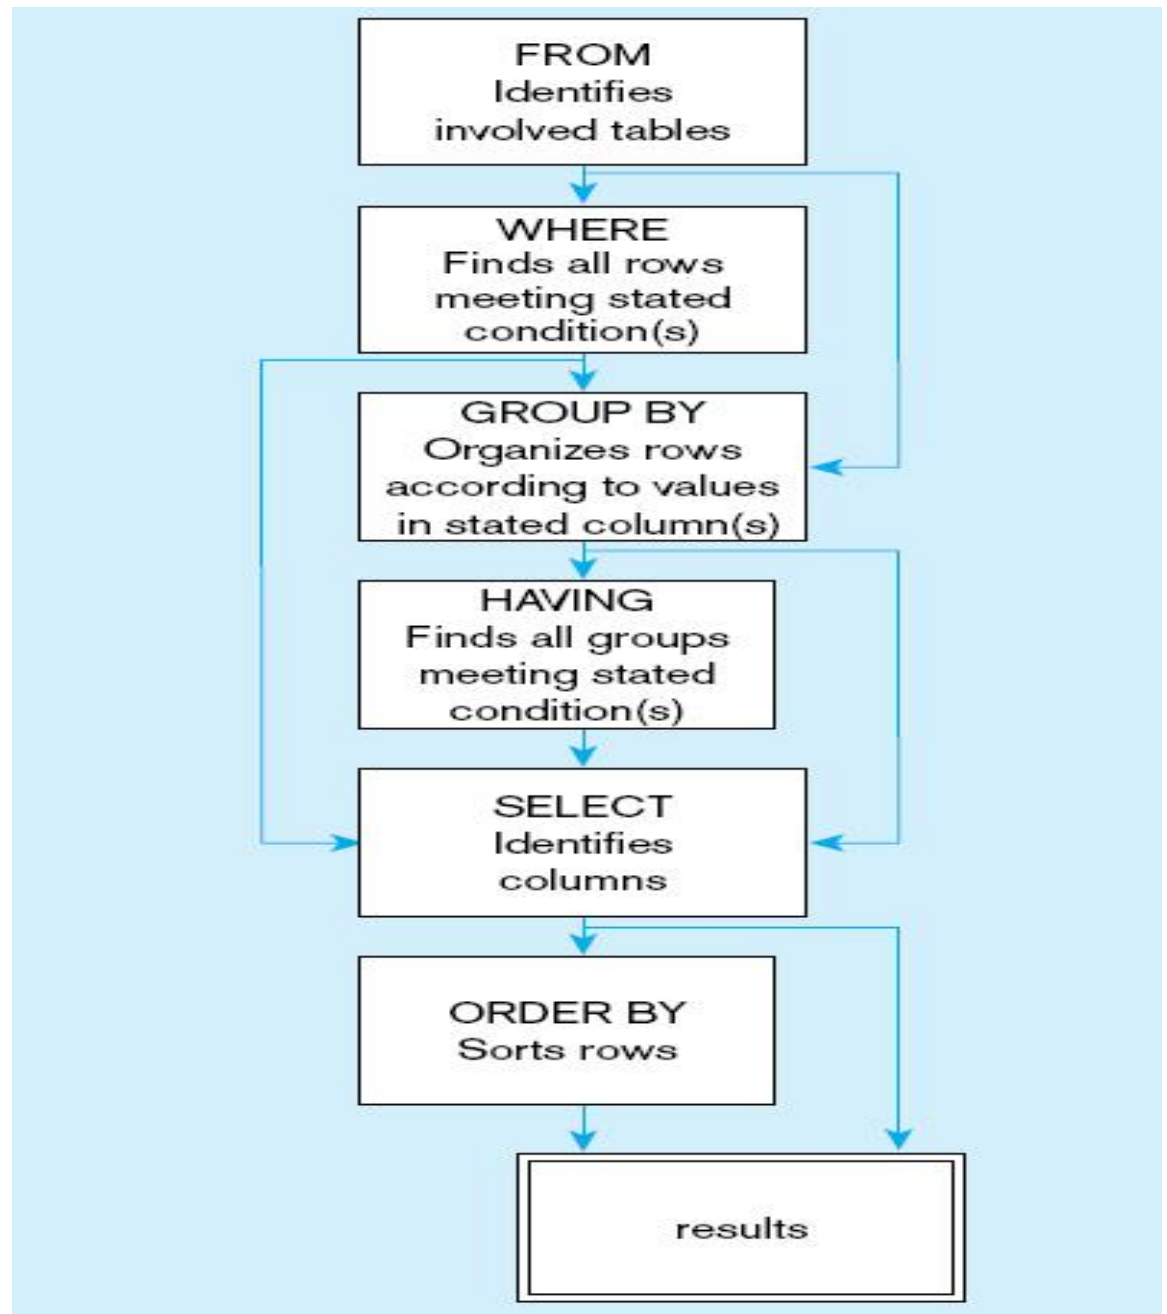


Table 6-3: Comparison Operators in SQL

- Find products with standard price less than \$275:

```
SELECT ProductDescription, ProductStandardPrice  
FROM Product_T  
WHERE ProductStandardPrice < 275;
```

TABLE 6-3 Comparison Operators in SQL

Operator	Meaning
=	Equal to
>	Greater than
>=	Greater than or equal to
<	Less than
<=	Less than or equal to
<>	Not equal to
!=	Not equal to

SELECT Example – Using Alias

- Alias is an alternative column or table name:

```
SELECT Cust.CustomerName AS Name,  
      Cust.CustomerAddress  
FROM Customer_V Cust  
WHERE Name = 'Home Furnishings';
```



UNIVERSITY OF
MARYLAND

SELECT Example – Using a Function

- Using the COUNT aggregate function to find totals:
SELECT COUNT(*) FROM Orderline_T
WHERE OrderID = 1004;
- Note: With aggregate functions you can't have single-valued columns included in the **SELECT** clause, unless they are included in the **GROUP BY** clause.



UNIVERSITY OF
MARYLAND

SELECT Example – Boolean Operators

- **AND, OR, and NOT** Operators for customizing conditions in WHERE clause:

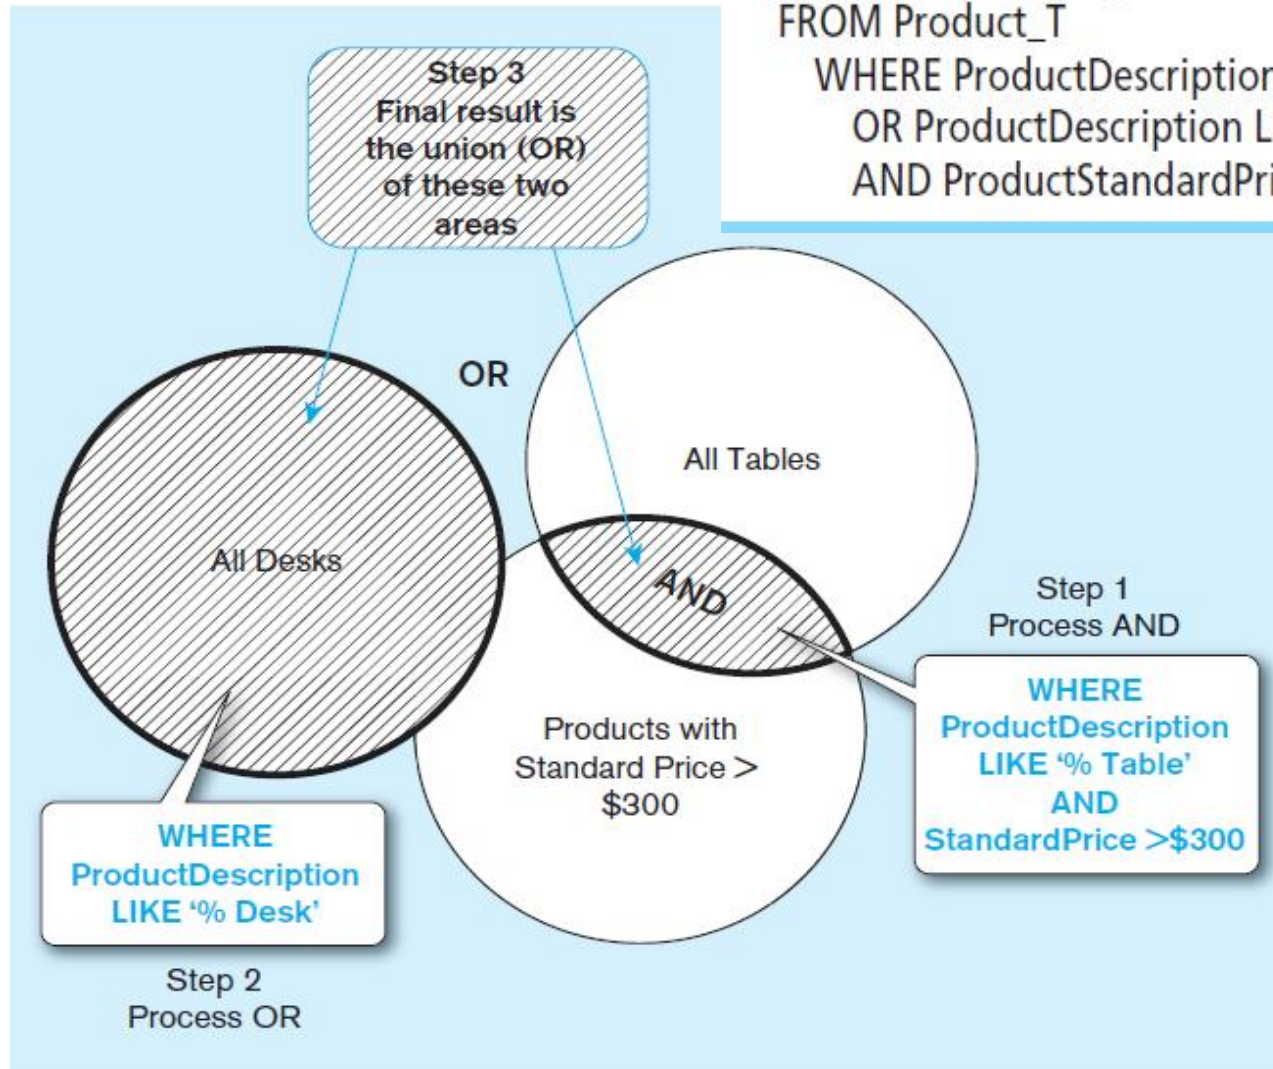
```
SELECT ProductDescription, ProductFinish, ProductStandardPrice  
FROM Product_T  
WHERE ProductDescription LIKE '%Desk'  
OR ProductDescription LIKE '%Table'  
AND ProductStandardPrice > 300;
```

- Note: The **LIKE** operator allows you to compare strings using wildcards. For example, the % wildcard in '%Desk' indicates that all strings that have any number of characters preceding the word 'Desk' will be allowed.



UNIVERSITY OF
MARYLAND

Figure 6-8: Boolean Query without Use of Parentheses



```
SELECT ProductDescription, ProductFinish, ProductStandardPrice  
FROM Product_T  
WHERE ProductDescription LIKE '%Desk'  
OR ProductDescription LIKE '%Table'  
AND ProductStandardPrice > 300;
```

By default,
processing order of
boolean operators
is NOT, then AND,
then OR.



UNIVERSITY OF
MARYLAND

ROBERT H. SMITH
SCHOOL OF BUSINESS

SELECT Example – BOOLEAN Operators

- With parentheses...these override the normal precedence of Boolean operators:

```
SELECT ProductDescription, ProductFinish, ProductStandardPrice  
FROM Product_T;  
WHERE (ProductDescription LIKE '%Desk'  
      OR ProductDescription LIKE '%Table')  
      AND ProductStandardPrice > 300;
```

- With parentheses, you can override normal precedence rules. In this case parentheses make the **OR** take place before the **AND**.

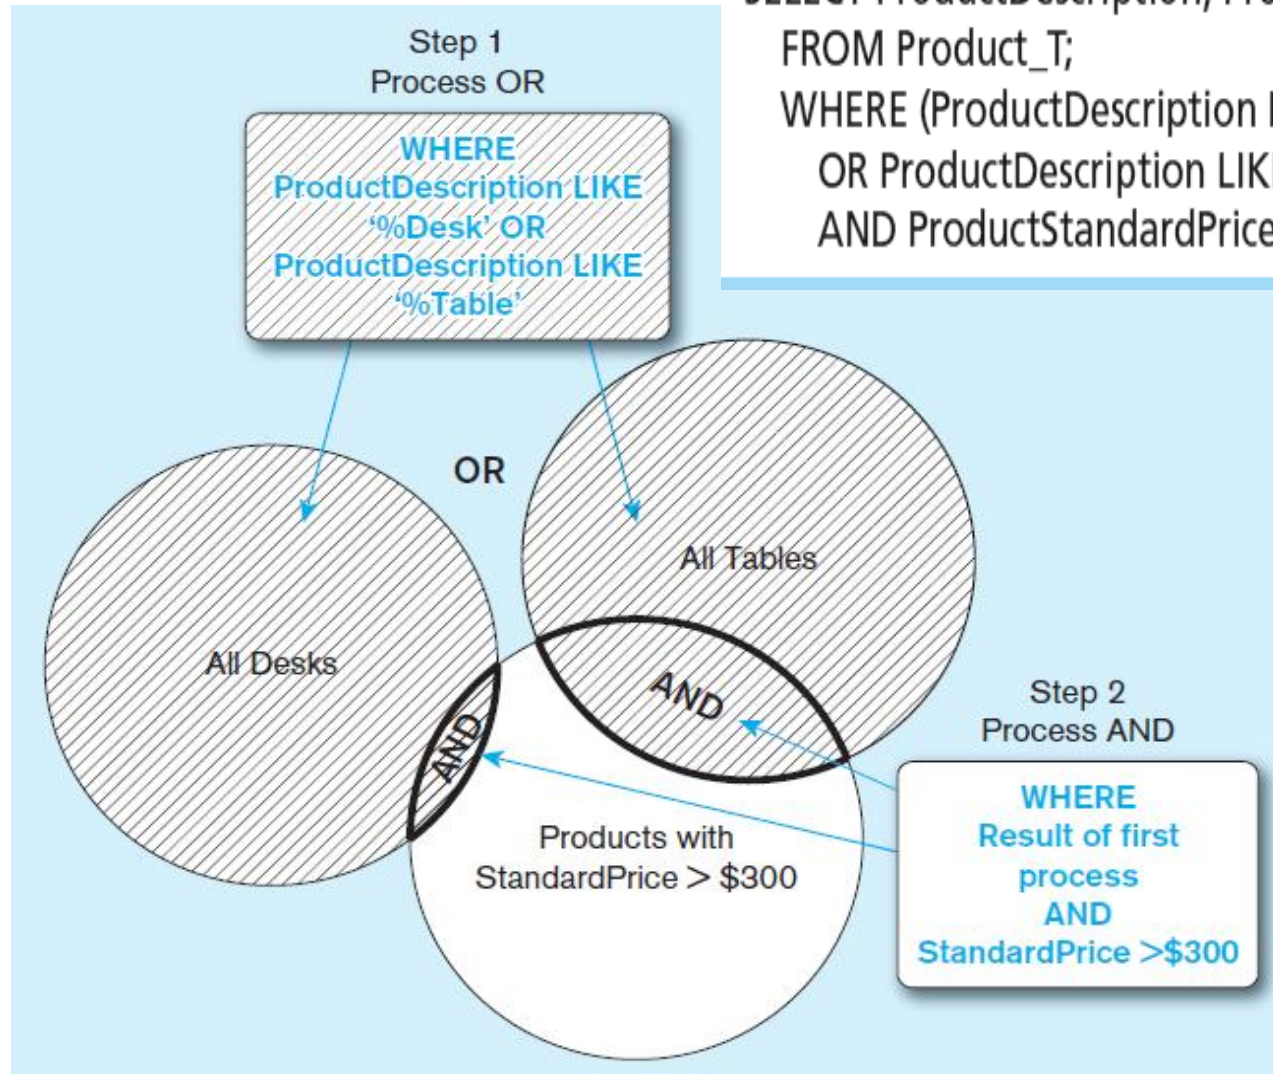


UNIVERSITY OF
MARYLAND

ROBERT H. SMITH
SCHOOL OF BUSINESS

Figure 6-9: Boolean Query with Use of Parentheses

```
SELECT ProductDescription, ProductFinish, ProductStandardPrice
FROM Product_T;
WHERE (ProductDescription LIKE '%Desk'
      OR ProductDescription LIKE '%Table')
      AND ProductStandardPrice > 300;
```



Sorting Results with ORDER BY Clause

- Sort the results first by CustomerState, and within a state by the CustomerName

```
SELECT CustomerName, CustomerCity, CustomerState  
FROM Customer_T  
WHERE CustomerState IN ('FL', 'TX', 'CA', 'HI')  
ORDER BY CustomerState, CustomerName;
```

- Note: The **IN** operator in this example allows you to include rows whose CustomerState value is either FL, TX, CA, or HI. It is more efficient than separate **OR** conditions.

Categorizing Results Using GROUP BY Clause

- For use with aggregate functions:
 - **Scalar aggregate:** Single value returned from SQL query with aggregate function.
 - **Vector aggregate:** Multiple values returned from SQL query with aggregate function (via GROUP BY).

```
SELECT CustomerState, COUNT (CustomerState)
FROM Customer_T
GROUP BY CustomerState;
```

- You can use single-value fields with aggregate functions if they are included in the **GROUP BY** clause.



UNIVERSITY OF
MARYLAND

ROBERT H. SMITH
SCHOOL OF BUSINESS

Qualifying Results by Categories Using the HAVING Clause

- For use with **GROUP BY**

```
SELECT CustomerState, COUNT (CustomerState)
FROM Customer_T
GROUP BY CustomerState
HAVING COUNT (CustomerState) > 1;
```

- Like a **WHERE** clause, but it operates on groups (categories), not on individual rows.
- Here, only those groups with total numbers greater than 1 will be included in final result.



UNIVERSITY OF
MARYLAND

ROBERT H. SMITH
SCHOOL OF BUSINESS

Using and Defining Views

- Views provide users controlled access to tables.
- **Base Table** – table containing the raw data.
- **Dynamic View:**
 - A “virtual table” created dynamically upon request by a user.
 - No data actually stored; instead data from base table made available to user.
 - Based on SQL **SELECT** statement on base tables or other views.
- **Materialized View:**
 - Copy or replication of data.
 - Data actually stored.
 - Must be refreshed periodically to match corresponding base tables.



UNIVERSITY OF
MARYLAND

ROBERT H. SMITH
SCHOOL OF BUSINESS

Sample CREATE VIEW

```
CREATE VIEW ExpensiveStuff_V
AS
    SELECT ProductID, ProductDescription, ProductStandardPrice
    FROM Product_T
    WHERE ProductStandardPrice > 300
    WITH CHECK OPTION;
```

- View has a name.
- View is based on a **SELECT** statement.
- **CHECK OPTION** works only for updateable views and prevents updates that would create rows not included in the view.



UNIVERSITY OF
MARYLAND

Advantages and Disadvantages of Views

- Advantages of views
 - Simplify query commands.
 - Assist with data security (but don't rely on views for security, there are more important security measures).
 - Enhance programming productivity.
 - Contain most current base table data.
 - Use little storage space.
 - Provide customized view for user.
 - Establish physical data independence.
- Disadvantages of views:
 - Use processing time each time view is referenced.
 - May or may not be directly updateable.



UNIVERSITY OF
MARYLAND

ROBERT H. SMITH
SCHOOL OF BUSINESS