

Data Processing and Analysis in Python

Lecture 14

Tabular Data and pandas



UNIVERSITY OF
MARYLAND

ROBERT H. SMITH
SCHOOL OF BUSINESS

DR. ADAM LEE

Pandas



- Panel Data System
- Built for the Python programming language
- An open source, BSD-licensed library (module)
- High-performance, easy-to-use data structures and data analysis tools
- Key components:
 - Series
 - DataFrame



UNIVERSITY OF
MARYLAND

ROBERT H. SMITH
SCHOOL OF BUSINESS

Use Pandas

```
# Import the whole module
```

```
>>> import pandas
```

```
# Import the whole module using alias
```

```
>>> import pandas as pd
```

```
# Import the whole module with all libraries
```

```
>>> from pandas import *
```

```
http://pandas.pydata.org/
```

```
>>> dir(pandas)
```

```
>>> dir(pandas.Series)
```

```
>>> dir(pandas.DataFrame)
```

```
>>> help(pandas)
```

```
>>> help(pandas.Series)
```

```
>>> help(pandas.DataFrame)
```



UNIVERSITY OF
MARYLAND

ROBERT H. SMITH
SCHOOL OF BUSINESS

Data Type: Series

- One-dimensional array like object containing data and labels (or index)
- Lots of ways to build a Series:

```
>>> srs=pd.Series(['a','b','c','d','e','f'])
```

```
>>> srs
```

```
0      a
```

```
1      b
```

```
2      c
```

```
3      d
```

```
4      e
```

```
5      f
```

```
dtype: object
```



UNIVERSITY OF
MARYLAND

ROBERT H. SMITH
SCHOOL OF BUSINESS

Series from Dictionary

- A dict-like container:

```
>>> dct={'a':100,'b':150,'c':200}
```

```
>>> srs=pd.Series(dct)
```

```
>>> srs
```

```
a      100
```

```
b      150
```

```
c      200
```

```
dtype: int64
```

```
>>> ses['b']
```

```
150
```



UNIVERSITY OF
MARYLAND

Series with Index

- A Series index can be specified:

- Single values can be selected by index
- Multiple values can be selected with multiple indexes

```
>>> srs=pd.Series([2,4,6],index=['a','b','c'])
```

```
>>> srs
```

```
a      2
```

```
b      4
```

```
c      6
```

```
dtype: int64
```

```
>>> srs['b']
```

```
4
```

```
>>> srs[['b','c']]
```

```
b      4
```

```
c      6
```

```
dtype: int64
```



UNIVERSITY OF
MARYLAND

ROBERT H. SMITH
SCHOOL OF BUSINESS

Series versus Dictionary

- Think of a Series as a fixed-length order dictionary
- However, unlike dictionary, index items do not have to be unique:

```
>>> srs=pd.Series([2,4,6,8],  
index=['a','b','a','b'])
```

```
>>> s['a']
```

```
a      2
```

```
a      6
```

```
dtype: int64
```

```
>>> s['b']
```

```
b      4
```

```
b      8
```

```
dtype: int64
```

dict

a	2
b	4

 →

a	6
b	8

pandas.Series

a	2
b	4

 →

a	2
b	4
a	6
b	8

key:value

Data: DataFrame

- Tabular data
- Rectangular data structure, a lot like an array
- Columns (Series) of different types
- Rows and columns act differently
- Can index by (column) labels as well as positions
- Handles missing data (NaN)
- Convenient plotting
- Fast operations with keys
- Lots of facilities for input/output



UNIVERSITY OF
MARYLAND

ROBERT H. SMITH
SCHOOL OF BUSINESS

Series to DataFrame

Series		Series		Series		DataFrame			
	pop		state		year		pop	state	year
0	18.8	0	FL	0	2010	0	18.8	FL	2010
1	19.1	1	FL	1	2011	1	19.1	FL	2011
2	9.7	2	GA	2	2008	2	9.7	GA	2008
3	9.7	3	GA	3	2010	3	9.7	GA	2010
4	9.8	4	GA	4	2011	4	9.8	GA	2011

index
(key)

columns
(values)

	A	B	C
1	pop	state	year
2	18.8	FL	2010
3	19.1	FL	2011
4	9.7	GA	2008
5	9.7	GA	2010
6	9.8	GA	2011

Create Data

- Creation with dict of equal-length lists:

```
>>> data={"state":["FL","FL","GA","GA","GA"],
          "year"  : [2010,2011,2008,2010,2011],
          "pop"   : [18.8,19.1, 9.7, 9.7, 9.8]}
```

```
>>> frame=pd.DataFrame(data)
```

```
>>> frame
```

	pop	state	year
0	18.8	FL	2010
1	19.1	FL	2011
2	9.7	GA	2008
3	9.7	GA	2010
4	9.8	GA	2011



UNIVERSITY OF
MARYLAND

Create Data

- Creation with dict of dicts:

```
>>> pop_data={"FL":{2010:18.8,
                    2011:19.1},
               "GA":{2008: 9.7,
                    2010: 9.7,
                    2011: 9.8}}
```

```
>>> pop=pd.DataFrame(pop_data)
```

```
>>> pop
```

	FL	GA
2008	NaN	9.7
2010	18.8	9.7
2011	19.1	9.8



UNIVERSITY OF
MARYLAND

Create List

- The initial set of baby names and birth rates:

```
>>> names = ["Bob", "Jane", "Mary", "John", "Mel"]
>>> births = [968, 155, 77, 578, 973]
>>> BabyDataSet = list(zip(names, births))
>>> BabyDataSet
[('Bob', 968), ('Jane', 155), ('Mary', 77),
 ('John', 578), ('Mel', 973)]
```



UNIVERSITY OF
MARYLAND

Create Data

- Create a Pandas DataFrame object df:

```
>>> df = pd.DataFrame(data=BabyDataSet,  
                        columns=["names", "births"])
```

```
>>> df
```

	names	births
0	Bob	968
1	Jane	155
2	Mary	77
3	John	578
4	Mel	973

- You can think of df holding the contents of BabyDataSet in a format similar to an SQL table or an Excel spreadsheet



UNIVERSITY OF
MARYLAND

Save Data into File

- `to_csv()` Export DataFrame to a comma-separated values (CSV) file:

```
>>> df.to_csv("filename.csv")
```

```
>>> df.to_csv("filename.csv", index=False,  
header=False)
```

- The file will be saved in the same location as your python file unless specified otherwise



UNIVERSITY OF
MARYLAND

Load Data from File

- `read_csv()` To import from the CSV file:

```
>>> filepath = "filename.csv"
```

```
>>> df = pd.read_csv(filepath, header=None)
```

```
>>> df
```

	0	1
0	Bob	968
1	Jane	155
2	Mary	77
3	John	578
4	Mel	973



UNIVERSITY OF
MARYLAND

Load Data from File

- To give columns specific names, set parameter "names":

```
>>> filepath = "filename.csv"
```

```
>>> df = pd.read_csv(filepath, names=["names",  
"births"])
```

```
>>> df
```

	names	births
0	Bob	968
1	Jane	155
2	Mary	77
3	John	578
4	Mel	973



UNIVERSITY OF
MARYLAND

Select Data

- `df.head()` returns the first 5 records
- `df.head(n)` returns the first n records
- `df.tail()` returns the last 5 records
- `df.tail(n)` returns the last n records



UNIVERSITY OF
MARYLAND

Prepare Data

- To find all unique records of "names" column:

```
>>> df["names"].unique()  
array(['Bob', 'Jane', 'Mary', 'John', 'Mel'],  
      dtype=object)
```

- To obtain all descriptive statistics:

```
>>> df["names"].describe()  
count          5  
unique          5  
top            Jane  
freq           1  
Name: Names, dtype: object
```



UNIVERSITY OF
MARYLAND

Analyze Data

- To find baby name with the highest birth rate:

```
>>> sorted = df.sort_values(["births"],  
ascending=False)
```

```
>>> sorted
```

	names	births
4	Mel	973
0	Bob	968
3	John	578
1	Jane	155
2	Mary	77

```
>>> sorted.head(1)
```

	names	births
4	Mel	973



UNIVERSITY OF
MARYLAND

Analyze Data

- `df["names"]` returns the entire list of baby names, i.e. the entire "Names" column
- `df["births"]` returns the entire list of births, i.e. the entire "Births" column
- `df["births"].max()` returns the maximum value found in the "Births" column
- `[df["births"] == df["births"].max()]` equals to [Find all records in the "Births" column, where it equals to the maximum, i.e. 973.]
- `df["names"][df["births"] == df["births"].max()]` equals to [Find all records in the "Names" column, where the "Births" column equals to the maximum, i.e. 973.]



UNIVERSITY OF
MARYLAND

Column Operations

```
>>> data = range(0,10)
>>> df = pd.DataFrame(list(data))
>>> df.columns = ["num"]
>>> df["new"] = 5
>>> df["new"] = df["new"] + 1
```

	num	new
0	0	6
1	1	6
2	2	6
3	3	6
4	4	6
5	5	6
6	6	6
7	7	6
8	8	6
9	9	6

Column Operations

```
>>> data = range(0,10)
>>> df = pd.DataFrame(list(data))
>>> df.columns = ["num"]
>>> df["three"] = 3
>>> df["col"] = df["num"]
```

	num	three	col
0	0	3	0
1	1	3	1
2	2	3	2
3	3	3	3
4	4	3	4
5	5	3	5
6	6	3	6
7	7	3	7
8	8	3	8
9	9	3	9

Column Operations

```
# ...  
>>> ind = ['a', 'b', 'c', 'd', 'e',  
           'f', 'g', 'h', 'i', 'j']  
>>> df.index = ind
```

	num	three	col
a	0	3	0
b	1	3	1
c	2	3	2
d	3	3	3
e	4	3	4
f	5	3	5
g	6	3	6
h	7	3	7
i	8	3	8
j	9	3	9

Column Operations

```
# ...  
>>> df.loc['a']  
Rev 0  
test 3  
col 0  
Name: a, dtype: int64
```

```
# loc[] both inclusive  
>>> df.loc['a':'d']
```

```
# iloc[] inclusive:exclusive  
>>> df.iloc[0:3]
```

	num	three	col
a	0	3	0
b	1	3	1
c	2	3	2
d	3	3	3

	num	three	col
a	0	3	0
b	1	3	1
c	2	3	2



UNIVERSITY OF
MARYLAND

Column Operations

```
# ...
```

```
>>> df[["num", "three"]]
```

```
# df.ix[rows, columns]
```

```
>>> df.ix[0:3, "num"]
```

```
>>> df.ix[5:8, "col"]
```

```
>>> df.ix[:3, ["col", "three"]]
```

a	0
b	1
c	2

	num	three
a	0	3
b	1	3
c	2	3

f	5
g	6
h	7
i	8
j	9

	num	three
a	0	3
b	1	3
c	2	3
d	3	3
e	4	3
f	5	3
g	6	3
h	7	3
i	8	3
j	9	3

GroupBy Operations

```
>>> data = {"letter": ['a', 'a', 'b', 'b', 'c'],  
            "one": [1, 1, 1, 1, 1],  
            "two": [2, 2, 2, 2, 2]}  
  
>>> df = pd.DataFrame(data)  
# create group object  
>>> letter = df.groupby("letter")  
# apply sum  
>>> letter.sum()
```

	one	two
letter		
a	2	4
b	2	4
c	1	2

	letter	one	two
0	a	1	2
1	a	1	2
2	b	1	2
3	b	1	2
4	c	1	2

GroupBy Operations

```
>>> letter_one =  
df.groupby(["letter", "one"]).sum()  
>>> letter_one =  
df.groupby(["letter", "one"], as_index=False).sum()  
( )
```

	letter	one	two
0	a	1	4
1	b	1	4
2	c	1	2

		two
letter	one	
a	1	4
b	1	4
c	1	2

Missing Values:

numpy.nan and pandas.NA

```
>>> import numpy as np
>>> df = pd.DataFrame(np.random.randn(5, 3),
                       index=['a', 'c', 'e', 'f', 'h'],
                       columns=["one", "two", "three"])
```

```
>>> df.loc['a', "two"] = np.nan
```

```
>>> df
```

	one	two	three
a	0.405857	NaN	2.291062
c	-0.491699	1.178464	-1.951439
e	2.649431	-1.938993	-0.827796
f	1.235505	-0.726769	-0.146558
h	2.837826	0.505976	-0.489474



UNIVERSITY OF
MARYLAND

Detect Missing Values

- `pandas.isnull()` Detect missing values for an array-like object
- `pandas.Series.isnull()` Detect missing values in a Series
- `pandas.DataFrame.isnull()` Detect missing values in a DataFrame

```
>>> df.isnull()
      one    two  three
a  False   True  False
c  False  False  False
e  False  False  False
f  False  False  False
h  False  False  False
```



UNIVERSITY OF
MARYLAND

Fill Missing Values

```
>>> df.fillna(0)
```

	one	two	three
a	0.405857	0.000000	2.291062
c	-0.491699	1.178464	-1.951439
e	2.649431	-1.938993	-0.827796
f	1.235505	-0.726769	-0.146558
h	2.837826	0.505976	-0.489474



UNIVERSITY OF
MARYLAND

Query Data

- Use the **query(expression)** function to embed boolean expressions on columns within quotes:

```
>>> df.query("one>0")
```

	one	two	three
a	0.405857	NaN	2.291062
e	2.649431	-1.938993	-0.827796
f	1.235505	-0.726769	-0.146558
h	2.837826	0.505976	-0.489474

```
>>> df.query("one>0 & two>0")
```

	one	two	three
h	2.837826	0.505976	-0.489474



UNIVERSITY OF
MARYLAND

ROBERT H. SMITH
SCHOOL OF BUSINESS

Query Data

- You can apply any function to the columns:

```
>>> df.apply(lambda x: x.max()-x.min())
one          3.329524
two          3.117457
three        4.242501
dtype: float64
```

- You can apply any function to element-wise:

```
>>> df.applymap(np.sqrt)
      one      two      three
a  0.637069      NaN  1.513626
c      NaN  1.085571      NaN
e  1.627707      NaN      NaN
f  1.111533      NaN      NaN
h  1.684585  0.711320      NaN
```



UNIVERSITY OF
MARYLAND

Pandas - Visualization

<http://pandas.pydata.org/pandas-docs/stable/visualization.html>

- Use the standard convention for referencing the matplotlib API
- Provide the basics in pandas to easily create decent looking plots



UNIVERSITY OF
MARYLAND